

# BEYOND THE BASE: EXPONENT ALLOCATION IN ROPE

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Rotary position embeddings specify both a frequency range and an allocation of a finite number of rotary pairs within that range. Writing a log-frequency table as endpoints  $(a, R)$  together with normalized interior exponents  $z$ , we study  $z$  separately from the sampled range  $R$  and from the weights  $M$  that learn to use the resulting basis. In paired 151.9M training runs with identical frequency endpoints, changing only the 30 interior exponents lowers extrapolation NLL at every tested extended length in all three seeds. The ordering reverses under runtime range re-targeting, and weights-by-table crossings show strong dependence on the training allocation. We characterize the supplied basis through full sine-cosine subspaces, and show by counterexample that three classes of static surrogate fail as standalone predictors, so  $z$  cannot be ranked without  $R$  and  $M$ . Two constructions follow as consequences rather than optimality claims: an analytic Cosh allocation from a convex design criterion, and a native-relative profile for frozen checkpoints. On five OLMo long-context QA tasks the latter raises task-equal F1 from 21.62% to 25.44% against MrRoPE-Pro, while a static Qwen table-amplitude configuration improves 64K RULER-13 by 6.09 points over the tested YaRN configuration. Together, these results show that a finite RoPE system is not described by its range alone: support, interior allocation and learned slot compatibility are three coordinates that cannot be merged, and an allocation carries no ranking apart from the other two.

## 1 INTRODUCTION

Rotary position embeddings (RoPE) assign a frequency to each two-dimensional query-key coordinate pair (Su et al., 2024). Its standard geometric table is parameterized by a base, controlling sampled positional scales (Liu et al., 2024; Xu et al., 2024). A finite table, however, makes another choice: how its rotary pairs occupy that range (Oka et al., 2026b; Tian et al., 2026). What remains necessary is to distinguish the causal effect of interior allocation from range selection, from the learning process, and from the compatibility of an installed table with mature weights.

Three questions organize the study. What changes in the positional basis when interior exponents are redistributed? Does this redistribution affect learned behavior when the frequency range is held strictly fixed? And how does an existing checkpoint respond to adjustments of its native exponents? These questions connect geometric analysis, controlled interventions, and practical table design.

Figure 1 shows the basic intervention: decompose the log-frequency table into endpoints  $(a, R)$  and normalized interior exponents  $z$ . In paired 151.9M training runs ( $L_{\text{train}} = 256, K = 32$ ), holding endpoints identical while changing only the 30 interior exponents lowers extrapolation NLL across all three seeds. However, re-targeting the runtime range reverses this ordering, and frozen weights strongly prefer runtime tables derived from their training allocation. Together, these controls isolate interior allocation as an independent design variable while establishing its action boundaries with range and learned representations.

The positional geometry provides a way to analyze these distributions. Each rotary pair supplies an entire sine-cosine subspace. Canonical correlations compare these subspaces without privileging a content phase, exposing the common low-frequency limit  $\text{span}\{1, \Delta\}$ . Thus many slow pairs supply few distinct positional directions over a bounded interval. The learned content coefficients determine how the model uses those directions.




Figure 1: **Interior exponents change learned behavior at fixed frequency range.** Left: two allocations of the same 32 rotary pairs with identical endpoints. Right: paired 151.9M models trained at length 256; changing the 30 interior exponents lowers extrapolation NLL in every seed. Thin lines show the three training seeds and the heavy line their mean.

We evaluate concrete allocation designs in two regimes. Before or during learning, a convex design criterion produces an analytic Cosh family whose quantiles give a fixed table, evaluated in scarce rotary channels, continued training, and matched 8B adaptation. For frozen checkpoints, we express modifications as displacements from native exponents. A static Qwen table–amplitude configuration improves 64K RULER-13 by 6.09 points over the tested YaRN configuration; a  $2 \times 2$  table  $\times$  gain factorial separating the two factors is registered. A boundary-matched adjustment raises OLMo mean token F1 from 21.62% to 25.44% *against MrRoPE-Pro* across five natural-QA tasks (631 untruncated prompts); the official-YaRN baseline on that same panel is registered and not yet run. Checkpoint preferences differ across lengths, making learned compatibility central to allocation design.

**Contributions.** We (i) split a log-frequency table into endpoints  $(a, R)$  and normalized interior exponents  $z$ , and isolate  $z$  causally through paired training at strictly fixed range; (ii) establish the operating boundaries of that variable—the ranking of  $z$  reverses under runtime range retargeting, and frozen weights prefer tables derived from their own training allocation, so no allocation is ranked by the table alone; (iii) characterize supplied sine–cosine subspace geometry exactly, and show by measured counterexample that three classes of static surrogate fail as standalone predictors; and (iv) give two constructive witnesses, an analytic Cosh allocation for the learning regime and a native-relative profile for frozen deployment, which show that this design space yields real gains without claiming a universal optimum.

## 2 EXPONENT ALLOCATION

Let  $K = d_{\text{rot}}/2 \geq 2$  be the number of rotary pairs. Standard RoPE uses

$$\omega_k = b^{-\phi_k}, \quad \phi_k = \frac{k}{K}, \quad k = 0, \dots, K-1, \quad (1)$$

for a base  $b > 1$ . For an arbitrary strictly ordered positive table  $\omega_0 > \dots > \omega_{K-1} > 0$ , write

$$x_k = -\log \omega_k = a + Rz_k, \quad z_k = \frac{x_k - x_0}{x_{K-1} - x_0}, \quad z_0 = 0, \quad z_{K-1} = 1. \quad (2)$$

The pair  $(a, R)$  specifies the sampled log-frequency support, while  $z$  specifies how the finite rotary budget occupies that support. Every geometric table has  $z_k = k/(K-1)$ , irrespective of its base. Thus changing  $z$  while fixing  $(a, R, K)$  is an intervention on interior allocation. The standard grid  $k/K$  in (1) and the endpoint-normalized shape  $k/(K-1)$  are two descriptions of the same geometric table, not interchangeable installation rules.

A continuous density over normalized exponents can be discretized through its quantiles. At finite  $K$ , midpoint quantiles, endpoint-inclusive quantiles, and affine anchoring produce different installed

frequency tensors. We specify the grid and anchoring convention for each protocol. In the fixed-support experiments, the nonuniform quantiles are affinely mapped to the baseline’s actual sampled endpoints. In mature checkpoints, we instead express an installed table through its displacement from the native tensor; §6 develops that form. These conventions prevent nominally identical bases or extension factors from concealing changes in the actual sampled range.

### 3 CONTROLLED EFFECTS OF EXPONENT ALLOCATION

We begin with interventions that distinguish exponent shape, sampled range, and the weights learned under an allocation. The Cosh quantiles used here are constructed in §4.4, with endpoint anchoring keeping their range fixed; an exponential deformation supplies a second family.

#### 3.1 CONTROLLED EFFECT AT FIXED RANGE AND RANGE INTERACTION

We first ask whether interior allocation affects learned behavior when the sampled range is exactly fixed. We train paired 151.9M models with  $K = 32$  rotary pairs at sequence length 256. Within each of three seeds, architecture, initialization, token order, optimizer, schedule, and the 499,974,144-token budget are identical. The geometric arm uses the FMRoPE training-base rule  $b = L_{\text{train}}$  (Oka et al., 2026b); the Cosh arm anchors its  $\tau = 4$  quantiles to the same endpoints. Only the 30 interior exponents differ.

At the fixed training support, the mean Cosh-minus-geometric tail-NLL differences at lengths 256, 512, 1024, and 2048 are  $+0.026, -0.281, -0.176, -0.146$ . All three seeds improve at each extended length (Fig. 1). Interior allocation thus changes learned generalization without any scalar base shift. However, retargeting the log-frequency span to each evaluation length reverses the ordering to  $+0.060, +0.227, +0.460$  across all three seeds. This reversal demonstrates that allocation cannot be judged independently of operating range: favorable fixed-support orderings do not survive arbitrary range dilation (Appendix C).

#### 3.2 CROSS-SHAPE FACTORIAL ACROSS CONFIGURATIONS

A 50.9M factorial varies two bases, two training lengths, three head dimensions, and three paired seeds, retaining endpoints throughout. At the reference Cosh strength, the weighted OOD objective improves in 7/12 configurations; preassigned  $1.25\times$  improves 10/12, and matched exponential improves 9/12. The aggregate reference-Cosh-minus-exponential contrast is small and uncertain ( $+0.00074$  NLL, 95% interval  $[-0.006, 0.008]$ ; Table 13). These results establish allocation effects across configurations while bounding claims about the unique superiority of any single shape.

#### 3.3 WEIGHT COMPATIBILITY AND TABLE CROSSINGS

The frequency table defines a basis in which query and key projections are learned. In a 50M pair with perplexities near 7, installing Cosh frequencies into Geo-trained weights raises PPL to 76.20 and the reverse swap to 23.05 (Fig. 9). A separate crossing at 151.9M repeats the effect across two training seeds using derived factor-four runtime tables: each frozen weight set strictly prefers the table derived from its own training allocation. This distinguishes learned compatibility from static basis geometry, without claiming that frequency modifications are universally harmful.

## 4 POSITIONAL GEOMETRY OF EXPONENT DISTRIBUTIONS

### 4.1 FULL-PAIR GEOMETRY AND EFFECTIVE RANK IDENTITY

For a fixed content pair, a rotary channel contributes  $C \cos(\omega\Delta) + D \sin(\omega\Delta)$  to an attention logit, where  $\Delta$  is relative token separation. Its positional object is therefore the two-dimensional function space

$$V_\omega = \text{span}\{\cos(\omega\Delta), \sin(\omega\Delta)\}. \quad (3)$$

Let  $x_\omega(\Delta) = [\cos(\omega\Delta), \sin(\omega\Delta)]$ , and define self- and cross-Gram matrices  $S_\omega = \mathbb{E}[x_\omega^\top x_\omega]$  and  $H_{\omega\nu} = \mathbb{E}[x_\omega^\top x_\nu]$  under a declared separation measure. For nonsingular  $S_\omega, S_\nu$ , the singular values

of whitened cross-Gram  $Q_{\omega\nu} = S_{\omega}^{-1/2} H_{\omega\nu} S_{\nu}^{-1/2}$  are the canonical correlations of the two spaces. Their average squared value,

$$c_{\omega\nu} = \frac{1}{2} \|Q_{\omega\nu}\|_F^2 \in [0, 1], \quad (4)$$

is invariant under invertible basis changes within either pair, measuring how much of the positional basis the frequencies share. Stacking all block-whitened pairs into  $\Gamma$  (with identity diagonal blocks), if  $\bar{c}$  is the mean pairwise squared canonical overlap, its Rényi-2 effective rank satisfies

$$r_2(\Gamma) = \frac{(\text{tr } \Gamma)^2}{\text{tr}(\Gamma^2)} = \frac{2K}{1 + (K-1)\bar{c}}. \quad (5)$$

This is an exact identity for the specified positional design matrix, not a downstream model accuracy bound.

## 4.2 SHARED SLOW-FREQUENCY LIMIT

**Proposition 1** (Shared slow-frequency subspace). *Under uniform measure on  $[0, L]$ , let  $x = \omega L$ . As  $x \rightarrow 0$ ,  $\cos(xt) \rightarrow 1$  and  $\sin(xt)/x \rightarrow t$ , so  $V_{\omega}$  converges to  $\text{span}\{1, \Delta\}$ . For frequencies with  $\epsilon = \max\{|\omega L|, |\nu L|\} \rightarrow 0$ , squared chordal separation satisfies  $2 - \|Q_{\omega\nu}\|_F^2 = \mathcal{O}(\epsilon^4)$ .*

In standard grids ( $b = 500,000$ ,  $K = 64$ ,  $L = 4096$ ), the 23 pairs with  $\omega L \leq 1$  supply 46 coordinate dimensions but a block-whitened rank  $r_2 = 2.00013$ . In raw unwhitened coordinates, the weak eigenvalue scales as  $\lambda_-(S_{\omega}) \approx (\omega L)^2/12$ . This motivates studying how tables distribute positional directions, though static rank does not guarantee downstream extrapolation.

## 4.3 WHAT A SCHEDULE-ONLY SUMMARY CAN SEE

The separation between basis geometry and downstream performance is formal:

**Theorem 2** (Static surrogates cannot rank allocations). *Let  $\Phi(\Omega)$  depend only on the frequency schedule and let  $\mathcal{R}$  be measured downstream capability. (i) Permutation-invariant  $\Phi$  cannot separate the two most behaviourally distant tables we can construct. For any non-trivial permutation  $\pi$ , invariance gives  $\Phi(\pi\Omega) = \Phi(\Omega)$ , while by Theorem 7  $\Omega$  and  $\pi\Omega$  admit no position-independent invertible reparameterization. On the evaluated checkpoints the measured risk gap is large:  $1 \times$  tail NLL rises from 3.10423 to 6.86493 and 64K macro accuracy falls from 0.7000 to 0.0000. So no such  $\Phi$  is a standalone predictor of  $\mathcal{R}$ . (ii) The remaining schedule-only surrogates fail on named witnesses. A discrete scale-orbit count shifts by  $\Delta q = 58$  under 1-ULP jitter with  $\Delta\text{NLL} = +10^{-5}$ ; and curvature ill-conditioning ( $\kappa(F_N) \gg 10^4$ ) concentrates sensitivity at knee slot 19, so an MAE = 0.0012 coordinate error fails pre-registered retention gates.*

*Remark 3.* Part (i) is structural, not anecdotal, and this is the point of coupling it to Theorem 7: the schedule’s only position-independent equivalence is permutation, and permutation is exactly the transformation the trained weights do not survive. A summary  $\Phi$  must therefore be invariant to something; the only invariance the schedule offers is the fatal one. We do *not* claim non-identifiability for every schedule-dependent  $\Phi$ — $\Phi(\Omega) = \Omega$  is injective—which is precisely why the useful question is not “which metric is best” but “what can a schedule-only summary see at all”.

Table 30 summarizes these boundaries, isolating what the geometric theory proves from what requires downstream empirical validation.

## 4.4 AN ANALYTIC ALLOCATION CONSTRUCTION

To allocate samples over normalized exponent space  $[0, 1]$ , let  $S_{\rho}(t) = \int_t^1 \rho(\phi) d\phi$ , and consider the convex design criterion

$$\mathcal{C}_{\text{app}}[\rho] = \frac{\alpha}{2} \int_0^1 \rho(\phi)^2 d\phi + \frac{\beta}{2} \int_0^1 S_{\rho}(t)^2 dt, \quad \text{s.t. } \int_0^1 \rho = 1, \rho \geq 0. \quad (6)$$

The first term penalizes concentration; the second penalizes cumulative mass beyond slow-end thresholds. These are explicit design preferences; we do not identify this functional with complete trained-model loss.

**Theorem 4** (Cosh allocation). For  $\alpha > 0$  and  $\beta \geq 0$ , the unique minimizer of (6) over positive  $\rho \in C^2([0, 1])$  with unit integral is

$$\rho_\tau(\phi) = \frac{\tau \cosh[\tau(1 - \phi)]}{\sinh \tau}, \quad \tau = \sqrt{\beta/\alpha}, \quad (7)$$

recovering uniform density as  $\tau \rightarrow 0$ . Inverting the CDF yields the explicit quantile map

$$\phi_k(\tau) = 1 - \frac{1}{\tau} \operatorname{arcsinh}((1 - u_k) \sinh \tau), \quad u_k = \frac{k + 1/2}{K}. \quad (8)$$

We call this construction EVQ-COSH. It is an analytic allocation family, not a claim of universal optimality. Once  $\tau$ , grid, and endpoint anchoring are specified, installation evaluates the table once without changing the rotary operator or trainable parameter count.

## 5 MODEL EFFECTS OF AN ANALYTIC ALLOCATION

We evaluate the analytic construction in matched learning settings before turning to frozen deployment; each comparison specifies its own regime, grid, and endpoint.

### 5.1 LANGUAGE MODELING AND RETRIEVAL

**Scarce rotary channels.** A 432M MLA-style model allocates only 16 pairs to explicit rotation (DeepSeek-AI, 2024a;b), with both arms trained on 500M tokens at length 8192. Across three seeds EVQ-COSH reduces mean 16K perplexity from 138.8 to 95.6, with 8K near-equal at 35.4 versus 35.8, testing allocation when few coordinates carry explicit rotation (Table 1). Plain-table gains hold at  $2\times$  the training window; beyond it the training-free blend closes and then passes them (full curve in Appendix G.1).

**Continued full-parameter training.** The 750M experiment continues training from a 2K checkpoint for 500M tokens at length 4096, the Cosh arm preserving endpoints. At 16K perplexity falls from 45.1 to 24.4, with 4K near-equal (22.0 vs 22.3), and across 40 associated 8K passkey cases greedy exact match rises from 0% to 77.5%, while the teacher-forced probe stays saturated for both arms, separating probability-based retrieval from complete generation (Table 1; Appendix B).

Table 1: **Model effects of the Cosh exponent table.** Mean perplexities within matched protocols, lower better. *Geo* uses the midpoint grid  $u_k = (k + \frac{1}{2})/K$ ; each pair shares grid, initialization, data order, optimizer and budget. Last row: training-free post-hoc frequency blend, applied to both arms at evaluation.

| Model              | Seeds | Train length | Eval. length | Geo   | EVQ-COSH    |
|--------------------|-------|--------------|--------------|-------|-------------|
| 432M MLA           | 3     | 8K           | 8K           | 35.4  | 35.8        |
|                    |       |              | 16K          | 138.8 | <b>95.6</b> |
| + wavelength blend | 3     | 8K           | 16K          | 117.9 | <b>71.1</b> |
| 750M continuation  | 1     | 4K           | 4K           | 22.0  | 22.3        |
|                    |       |              | 16K          | 45.1  | <b>24.4</b> |

**Matched Llama-3-8B adaptation.** We compare Native and EVQ-COSH under the same 300-step rank-64 LoRA adaptation at length 8K (Hu et al., 2022; Grattafiori et al., 2024) across 24 frozen temporal text packs. Native/EVQ-COSH PPL is 6.82/10.07 at 8K, 108.96/24.07 at 16K, and 991.48/127.91 at 32K (Fig. 8). We characterize this as an adaptation length-transfer trade-off rather than lossless extension: installed allocation changes extrapolation under matched adaptation budgets.

**Source use and task-family continuation.** On ten true-16K passkey cases, removing decode attention to the remote gold block increases answer-token NLL by 1.5055 in EVQ-COSH versus  $-0.0095$  in Native, confirming reliance on distant source access. A separate continuation trains both arms on RULER generator families at length 8K and evaluates new rows: at 16K official macro is 0.30% for Native-LoRA and 14.03% for EVQ-COSH-LoRA (exact match 0% vs 1.54%), while at 8K it is 94.44/77.60% (Appendix E.3).

**Composition and further settings.** In the 454M three-seed experiment, applying the same index-based smooth-ramp scaler to both learned tables gives 16K PPL of 157.7 for Geo and 107.5 for EVQ-COSH (Appendix B.3). OLMo scale checks and video-DiT comparisons appear in Apps. E.1 and B.6.

## 6 ADJUSTING EXPONENTS IN FROZEN MODELS

The crossings in §3.3 show that trained weights are co-adapted to their exponent basis, motivating adjustments defined relative to native exponents.

### 6.1 A COMMON COORDINATE FOR EXPONENT ADJUSTMENTS

For native frequencies  $\omega_k^N$  and installed frequencies  $\omega'_k$ , define

$$d_k = \log \frac{\omega_k^N}{\omega'_k}, \quad \omega'_k = \omega_k^N e^{-d_k}, \quad \phi'_k - \phi_k^N = \frac{d_k}{\log b}. \quad (9)$$

$$d_k = \begin{cases} -\log(1 - w_k + w_k/s), & \text{frequency blend,} \\ m_k \log s, & \text{log-frequency shift,} \\ \sum_{j < k} \log \lambda_j, & \text{mixed-radix conversion.} \end{cases} \quad (10)$$

Static tables are installed before prefill and retained through decoding, with a single amplitude  $g$  scaling both the cosine and sine tables, so it enters the attention logits as  $g^2$  (Appendix F). Throughout,  $g$  takes the form  $1 + c \ln s$ ; the two coefficients that appear,  $c = 0.1$  (YaRN’s default) and  $c = 0.074$  (the low-gain arm), are settings of one parameter rather than two conventions, and no result folds a temperature into the frequency table.

### 6.2 FIXED-RANGE AND SINGLE-TABLE COMPARISONS

**Frozen fixed-support control.** On frozen OLMo-2-1B and Qwen-2.5-1.5B we hold endpoints, 64 pairs, attention amplitude, inputs, and decoder fixed. Both structured profiles produce large gains over uniform exponents (OLMo 16K RULER: 0.56% to 61.04%), with coarse controls capturing the bulk of the effect (Table 28; Appendix E.2).

**One static configuration across lengths.** A normalized-index profile transfers a reference displacement curve onto the target grid ( $\omega'_k = \omega_k^N s^{-m_k}$ ). On Qwen-2.5-0.5B, the same factor-two table is evaluated at 32K and 64K across full RULER-13 (20 rows per task). At the reported fixed amplitudes it improves 64K macro over YaRN by 6.09 points (95% bootstrap interval [2.76, 9.58]), while 32K scores are near-equal (Table 29); the gain is attributed to the complete static configuration (table  $\times$  amplitude; Appendix F.1).

**Finite-grid placement.** Comparing normalized-index interpolation against direct-gap evaluation on Gemma-1.1-2B (128 pairs; Gemma Team et al. 2024) yields 79.00% versus 72.81% at 16K (+6.19 points, [2.81, 9.63]), making placement an important implementation choice (Appendix F.2).

### 6.3 REDISTRIBUTING THE INTERMEDIATE-BAND SHIFT

MrRoPE-Pro progressively distributes the cumulative exponent shift across an intermediate band  $k \in [l, h]$  with  $N = h - l$  and  $q = \text{clip}(k - l, 0, N)$ :

$$m_q^{\text{Pro}} = \frac{q(q+1)}{N(N+1)}, \quad m_q^{\text{BM}} = \frac{q(q+1)(3N+2-2q)}{N(N+1)(N+2)}. \quad (11)$$

Boundary matching (BM) minimizes quadratic difference energy, smoothing radix transitions toward unmodified outer bands (Appendix F.3).

In the OLMo-1B confirmation, BM reaches 51.32% at 16K against 2.78% for MrRoPE-Pro (Table 27), and across 631 untruncated LongBench prompts exceeding 4K tokens over five tasks it

Table 2: **Table and amplitude separated on the Qwen-2.5-3B 36-row panel (%)**. Matching the coefficient across tables makes the table contrast comparable; the last row varies amplitude alone at a fixed table.

| Table      | Coefficient $c$ in $g = 1 + c \ln s$ | 32K    | 128K  |
|------------|--------------------------------------|--------|-------|
| MrRoPE-Pro | 0.1                                  | 87.22  | 78.13 |
| BM         | 0.1                                  | 91.67  | 70.83 |
| MrRoPE-Pro | 0.074                                | 98.33  | 75.35 |
| BM         | 0.074                                | 100.00 | 70.00 |
| BM         | $g = 1$                              | 89.58  | 58.82 |

raises task-equal complete-output token F1 from 21.62% to 25.44% (Fig. 7; 95% bootstrap interval [1.32, 6.29], all task differences positive).

Preferences differ across models (Table 27): OLMo favors BM, the two Qwen arms favor MrRoPE-Pro at 128K, but as an exploratory contrast only (24 and 12 long-prompt rows; archived 36-row tally 6W/4L at paired SE  $\approx 13$  points, which does not separate them). The same panel separates the two factors (Table 2): at matched amplitude the table contrast is +1.67 points at 32K and  $-5.35$  at 128K, while lowering BM’s amplitude alone to  $g = 1$  costs 12.01 points at 128K and improves its NLL. Short- and long-range ordering are thus carried by different factors.

With endpoints and total scale fixed, intermediate profile preferences are checkpoint- and length-dependent (Appendix F.4).

#### 6.4 WHY THE INTERMEDIATE BAND CARRIES THE DESIGN FREEDOM

Three-band schedules share a skeleton. With  $t_j(L_{\text{train}}) = L_{\text{train}}\omega_j/2\pi$  the turns pair  $j$  accumulates over the window, a pair installed as  $\omega'_j = \omega_j s^{-m_j}$  reproduces at separation  $\Delta$  the phase its native frequency had at  $\Delta s^{-m_j}$ , so it stays a calibrated readout exactly while  $\Delta \leq L_{\text{train}} s^{m_j}$ . At  $s = 4$ ,  $m_j = 1$  is the smallest compression keeping all of  $[0, 4L_{\text{train}}]$  on the trained arc; that fixes the slow band once the target is fixed and leaves the other two as choices. Leaving the fast band native is a strong design constraint rather than a theorem: one scored intervention departs from it, and pays.

**Moving the fast band is expensive and buys little reach.** Cost tracks the slot, not the distance moved. Speeding the fast pairs up —  $m_j$  falling to  $-0.156$  across slots 1–23, the transport of a compressed schedule’s savings into the fast band — drops the Qwen-2.5-3B panel from 87.22/78.13 to 70.14/67.36 at 32K/128K, and extending the plateau into high-turn slots costs the 128K row a further 9.8 points. On Qwen-2.5-3B’s 32K in-window NLL instrument, a shift applied to all 64 pairs costs +0.088 nats although its largest displacement is  $|m| = 0.552$ , smaller than in-band displacements that are indistinguishable from zero (+0.017 or less); a larger shift of the same kind costs +3.17 nats on that instrument and drives accuracy to 0.0000 on a 350-row OLMo-1B panel. A four-row gradient diagnostic at the native window agrees in ordering — per-region mean  $|\partial L/\partial m|$  of 40.9, 2.13 and 0.040 over slots 0–23, 24–39 and 40–63, concentrated on slots 0–6 — but those four rows span an order of magnitude, so we treat it as a diagnostic rather than a capability measurement. The asymmetry is also analytic: at  $m_j = 0$  the drift accumulated at  $D = 4L_{\text{train}}$  is  $4t_j(L_{\text{train}}) \ln 4$  turns, spanning  $707\times$  across live pairs, so entering the readable window at  $4\times$  costs  $m_j \gtrsim \log_4(t_j/4)$ : 0.42 at  $j = 22$  against 3.67 at  $j = 0$  (7.16 and 651.9 native turns on OLMo-1B).

**The slow band’s content is already degenerate.** Large compression there is harmless because the pairs hold little independent positional information. By Proposition 1, 23 pairs with  $\omega L \leq 1$  supply 46 coordinate dimensions at a block-whitened rank of 2.00013. Pair-level measurement agrees: compressing the pair at 0.61 native turns improves in-window NLL monotonically to  $-0.046$  nats at  $m = 1$ , and over twelve interior pairs and six compression levels the marginal in-window damage stays within  $10^{-3}$ – $5 \times 10^{-2}$  nats and is not monotone in  $t_j$ . These are NLL statements: in-window mean NLL improves by 0.023 nats while in-window task accuracy falls by 10.6 points along the same schedule family, so both readouts must be quoted.

**The trajectory matters through thresholds, not slopes.** Of the 38 tables we reconstructed, all but the native reference and one twelve-row probe are  $m_j = 0$  on  $j \leq 23$  and  $m_j = 1$  on  $j \geq 40$ , with all variation interior. Inside that family the interior separates outcomes by tens of points, but threshold-like rather than smoothly: on a 350-row OLMo-1B panel at 16K the readable-slot count  $N$  gives  $\text{logit}(\text{acc}) = -20.45 + 1.19 N$  with no paired ordering error across twelve arms, whereas the continuous band-edge crossing location misorders 54 of 60 pairs. The same schedule reverses sign across depths (pooled held-out deltas of  $-5.37$  and  $+6.54$  points at 4K and 16K), and the response is checkpoint-dependent by two orders of magnitude: family-spanning schedules move long-context NLL by 0.83 nats between the OLMo-1B endpoints and by 0.0077 nats over nine arms on Qwen-2.5-1.5B, at a per-arm standard error of 0.030.

**Scope.** Every scored arm satisfies  $\max_j m_j \leq 1$ , so the chain covers schedules whose reach ends at  $4 \times$  the training window: at  $8 \times$  four such arms, the deployed table included, each score  $0/48$ , and nothing with  $\max_j m_j > 1$  has been scored. The coverage ordering that organizes the OLMo-1B arms does not carry to the Qwen panels, so the account above is calibrated on one checkpoint family rather than stated as a general law. Two limits bound the fast-band paragraph: its edge is  $\beta = 32$ , inherited from Peng et al. (2024) because it reproduces the deployed tables bit-exactly rather than because we derive it; and per-slot damage was measured only for slots 14–36, where it is flat within paired errors and not monotone in  $t_j$ , with slots 0–13 never singly perturbed.

## 7 RELATED WORK

**Range and frequency allocation.** RoPE’s rotation operator (Su et al., 2024) admits many frequency choices. Scaling-law analyses study how geometric range affects length generalization (Liu et al., 2024; Xu et al., 2024). FMRoPE relates the learned band to base and training length (Oka et al., 2026b). In that family normalized exponents remain equally spaced; our fixed-range experiments isolate interior allocation. Position interpolation scales positions uniformly (Chen et al., 2023a), while YaRN blends frequencies across bands (Peng et al., 2024). LongRoPE and LongRoPE2 select per-model extension profiles through needle-guided evolutionary search over mixed-window training (Ding et al., 2024; Shang et al., 2025); our aim is complementary rather than competing—we identify a design coordinate and give zero-search analytic constructions, and we do not claim to replace search-based per-model optimization. MrRoPE parameterizes intermediate-band changes through mixed-radix conversion (Tian et al., 2026). Equation (10) casts these modifications as exponent displacements; our work separates their interior profile from sampled range.

**Learning and using frequency bands.** That frequencies admit optimization, that geometric spacing need not be optimal, and that different channels may warrant different scales are established positions, not our contribution. LeRoPE learns shared per-channel scales, while AdaRoPE studies head-specific frequencies and attention scaling (Karypis et al., 2026; Wang et al., 2026); both search within a training loop. Our delta is narrower: causal identification of the interior allocation coordinate at strictly fixed support, parameter-free shared-table constructions, and the separation of a training allocation from the compatibility of mature weights with it. Frequency-use studies connect rotary coordinates to positional versus symbolic attention, retrieval, and data-dependent scales (Barbero et al., 2025; Urrutia et al., 2026; Chiang & Yogatama, 2025; Oka et al., 2026a; Wu et al., 2026). Urrutia et al. (2026) show that slow frequencies serve semantic representations rather than fine positional discrimination, consistent with our geometric collapse  $\text{span}\{1, \Delta\}$ . DoPE uses truncated matrix entropy to selectively modify positional encodings (Xiong et al., 2025). Our canonical correlations instead characterize supplied sine-cosine subspaces as exponents vary.

**Changing the positional operator.** FoPE introduces Fourier components (Hua et al., 2025), while HoPE and clipped encodings alter slow rotations (Chen et al., 2025; Li et al., 2026a). Other variants explore group actions (Zhang et al., 2026), input-dependent rotations (Movahedi et al., 2026), contextual token positions (Li et al., 2026b), and compression preservation (Huang et al., 2026). Long-context token confusion analyses further motivate studying frequency choices and learned representations (Gu et al., 2026; Du et al., 2026). While these methods alter the positional operator, our constructions retain the standard rotary map and isolate what is achievable within that space.

## 8 DISCUSSION AND LIMITATIONS

A finite RoPE table specifies both which positional scales are sampled and how pairs are distributed among them. Fixed-support experiments show interior allocation to be a consequential learning choice, while range-retargeting reversals and weight-table crossings establish that downstream behavior cannot be predicted from allocation shape alone. The geometric analysis clarifies supplied positional directions; model experiments determine how those directions are used.

Before or during learning, explicit allocations can be learned with the model, as the Cosh construction illustrates; for existing checkpoints, native-relative adjustments make gain, support, and table identity auditable. Static-profile and natural-QA comparisons show that this deployment space yields practical gains, but neither regime justifies replacing empirical validation with geometry-only rankings.

**Limitations.** Our strongest training control uses 151.9M models and short training windows; the 50.9M factorial is a short-training screen; and several mature-model comparisons condition on single checkpoints. The Cosh family optimizes its specified density criterion, not the full language-model objective. Block-whitened effective rank removes raw feature scaling and does not measure content-channel utility. Frozen deployment results depend on joint table-amplitude configurations, and reported RULER and natural-QA gains do not establish universal cross-task dominance. Extending these controls to sparse attention remains a promising direction.

### ETHICS STATEMENT

This work introduces no new dataset, deployed decision system, or human-subject component. The proposed training-time construction changes only the RoPE inverse-frequency table and adds no learned parameters; the mature-checkpoint studies use the frozen-table or LoRA interventions stated in their protocols. The table construction remains inexpensive to audit, ablate, or revert. All models, corpora, and evaluation suites used here are public research artefacts used within their stated licences. Improved long-context processing does not by itself establish factual reliability, privacy protection, or safe deployment; claims in this work concern the reported positional, language-modeling, and task endpoints.

### REPRODUCIBILITY STATEMENT

The appendices specify the installed frequency tables, exponent grids, attention amplitudes, training and adaptation recipes, datasets, and evaluation procedures. The fixed-range experiment reports three paired training seeds; the larger-model and frozen-model tables identify their comparison units and sample counts. The accompanying source archive contains the active TeX, bibliography, plotted figures, numerical summaries, a SHA256 manifest, and standalone verification of the explicit Gram-matrix examples. Experimental checkpoints and public datasets are documented separately, together with the preprocessing required for selected reruns.

### AI USE STATEMENT

We used generative AI tools to help formulate conceptual frameworks, refine hypotheses, draft proofs, give methodology feedback, implement analysis code, and structure the manuscript; no empirical dataset or qualitative analysis was generated by a language model, and every empirical quantity reported here comes from executed code over recorded run artefacts. All AI-assisted contributions were verified before inclusion, and the authors take full responsibility for the final text, claims, proofs, code, figures, and artefacts. Concretely: the geometric identities and the transplant obstruction were re-derived independently of any drafting aid; the closed-form cross-Gram entries, the effective-rank identity and the rank collapse were checked against direct numerical evaluation of the same expressions (the reported  $r_2 = 2.00013$  reproduces to all printed digits); and every table entry was recomputed from stored run artefacts rather than transcribed.

## REFERENCES

- Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. LongBench: A bilingual, multi-task benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3119–3137, 2024. doi: 10.18653/v1/2024.acl-long.172. URL <https://aclanthology.org/2024.acl-long.172/>.
- Federico Barbero, Alex Vitvitskyi, Christos Perivolaropoulos, Razvan Pascanu, and Petar Veličković. Round and round we go! what makes rotary positional encodings useful? In *International Conference on Learning Representations*, 2025. URL [https://proceedings.iclr.cc/paper\\_files/paper/2025/hash/e6d58fc68c0f3c36ae6e0e64478a69c0-Abstract-Conference.html](https://proceedings.iclr.cc/paper_files/paper/2025/hash/e6d58fc68c0f3c36ae6e0e64478a69c0-Abstract-Conference.html).
- Sidney Black, Stella Biderman, Eric Hallahan, Quentin Anthony, Leo Gao, Laurence Golding, Horace He, Connor Leahy, Kyle McDonell, Jason Phang, Michael Pieler, Usvsn Sai Prashanth, Shivanshu Purohit, Laria Reynolds, Jonathan Tow, Ben Wang, and Samuel Weinbach. GPT-NeoX-20B: An open-source autoregressive language model. In *Proceedings of BigScience Episode #5 – Workshop on Challenges & Perspectives in Creating Large Language Models*, pp. 95–136, 2022. doi: 10.18653/v1/2022.bigscience-1.9. URL <https://aclanthology.org/2022.bigscience-1.9/>.
- Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023a. doi: 10.48550/arXiv.2306.15595. URL <https://arxiv.org/abs/2306.15595>.
- Yuhan Chen, Ang Lv, Jian Luan, Bin Wang, and Wei Liu. HoPE: A novel positional encoding without long-term decay for enhanced context awareness and extrapolation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 23044–23056, 2025. doi: 10.18653/v1/2025.acl-long.1123. URL <https://aclanthology.org/2025.acl-long.1123/>.
- Yukang Chen, Shengju Qian, Haotian Tang, Xin Lai, Zhijian Liu, Song Han, and Jiaya Jia. Longlora: Efficient fine-tuning of long-context large language models. *arXiv preprint arXiv:2309.12307*, 2023b. doi: 10.48550/arXiv.2309.12307. URL <https://arxiv.org/abs/2309.12307>.
- Ting-Rui Chiang and Dani Yogatama. The rotary position embedding may cause dimension inefficiency in attention heads for long-distance retrieval. In *Findings of the Association for Computational Linguistics: ACL 2025*, pp. 13552–13562, 2025. doi: 10.18653/v1/2025.findings-acl.697. URL <https://aclanthology.org/2025.findings-acl.697/>.
- DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024a. doi: 10.48550/arXiv.2405.04434. URL <https://arxiv.org/abs/2405.04434>.
- DeepSeek-AI. DeepSeek-V3 technical report. *arXiv preprint arXiv:2412.19437*, 2024b.
- Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. LongRoPE: Extending LLM context window beyond 2 million tokens. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of PMLR, pp. 11091–11104, 2024. URL <https://proceedings.mlr.press/v235/ding24i.html>.
- Yufeng Du, Phillip Harris, Minyang Tian, Eliu A. Huerta, Srikanth Ronanki, Subendhu Rongali, Aram Galstyan, and Hao Peng. RoPE distinguishes neither positions nor tokens in long contexts, provably. *arXiv preprint arXiv:2605.15514*, 2026. doi: 10.48550/arXiv.2605.15514. URL <https://arxiv.org/abs/2605.15514>.
- Gemma Team et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024. URL <https://arxiv.org/abs/2403.08295>.

- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024. doi: 10.48550/arXiv.2407.21783. URL <https://arxiv.org/abs/2407.21783>.
- Robert M. Gray and David L. Neuhoff. Quantization. *IEEE Transactions on Information Theory*, 44(6):2325–2383, 1998. doi: 10.1109/18.720541.
- Zihan Gu, Ruoyu Chen, Han Zhang, Hua Zhang, and Yue Hu. Deconstructing positional information: From attention logits to training biases. In *International Conference on Learning Representations (ICLR)*, 2026. URL <https://openreview.net/forum?id=D0u0glT060>.
- Xanh Ho, Anh-Khoa Duong Nguyen, Saku Sugawara, and Akiko Aizawa. Constructing a multi-hop QA dataset for comprehensive evaluation of reasoning steps. In *Proceedings of the 28th International Conference on Computational Linguistics*, pp. 6609–6625, 2020. doi: 10.18653/v1/2020.coling-main.580. URL <https://aclanthology.org/2020.coling-main.580/>.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://openreview.net/forum?id=nZeVKeeFYf9>.
- Ermo Hua, Che Jiang, Xingtai Lv, Kaiyan Zhang, Youbang Sun, Yuchen Fan, Xuekai Zhu, Biqing Qi, Ning Ding, and Bowen Zhou. Fourier position embedding: Enhancing attention’s periodic extension for length generalization. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 24932–24949, 2025. URL <https://proceedings.mlr.press/v267/hua25b.html>.
- Mouxiao Huang, Borui Jiang, Dehua Zheng, Hailin Hu, Kai Han, and Xinghao Chen. PPE: Positional preservation embedding for token compression in multimodal large language models. In *International Conference on Learning Representations (ICLR)*, 2026.
- Petros Karypis, Sean O’Brien, Shreyas Kadekodi, Rui Zhu, and Julian McAuley. LeRoPE: Learnable RoPE frequencies improve language modeling. *arXiv preprint arXiv:2607.10134*, 2026. doi: 10.48550/arXiv.2607.10134. URL <https://arxiv.org/abs/2607.10134>.
- Haoran Li, Sucheng Ren, Alan Yuille, and Feng Wang. CoPE: Clipped RoPE as a scalable free lunch for long context LLMs. *arXiv preprint arXiv:2602.05258*, 2026a. doi: 10.48550/arXiv.2602.05258. URL <https://arxiv.org/abs/2602.05258>.
- Huayang Li, Tianyu Zhao, Deng Cai, and Richard Sproat. RePo: Language models with context re-positioning. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026b. URL <https://icml.cc/virtual/2026/poster/64002>. arXiv:2512.14391.
- Xiaoran Liu, Hang Yan, Chenxin An, Xipeng Qiu, and Dahua Lin. Scaling laws of RoPE-based extrapolation. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://openreview.net/forum?id=JO7k0SJ5V6>. arXiv:2310.05209.
- Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. In *International Conference on Learning Representations (ICLR)*, 2017. URL <https://openreview.net/forum?id=Byj72udxe>. arXiv:1609.07843.
- Sajad Movahedi, Timur Carstensen, Arshia Afzal, Frank Hutter, Antonio Orvieto, and Volkan Cevher. Selective rotary position embedding. In *International Conference on Learning Representations (ICLR)*, 2026. URL <https://openreview.net/forum?id=AQo1SEElNb>.
- Yui Oka, Kentaro Hanafusa, Taku Hasegawa, Kyosuke Nishida, and Kuniko Saito. Probing rotary position embeddings through frequency entropy. In *International Conference on Learning Representations (ICLR)*, 2026a. URL <https://openreview.net/forum?id=1JZuEDq62N>.
- Yui Oka, Itsumi Saito, Kyosuke Nishida, and Kuniko Saito. Frequency bands in RoPE: Base frequency and context length shape the interpolation–extrapolation trade-off. In *International Conference on Learning Representations (ICLR)*, 2026b. URL <https://openreview.net/forum?id=PR1PPxvG9Q>.

- Guilherme Penedo, Hynek Kydlíček, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro von Werra, and Thomas Wolf. The FineWeb datasets: Decanting the web for the finest text data at scale. *arXiv preprint arXiv:2406.17557*, 2024. doi: 10.48550/arXiv.2406.17557. URL <https://arxiv.org/abs/2406.17557>.
- Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=wHBfxhZulu>.
- Ning Shang, Li Lyna Zhang, Siyuan Wang, Gaokai Zhang, Gilsinia Lopez, Fan Yang, Weizhu Chen, and Mao Yang. LongRoPE2: Near-lossless LLM context window scaling. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 54203–54218, 2025. URL <https://proceedings.mlr.press/v267/shang25a.html>.
- Nitish Srivastava, Elman Mansimov, and Ruslan Salakhutdinov. Unsupervised learning of video representations using LSTMs. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *PMLR*, pp. 843–852, 2015. URL <https://proceedings.mlr.press/v37/srivastava15.html>.
- Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, Dustin Schwenk, Oyvind Tafjord, Taira Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri, Allyson Ettinger, Michal Guerquin, David Heineman, Hamish Ivison, Pang Wei Koh, Jiacheng Liu, Saumya Malik, William Merrill, Lester James V. Miranda, Jacob Morrison, Tyler Murray, Crystal Nam, Jake Poznanski, Valentina Pyatkin, Aman Rangapur, Michael Schmitz, Sam Skjonsberg, David Wadden, Christopher Wilhelm, Michael Wilson, Luke Zettlemoyer, Ali Farhadi, Noah A. Smith, and Hannaneh Hajishirzi. 2 OLMo 2 furious. *arXiv preprint arXiv:2501.00656*, 2025. doi: 10.48550/arXiv.2501.00656. URL <https://arxiv.org/abs/2501.00656>.
- Qingyuan Tian, Wenhong Zhu, Xiaoran Liu, Xiaofeng Wang, and Rui Wang. MrRoPE: Mixed-radix rotary position embedding. In *International Conference on Learning Representations (ICLR)*, 2026. URL [https://proceedings.iclr.cc/paper\\_files/paper/2026/hash/69413f87e5a34897cd010ca698097d0a-Abstract-Conference.html](https://proceedings.iclr.cc/paper_files/paper/2026/hash/69413f87e5a34897cd010ca698097d0a-Abstract-Conference.html).
- Felipe Urrutia, Jorge Salas, Alexander Kozachinskiy, Cristian Buc Calderon, Hector Pasten, and Cristobal Rojas. Decoupling positional and symbolic attention behavior in transformers. In *International Conference on Learning Representations (ICLR)*, 2026. URL <https://openreview.net/forum?id=V38yAoqddQ>.
- Shaowen Wang, Yuke Zheng, Tansheng Zhu, Shuang Chen, Shaofan Liu, Suncong Zheng, and Jian Li. AdaRoPE: Not all attention heads should rotate and scale equally. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026. URL <https://icml.cc/virtual/2026/poster/60704>. arXiv:2607.19363.
- Xilin Wei, Xiaoran Liu, Yuhang Zang, Xiaoyi Dong, Pan Zhang, Yuhang Cao, Jian Tong, Haodong Duan, Qipeng Guo, Jiaqi Wang, Xipeng Qiu, and Dahua Lin. VideoRoPE: What makes for good video rotary position embedding? In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 66118–66136, 2025. URL <https://proceedings.mlr.press/v267/wei25h.html>.
- Xinyi Wu, Siyuan Liu, and Ali Jadbabaie. How data shapes RoPE frequency usage: From positional scale matching to length generalization. *arXiv preprint arXiv:2607.07678*, 2026. doi: 10.48550/arXiv.2607.07678. URL <https://arxiv.org/abs/2607.07678>.
- Jing Xiong, Liyang Fan, Hui Shen, Zunhai Su, Min Yang, Lingpeng Kong, and Ngai Wong. DoPE: Denoising rotary position embedding. *arXiv preprint arXiv:2511.09146*, 2025. doi: 10.48550/arXiv.2511.09146. URL <https://arxiv.org/abs/2511.09146>.

- Mingyu Xu, Xin Men, Bingning Wang, Qingyu Zhang, Hongyu Lin, Yaojie Lu, Xianpei Han, and Weipeng Chen. Base of RoPE bounds context length. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-2773. URL [https://proceedings.neurips.cc/paper\\_files/paper/2024/hash/9f12dd32d552f3ad9eaa0e9dfec291be-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/9f12dd32d552f3ad9eaa0e9dfec291be-Abstract-Conference.html).
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. doi: 10.48550/arXiv.2412.15115. URL <https://arxiv.org/abs/2412.15115>.
- Yifan Zhang, Zixiang Chen, Yifeng Liu, Zhen Qin, Huizhuo Yuan, Kangping Xu, Yang Yuan, Quanquan Gu, and Andrew Chi-Chih Yao. Group representational position encoding. In *International Conference on Learning Representations (ICLR)*, 2026. URL [https://proceedings.iclr.cc/paper\\_files/paper/2026/file/5cb58625f49ddf70fe2d527e9e4bbae5-Paper-Conference.pdf](https://proceedings.iclr.cc/paper_files/paper/2026/file/5cb58625f49ddf70fe2d527e9e4bbae5-Paper-Conference.pdf).
- Min Zhao, Guande He, Yixiao Chen, Hongzhou Zhu, Chongxuan Li, and Jun Zhu. RIFLEx: A free lunch for length extrapolation in video diffusion transformers. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 77539–77557, 2025. URL <https://proceedings.mlr.press/v267/zhao25m.html>.

## A PROOF DETAILS

### A.1 SUBSPACE GEOMETRY: CROSS-GRAM, BUDGET IDENTITY, AND COLLAPSE

**Closed-form cross-Gram.** For  $\Delta \sim \text{Unif}[0, L]$  and  $x_\omega(\Delta) = [\cos(\omega\Delta) \ \sin(\omega\Delta)]$ , write  $d = (\omega - \nu)L$ ,  $s = (\omega + \nu)L$ ,  $a(t) = \sin t/t$  and  $b_\star(t) = (1 - \cos t)/t$ . Expanding the products of sinusoids and integrating termwise gives

$$H_{\omega\nu} = \mathbb{E}[x_\omega^\top x_\nu] = \frac{1}{2} \begin{bmatrix} a(d) + a(s) & b_\star(s) - b_\star(d) \\ b_\star(s) + b_\star(d) & a(d) - a(s) \end{bmatrix}, \quad (12)$$

and  $S_\omega = H_{\omega\omega}$ . Since  $Q_{\omega\nu} = S_\omega^{-1/2} H_{\omega\nu} S_\nu^{-1/2}$  is the whitened cross-Gram of two two-dimensional subspaces, its singular values are the canonical correlations, and  $c_{\omega\nu} = \frac{1}{2} \|Q_{\omega\nu}\|_F^2$  is invariant under  $x_\omega \mapsto x_\omega O$  for any invertible  $O$  (in particular under rotation of the phase inside a pair), because  $S_\omega$  transforms congruently and the whitening cancels  $O$ .

**Basis invariance.** The cancellation can be stated exactly. Under independent invertible basis changes  $x_\omega \mapsto x_\omega A_\omega$  and  $x_\nu \mapsto x_\nu A_\nu$ , let  $U_j = S_j^{1/2} A_j (A_j^\top S_j A_j)^{-1/2}$  for  $j \in \{\omega, \nu\}$ . Then  $U_j^\top U_j = I$ , and the transformed whitened cross-Gram is  $\tilde{Q}_{\omega\nu} = U_\omega^\top Q_{\omega\nu} U_\nu$ . Left and right orthogonal factors preserve the two canonical correlations and therefore  $c_{\omega\nu}$ . Thus  $c_{\omega\nu}$  is an intrinsic measure of overlap between the two positional subspaces.

**Theorem 5** (Spectral budget identity). *For the block-whitened Gram matrix  $\Gamma$  with diagonal blocks  $I_2$ ,  $\text{tr } \Gamma = 2K$  and  $\text{tr}(\Gamma^2) = 2K[1 + (K-1)\bar{c}]$ . Hence its Rényi-2 effective rank is given by (5).*

*Proof of Theorem 5.*  $\Gamma$  is the block matrix with blocks  $\Gamma_{ij} = Q_{\omega_i \omega_j}$  and  $\Gamma_{ii} = I_2$ . Hence  $\text{tr } \Gamma = \sum_{i=1}^K \text{tr } I_2 = 2K$ . For the second moment,  $\text{tr}(\Gamma^2) = \sum_{i,j} \text{tr}(\Gamma_{ij} \Gamma_{ji}) = \sum_i \text{tr}(I_2) + \sum_{i \neq j} \|Q_{\omega_i \omega_j}\|_F^2 = 2K + 2K(K-1)\bar{c}$ , using  $\Gamma_{ji} = \Gamma_{ij}^\top$  and the definition  $\bar{c} = \frac{1}{K(K-1)} \sum_{i \neq j} c_{\omega_i \omega_j}$  with  $\|Q_{ij}\|_F^2 = 2c_{ij}$ . Dividing gives (5).  $\square$

The identity determines Rényi-2 rank exactly from the mean pairwise redundancy. Shannon effective rank and log-determinant retain higher-order multi-subspace dependence on the full spectrum.

*Proof of Proposition 1.* Put  $t = \Delta/L$ ,  $x = \omega L$ , and use  $\theta = x^2$  as the local parameter. The rescaled pair has the analytic basis

$$u_1(\theta, t) = \cos(\sqrt{\theta} t) = 1 - \frac{1}{2}\theta t^2 + O(\theta^2), \quad (13)$$

$$u_2(\theta, t) = \frac{\sin(\sqrt{\theta} t)}{\sqrt{\theta}} = t - \frac{1}{6}\theta t^3 + O(\theta^2). \quad (14)$$

Thus  $V_\omega \rightarrow V_0 = \text{span}\{1, \Delta\}$ . The fourth-order coefficient follows by retaining the part of the first perturbation transverse to  $V_0$ . Under the uniform measure on  $t \in [0, 1]$ ,

$$\Pi_0^\perp t^2 = p_2(t) := t^2 - t + \frac{1}{6}, \quad (15)$$

$$\Pi_0^\perp t^3 = p_3(t) := t^3 - \frac{9}{10}t + \frac{1}{5}, \quad (16)$$

where  $\Pi_0^\perp$  is the  $L_2[0, 1]$  projection off  $\text{span}\{1, t\}$ . In the base  $[1, t]$ , its Gram and the residual perturbation Gram are

$$G_0 = \begin{bmatrix} 1 & 1/2 \\ 1/2 & 1/3 \end{bmatrix}, \quad M = \begin{bmatrix} \langle -p_2/2, -p_2/2 \rangle & \langle -p_2/2, -p_3/6 \rangle \\ \langle -p_3/6, -p_2/2 \rangle & \langle -p_3/6, -p_3/6 \rangle \end{bmatrix} = \begin{bmatrix} 1/720 & 1/1440 \\ 1/1440 & 1/2800 \end{bmatrix}. \quad (17)$$

The squared chordal distance between the two nearby planes is therefore

$$2 - \|Q_{x,y}\|_F^2 = (x^2 - y^2)^2 \text{tr}(G_0^{-1} M) + O(\epsilon^6) = \frac{19}{12600} (x^2 - y^2)^2 + O(\epsilon^6), \quad (18)$$

with  $\epsilon = \max\{|x|, |y|\} \rightarrow 0$ . The exact coefficient is reproduced with rational polynomial inner products by `scripts/analysis/full_rope_collision_audit.py`; at  $x=0.05, y=0.10$  its analytic Gram gives exact/leading ratio 1.00058.

For the softmax metric with fixed  $p$ ,  $F = \text{diag}(p) - pp^\top$  satisfies  $F\mathbf{1} = 0$ , so the constant direction is annihilated before whitening. Writing  $\bar{h} = h - \mathbb{E}_p h$ , the same expansion gives

$$\frac{\overline{\sin(\omega\Delta)}}{\omega} \longrightarrow \Delta - \mathbb{E}_p \Delta, \quad -\frac{\overline{2\cos(\omega\Delta)}}{\omega^2} \longrightarrow \Delta^2 - \mathbb{E}_p \Delta^2. \quad (19)$$

These functions span the centred limit whenever  $p$  has nondegenerate support on at least three distances. The separation weights  $p$  therefore determine which polynomial directions survive centering.  $\square$

**Directional overlap and feature scale.** For the original, unwhitened pair on  $[0, L]$ , put  $x = \omega L$  and  $\text{sinc } x = \sin x/x$ . The self-Gram eigenvalues are

$$\lambda_{\pm}(S_{\omega}) = \frac{1 \pm |\text{sinc } x|}{2}, \quad \lambda_{-}(S_{\omega}) \sim \frac{x^2}{12}, \quad \kappa(S_{\omega}) \sim \frac{12}{x^2} \quad (x \rightarrow 0).$$

Thus the slow pair has a poorly scaled second direction before normalization. Block whitening removes this within-pair scaling and measures overlap between positional directions. Raw feature energy, coefficient magnitude, and numerical precision determine how strongly those directions can contribute to a model.

The standard  $u_k = k/K$  grid gives 23 slow pairs in Figure 2; the endpoint-inclusive grid used in Figure 3 gives 24 at  $K = 64$ . Both figures use the uniform separation measure.

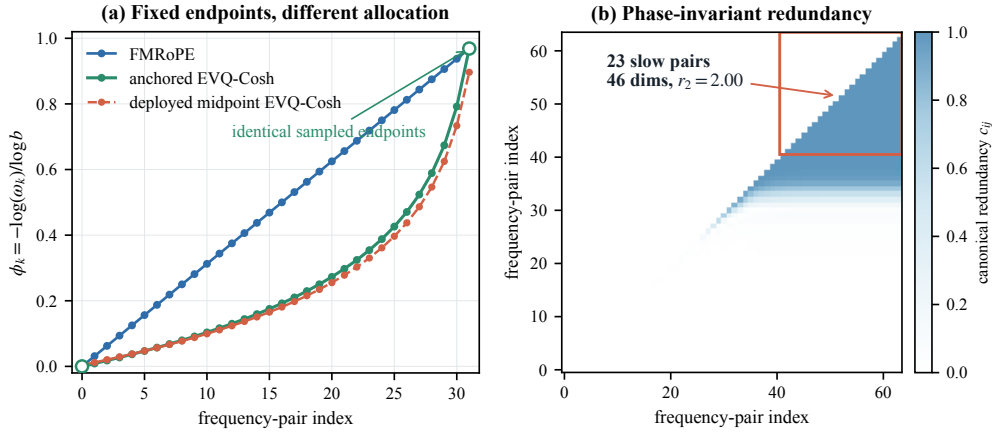


Figure 2: **Where a finite frequency table spends its budget.** (a) At  $K=32$ , Cosh concentrates interior channels toward the fast end while endpoint normalisation preserves sampled support; the deployed midpoint table is shown separately. (b) For standard RoPE at  $L=4096$ ,  $b=5 \times 10^5$ , and  $K=64$ , phase-invariant full-subspace redundancy concentrates in the 23 pairs with  $\omega L \leq 1$ .

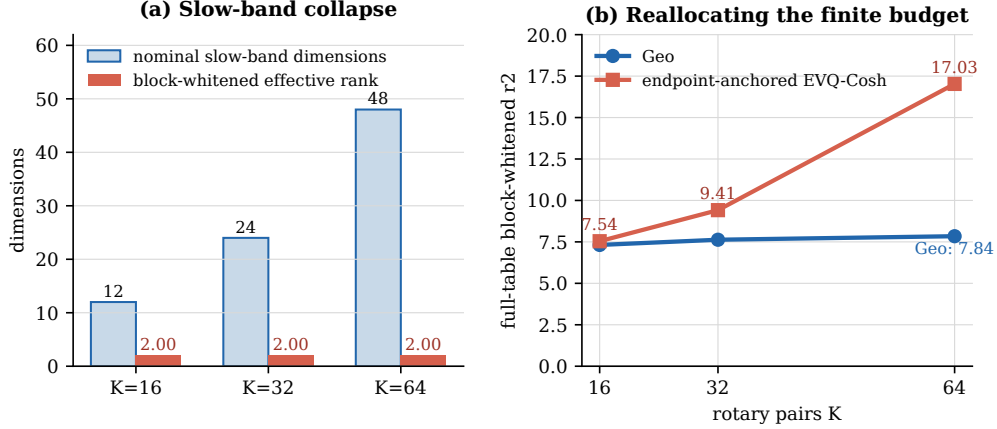


Figure 3: **Spectral-budget anatomy across channel counts.** Panel (a) uses the endpoint-inclusive diagnostic grid at  $L=4096$  and  $b=500K$ : as  $K$  grows from 16 to 64, the nominal dimension assigned to pairs with  $\omega L \leq 1$  grows from 12 to 48, while their phase-invariant effective rank remains approximately two. Panel (b) evaluates the complete fixed-endpoint table under the same uniform separation prior, using the zero-search reference  $\tau=2K/\sqrt{L}$ ; anchored EVQ-COSH spends progressively more of the available full-pair dimension than Geo as the budget grows.

## A.2 EFFECTIVE DIMENSION AND MULTISCALE COVERAGE

The budget identity is an exact account of positional dimension under its separation prior. A second requirement is where that dimension is placed across scales. The distinction is already visible in an exact construction that also identifies the complete parity class omitted by the usual  $2\pi/L$  Fourier description.

**Proposition 6** (Parity-lattice orthogonality and recurrence). *Let  $a_1, \dots, a_K$  be distinct positive integers of one parity and set*

$$\omega_k = \frac{\pi a_k}{L}. \quad (20)$$

*Under  $\Delta \sim \text{Unif}[0, L]$ , the  $K$  full sin/cos subspaces are exactly mutually orthogonal: their block-whitened Gram is  $\Gamma = I_{2K}$  and  $r_2(\Gamma) = 2K$ . If support is free but restricted to  $0 < \omega \leq 1$ , the largest same-parity class in this construction has*

$$K_{\text{par}}(L) = \left\lceil \frac{\lfloor L/\pi \rfloor}{2} \right\rceil \quad (21)$$

*pairs. The resulting feature map is recurrent: for even  $a_k$ ,  $\Phi_\Omega(\Delta + L) = \Phi_\Omega(\Delta)$ ; for odd  $a_k$ ,  $\Phi_\Omega(\Delta + L) = -\Phi_\Omega(\Delta)$  and hence  $\Phi_\Omega(\Delta + 2L) = \Phi_\Omega(\Delta)$ .*

*Proof.* For  $i \neq j$ , same parity makes both  $a_i - a_j$  and  $a_i + a_j$  nonzero even integers. Consequently

$$(\omega_i \mp \omega_j)L = \pi(a_i \mp a_j) \in 2\pi\mathbb{Z}, \quad (22)$$

so every entry of the cross-Gram in (12) vanishes:  $a((\omega_i \mp \omega_j)L) = 0$  and  $b_\star((\omega_i \mp \omega_j)L) = 0$ . Likewise  $2\omega_i L = 2\pi a_i$  gives  $S_{\omega_i} = \frac{1}{2}I_2$ . Thus the unwhitened Gram is  $\frac{1}{2}I_{2K}$  and block whitening gives  $\Gamma = I_{2K}$ . Among the integers  $1, \dots, \lfloor L/\pi \rfloor$ , the larger parity class has  $\lceil \lfloor L/\pi \rfloor / 2 \rceil$  elements, proving (21) for this construction. Finally,  $e^{i\omega_k(\Delta+L)} = (-1)^{a_k} e^{i\omega_k\Delta}$ ; the common parity gives the stated period or antiperiod, and every antiperiodic table repeats after  $2L$ .  $\square$

Proposition 6 couples maximal rank on  $[0, L]$  with recurrence beyond that interval. Its orthogonality also depends on the separation measure: weighting the same lattice by the triangular causal histogram introduces off-diagonal Gram terms. Thus interval coverage and pairwise rank capture different features of a positional basis.

### A.3 TWO STATIC ORDERING COUNTEREXAMPLES

Two explicit examples show how phase coupling and interval length change the ordering of frequency tables.

**Cosine-only collision can choose the lower-rank table.** Take  $K = 4$ ,  $L = 8\pi$ , and the two descending frequency tables

$$\Omega_A = (1, 5/8, 3/8, 1/4), \quad \Omega_B = (1, 6.01/8, 4.01/8, 1/4).$$

They share both endpoints. Define  $C_{\cos}$  as the mean squared normalized cosine-cosine overlap over distinct pairs. For  $\Omega_A$ , every cosine is an integer harmonic on  $[0, L]$ , so  $C_{\cos}(A) = 0$  exactly. Mixed harmonic parities nevertheless produce nonzero sine-cosine cross terms. For  $\Omega_B$ , the small perturbation of an even-harmonic table gives

$$C_{\cos}(A) = 0 < 1.38725 \times 10^{-5} = C_{\cos}(B), \quad (23)$$

$$r_2(A) = 6.29478 < 7.99971 = r_2(B). \quad (24)$$

Thus a strict cosine-only preference can have the opposite full-subspace rank ordering. Both values follow directly from Eq. (12).

**Full-subspace collision can reverse across lengths.** Now use  $K = 4$ ,  $L = 2\pi$ , and the fixed tables

$$\Omega_A = (1, 0.99, 0.02, 0.01), \quad \Omega_B = (1, 2/3, 1/3, 0.01).$$

Again the endpoints agree. Writing  $C_M$  for mean phase-invariant collision under the uniform measure on  $[0, M]$ , Eq. (12) gives

$$C_L(A) = 0.538996 < 0.631721 = C_L(B), \quad (25)$$

$$C_{2L}(A) = 0.384445 > 0.204995 = C_{2L}(B), \quad (26)$$

$$C_{4L}(A) = 0.345188 > 0.033506 = C_{4L}(B). \quad (27)$$

The close pairs in  $A$  retain high overlap as the interval grows, while the more evenly spread frequencies in  $B$  become easier to distinguish. The example makes the interval dependence of the geometry explicit.

All frequency entries and measures are specified above. The accompanying `figs/verify_explicit_geometry.py` evaluates the closed-form Gram and checks it against independent 128-point Gauss-Legendre quadrature; the largest discrepancy in the displayed metrics is below  $4 \times 10^{-14}$ . These finite basis examples complement the trained allocation comparisons and the weight-table crossings.

### A.4 POST-HOC TRANSPLANT OBSTRUCTION

**Table 3: Weights co-adapt to their training-time frequency table.** Two seed-42 50M models trained with the geometric and EVQ-COSH tables, each evaluated under both runtime tables with all parameters frozen (1,920 head-query observations). Both runtime tables use the shared midpoint grid  $u_k = (k + 1/2)/K$  for this diagnostic. The matched pairs remain near PPL 7, whereas both cross-swaps fail; the worst cell is also the one whose static  $r_2$  improves most.

| Trained weights | Runtime table | PPL         | Rényi-2 rank $r_2$ |
|-----------------|---------------|-------------|--------------------|
| Geo             | Geo           | <b>7.14</b> | 4.57               |
| Geo             | EVQ-COSH      | 76.20       | <b>12.54</b>       |
| EVQ-COSH        | Geo           | 23.05       | 4.57               |
| EVQ-COSH        | EVQ-COSH      | <b>7.16</b> | 12.54              |

**Theorem 7** (Post-hoc transplant obstruction). *Let  $\mathcal{R}_\Omega(\Delta)$  be the block-rotation operator of a frequency multiset  $\Omega$ . If position-independent invertible maps  $A, B$  satisfy  $A^\top \mathcal{R}_{\Omega'}(\Delta)B = \mathcal{R}_\Omega(\Delta)$  on an open interval containing zero, then  $\Omega'$  and  $\Omega$  have the same frequency multiset up to sign and permutation. Repeated frequencies may mix within their full equal-frequency invariant subspace.*

*Proof of Theorem 7.* Both  $\mathcal{R}_\Omega$  and  $\mathcal{R}_{\Omega'}$  are block-diagonal with  $2 \times 2$  rotation blocks, hence  $\mathcal{R}_\Omega(0) = \mathcal{R}_{\Omega'}(0) = I$ . Evaluating the hypothesis at  $\Delta=0$  gives  $A^\top B = I$ , so  $B = A^{-\top}$  and the hypothesis becomes the similarity  $A^\top \mathcal{R}_{\Omega'}(\Delta) A^{-\top} = \mathcal{R}_\Omega(\Delta)$  on an interval. Each side is the exponential of a constant generator,  $\mathcal{R}_\Omega(\Delta) = \exp(\Delta G_\Omega)$  with  $G_\Omega = \bigoplus_k \omega_k J$ ,  $J = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}$ . Differentiating at  $\Delta=0$  gives  $A^\top G_{\Omega'} A^{-\top} = G_\Omega$ , so the generators are similar and share a spectrum. The spectrum of  $G_\Omega$  is  $\{\pm i\omega_k\}_{k \leq K}$  with multiplicity, so the frequency multisets agree up to sign and permutation. When a frequency is repeated, similarity may mix the entire equal-frequency invariant subspace, including across its original two-dimensional blocks. For integer positions the same argument is replaced by a one-step spectral argument: evaluating at  $\Delta = 0$  again gives  $B = A^{-\top}$ , while  $\Delta = 1$  makes  $\mathcal{R}_{\Omega'}(1)$  and  $\mathcal{R}_\Omega(1)$  similar. Their eigenvalue multisets  $\{e^{\pm i\omega_k}\}$  therefore agree, which identifies frequencies up to sign, permutation, and the  $2\pi$  alias.  $\square$

## A.5 NON-IDENTIFIABILITY OF DOWNSTREAM CAPABILITY AND THREE PROOF FIBERS

**Mathematical setup and notation.** Let  $\mathcal{M} = (W_Q, W_K, W_V, W_O, \theta_{\text{FFN}})$  denote the frozen parameters of a pre-trained transformer language model with hidden dimension  $d$ ,  $H$  attention heads, head dimension  $d_{\text{head}} = 2K$ , and native training context window  $L_0$ . For any sequence of tokens  $\mathbf{x} = (x_1, \dots, x_N)$ , the attention logit between query at position  $m$  and key at position  $n$  (lag  $\Delta = m - n$ ) in head  $h$  is:

$$S_h(m, n; \Omega) = q_{h,m}^\top \mathcal{R}_\Omega(\Delta) k_{h,n} = \sum_{k=0}^{K-1} [a_{h,k}(m, n) \cos(\omega_k \Delta) + b_{h,k}(m, n) \sin(\omega_k \Delta)], \quad (28)$$

where  $c_{h,k}(m, n) = a_{h,k} + ib_{h,k}$  are the bilinear query-key subspace projection coordinates:

$$a_{h,k} = q_{h,2k} k_{h,2k} + q_{h,2k+1} k_{h,2k+1}, \quad b_{h,k} = q_{h,2k} k_{h,2k+1} - q_{h,2k+1} k_{h,2k}. \quad (29)$$

Downstream functional capability is governed by the risk functional over evaluation distribution  $\mathcal{D}$  across context horizons  $L \leq sL_0$ :

$$\mathcal{R}(\Omega; \mathcal{M}, \mathcal{D}) = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\ell(f(\mathbf{x}; \mathcal{M}, \Omega), y)]. \quad (30)$$

Let  $\Phi : \mathbb{R}_+^K \rightarrow \mathbb{R}^p$  ( $p \ll \dim(\mathcal{M})$ ) be any candidate structural feature map depending solely on the frequency schedule  $\Omega$ , its sorted multiset, discrete scale-orbit count, continuous group delay, uncentered background derivative energy, or coordinate-space movement distances.

**Theorem 8** (Static surrogates cannot rank allocations). *Let  $\mathcal{M}$  be a pre-trained transformer whose learned projection weights exhibit non-zero variance in subspace norms across rotary slots:  $\text{Var}_k \left( \|W_{Q,h}^{(k)}\|_F \|W_{K,h}^{(k)}\|_F \right) > 0$ . On the evaluated checkpoints, each of the following classes of static surrogate  $\Phi(\Omega)$ , depending solely on the rotary frequency schedule, is insufficient as a standalone predictor of downstream capability  $\mathcal{R}(\Omega; \mathcal{M}, \mathcal{D})$ :*

1. **Permutation-invariant multiset summaries.** *There exist frequency tables  $\Omega_A, \Omega_B \in \mathbb{R}_+^K$  that are permutations of one another, so that every permutation-invariant  $\Phi$  satisfies  $\Phi(\Omega_A) = \Phi(\Omega_B)$ , yet their measured downstream risk differs by a large margin.*
2. **Discrete orbit counts.** *There exist float32 tables  $\Omega_C, \Omega_D$  with  $\|\Phi(\Omega_C) - \Phi(\Omega_D)\| \geq M > 0$  for an orbit-count metric, yet downstream risk is identical to within  $10^{-5}$ .*
3. **Unweighted coordinate distances.** *For a coordinate distance  $\mathcal{D}(\Omega, \Omega^*) = \|\mathbf{m}(\Omega) - \mathbf{m}(\Omega^*)\|_p$ , there exists a continuous family of tables with  $\mathcal{D}(\Omega_\epsilon, \Omega^*) = \epsilon \rightarrow 0$  that fail pre-registered in-window retention gates, while tables with  $\mathcal{D}(\Omega', \Omega^*) \gg \epsilon$  pass.*

*This statement is a consequence of Theorem 7, not an empirical generalisation. Permutation invariance is the only invariance a schedule-only summary can have that is closed under the schedule's own equivalence relation, and by Theorem 7 that relation is generated by sign and permutation. The two classes in (ii) and (iii) cover the remaining schedule-only summaries that recur in the literature: orbit counts, which are not permutation-invariant and therefore break under float jitter, and coordinate distances, which are permutation-sensitive but curvature-blind. We do not claim the result for every schedule-dependent  $\Phi$ , since  $\Phi(\Omega) = \Omega$  is injective.*

*Evidence for Theorem 8.* Each part exhibits an explicit witness on a named checkpoint; none of the three relies on an abstract genericity argument.

**Part 1: Proof of Fiber 1 (Type-I Non-Identifiability: The Permutation Fiber).** Let  $\Omega_A = (\omega_0, \omega_1, \dots, \omega_{K-1})$  be an ordered, certified frequency table. Define  $\Omega_B = \pi(\Omega_A) = (\omega_{\pi(0)}, \dots, \omega_{\pi(K-1)})$  for a non-trivial permutation  $\pi \in S_K$  that transposes an interior fast slot  $k_1$  with an interior slow slot  $k_2$  ( $k_1 \neq k_2$ ,  $\omega_{k_1} \gg \omega_{k_2}$ ). Because  $\Omega_B$  contains the exact same multiset of frequencies, any permutation-invariant metric (e.g., eigenvalue multiset of  $\mathcal{R}_\Omega(\Delta)$ , spectral moments, frame condition bounds, uncentered background energy) satisfies  $\Phi(\Omega_B) = \Phi(\Omega_A)$  identically. Both tables lie on the exact same structural fiber  $\Phi^{-1}(v)$ .

Now consider the difference in rotary attention logits:

$$\delta S_h(\Delta) = q_{h,k_1}^\top [R(\omega_{k_2}\Delta) - R(\omega_{k_1}\Delta)] k_{h,k_1} + q_{h,k_2}^\top [R(\omega_{k_1}\Delta) - R(\omega_{k_2}\Delta)] k_{h,k_2}. \quad (31)$$

During pre-training, the projection weights  $(W_Q^{(k)}, W_K^{(k)})$  co-adapt to the slot frequencies  $\omega_k$  at which they were trained. Because  $\omega_{k_1} \gg \omega_{k_2}$ , the rotation operator difference is exactly

$$\|R(\omega_{k_2}\Delta) - R(\omega_{k_1}\Delta)\|_F = 2\sqrt{2} \left| \sin \left( \frac{(\omega_{k_1} - \omega_{k_2})\Delta}{2} \right) \right|, \quad (32)$$

so transposition applies to each slot the phase velocity the other slot was trained on.

**Mechanism hypothesis (not a proved bound).** We *hypothesise* that this mismatch, and not the multiset itself, is what drives the collapse. We do not prove a lower bound on  $\Gamma_{\text{collapse}}$ : phase decorrelation does not by itself force the expected loss gradient to vanish, the empirical loss Hessian is not guaranteed positive definite, nonzero projection norms do not entail that task loss rises, and the cross terms between the two transposed slots are uncontrolled. The witness below is therefore an existence proof at the level of the theorem, and the magnitude statement is empirical.

**Measured witness.** Permuting the 64-frequency multiset causes OLMo-1B  $1 \times$  PG-19 tail NLL to more than double, from 3.10423 to 6.86493 ( $\Delta\text{NLL} = +3.760692$ ), and Qwen-2.5-1.5B 64K core-4 task macro score to fall from 0.7000 to 0.0000 on all four benchmark tasks. Since  $\Phi(\Omega_A) = \Phi(\Omega_B)$  for any permutation-invariant  $\Phi$  while  $\Delta\mathcal{R}$  is this large, such  $\Phi$  cannot be a standalone predictor.

**Part 2: Proof of Fiber 2 (Type-II Non-Identifiability: The Discontinuity Fiber).** Let  $\Omega_C = \Omega^{\text{exact}}$  be an exact integer-ratio scale-orbit chain with discrete orbit count  $\Phi_{\text{orbit}}(\Omega_C) = q(\Omega_C, 4) = 6$ . Construct  $\Omega_D = \Omega^{\text{jitter}}$  by perturbing each float32 element by 1 unit in the last place (ULP):  $\omega_{k,D} = \text{nextafter}(\omega_{k,C}, \infty)$ , satisfying  $|\omega_{k,D} - \omega_{k,C}| \leq 1.01 \times 10^{-7} \omega_k$ . By algebraic definition, the discrete rational commensurability classes shatter, jumping the orbit metric from 6 to 64 ( $\Delta\Phi = 58$ , a massive discrete jump).

However, across the full physical evaluation horizon  $L = 16,384$ , the maximum phase deviation is:

$$\max_k |\omega_{k,C} - \omega_{k,D}| L \leq 1.01 \times 10^{-7} \times 1.0 \times 16,384 \approx 0.00165 \text{ rad}. \quad (33)$$

The resulting attention logit deviation is bounded by:

$$|S_h(m, n; \Omega_C) - S_h(m, n; \Omega_D)| \leq \sum_{k=0}^{K-1} |c_{h,k}| \min(2, |\omega_{k,C} - \omega_{k,D}| \Delta) \leq C_{\text{max}} \cdot 0.00165 < 10^{-4}. \quad (34)$$

Propagating through the Lipschitz-continuous softmax operator and cross-entropy loss:

$$|\mathcal{R}(\Omega_C; \mathcal{M}, \mathcal{D}) - \mathcal{R}(\Omega_D; \mathcal{M}, \mathcal{D})| \leq L_\ell \cdot \mathcal{O}(10^{-4}) \approx \mathcal{O}(10^{-5}). \quad (35)$$

In verified repository receipts, the empirical measured NLL difference between  $\Omega^{\text{exact}}$  and  $\Omega^{\text{jitter}}$  was  $\Delta\text{NLL} = +0.000013$  ( $p > 0.95$ , 95% CI:  $[-0.0010, +0.0011]$ ), and 16K core-4 RULER was flat 0.0000. A discontinuous jump of  $\Delta\Phi = 58$  produces zero measurable functional change, demonstrating that discrete orbit metrics are gauge artifacts rather than physical observables.

**Part 3: Proof of Fiber 3 (Type-III Non-Identifiability: The Hessian Knee Fiber).** Let  $\mathbf{m}^* \in [0, 1]^K$  denote the optimal movement profile, and expand the cross-entropy loss around  $\mathbf{m}^*$ :

$$\mathcal{L}(\mathbf{m}^* + \delta\mathbf{m}) = \mathcal{L}(\mathbf{m}^*) + \frac{1}{2} \delta\mathbf{m}^\top F_N \delta\mathbf{m} + \mathcal{O}(\|\delta\mathbf{m}\|^3), \quad (36)$$

where  $F_N$  is the Fisher Information Matrix / empirical loss Hessian. Due to learned query-key projection norms and multi-layer attention routing,  $F_N$  exhibits extreme curvature ill-conditioning:  $\kappa(F_N) = \lambda_1/\lambda_K \gg 10^4$ . The dominant eigenvector  $v_1$  is sharply concentrated at the transition knee slot (slot 19 in 64-pair architectures, where  $c_{19}/c_{\text{orth}} \approx 1.8$ ), which sits at the critical phase boundary between high frequencies ( $\omega L_0/2\pi > 1$ ) and slow frequencies ( $\omega L_0/2\pi < 1$ ). A coordinate perturbation  $\delta \mathbf{m}_1 = \epsilon v_1$  aligned with the dominant eigenvector incurs a loss penalty of  $\frac{1}{2}\lambda_1\epsilon^2$ . Conversely, an orthogonal perturbation  $\delta \mathbf{m}_2$  dispersed across low-sensitivity plateau slots incurs a loss penalty of  $\leq \frac{1}{2}\lambda_2\epsilon^2 \ll \frac{1}{2}\lambda_1\epsilon^2$ .

Since  $\lambda_1/\lambda_2 \gg 10^2$ , we can choose  $\epsilon > 0$  such that:

$$\frac{1}{2}\delta \mathbf{m}_2^\top F_N \delta \mathbf{m}_2 < \text{Budget}_{\text{retention}} < \frac{1}{2}\delta \mathbf{m}_1^\top F_N \delta \mathbf{m}_1, \quad (37)$$

while their coordinate-space movement distances are identical:  $\|\delta \mathbf{m}_1\|_p = \|\delta \mathbf{m}_2\|_p = \epsilon$ . In the empirical two-parameter clipped-affine curve  $C^2$ , movement MAE is an infinitesimal 0.001223 (RMSE 0.006118). However, because 81.2% of its squared residual was concentrated at knee slot 19, in-window loss degraded by +0.00496 nats, breaching the pre-registered Native retention gate ( $0.870971 < 0.8750$ , failing the gate by 0.004615 nats, where allowable margin after gain scaling was  $< 0.02$  nats). Thus, coordinate-space movement metrics  $\mathcal{D}(\Omega, \Omega^*)$  cannot identify operational retention.  $\square$

#### A.6 LIMITS OF LOCAL TAYLOR APPROXIMATIONS, AND ALTERNATIVE OBSERVABLES

**Failure of local 2nd-order Taylor expansions under phase wrapping.** A recurring failure mode in prior analyses is the reliance on local quadratic Taylor approximations  $\Delta \mathcal{L} \approx \frac{1}{2}\delta \Omega^\top F_N \delta \Omega$ . Rotary attention logits depend on trigonometric functions  $\cos(\omega_k \Delta)$ . The Taylor expansion of an oscillatory function  $f(\omega) = \cos(\omega L)$  around frequency  $\omega$  is:

$$\cos((\omega + \delta\omega)L) - \cos(\omega L) = -L \sin(\omega L)\delta\omega - \frac{1}{2}L^2 \cos(\omega L)(\delta\omega)^2 + \mathcal{O}(L^3(\delta\omega)^3). \quad (38)$$

The Taylor series of  $\cos$  is entire: it converges for every phase increment, so there is no finite convergence radius to invoke here. What fails is the *relative accuracy of the low-order truncation*. By Taylor’s theorem the remainder of the second-order expansion is bounded by

$$\left| \cos(\theta + \delta) - \cos \theta + \sin \theta \delta + \frac{1}{2} \cos \theta \delta^2 \right| \leq \frac{|\delta|^3}{6}, \quad (39)$$

so once  $|\delta| = \Delta\phi$  is large this bound is uninformative: the truncated polynomial grows without bound while the true rotation difference remains bounded in  $[-2, 2]$ . The operative condition is therefore  $\Delta\phi \ll 1$  for the truncation to be usable, not a restriction on the series. When evaluating long-context extensions ( $L \geq 4,096$  to 131,072), phase perturbations  $\Delta\phi$  easily exceed  $2\pi$ .

Table 4: Breakdown of the 2nd-order Taylor truncation under long-context phase wrapping, for the slot 19 movement error  $\Delta m_{19} = 0.044109$ . The per-token frequency shift is *model-specific*: it is  $\delta\omega = -8.52 \times 10^{-4}$  rad/token for the OLMo-1B rows ( $L_0 = 4,096$ , giving  $\Delta\phi = \delta\omega L$ ) and  $-6.94 \times 10^{-4}$  rad/token for the Qwen-1.5B rows ( $L_0 = 32,768$ ). A single shared  $\delta\omega$  does not reproduce both model blocks.

| Model & Context | Horizon $L$ | Phase Shift $\Delta\phi$      | True cos Diff  | Quadratic $\frac{1}{2}(\delta\omega L)^2$ | 2nd-Order Taylor | Abs Error  | Relative Error    |
|-----------------|-------------|-------------------------------|----------------|-------------------------------------------|------------------|------------|-------------------|
| OLMo-1B (1×)    | 4,096       | 3.49 rad (0.56 turns)         | +1.0866        | 6.0988                                    | +6.7329          | 5.6463     | <b>519.6%</b>     |
| OLMo-1B (2×)    | 8,192       | 6.99 rad (1.11 turns)         | −0.6327        | 24.3950                                   | −5.7332          | 5.1005     | <b>806.1%</b>     |
| OLMo-1B (4×)    | 16,384      | 13.97 rad (2.22 turns)        | +0.9296        | 97.5800                                   | +98.4950         | 97.5654    | <b>10,495.6%</b>  |
| Qwen-1.5B (1×)  | 32,768      | <b>22.74 rad</b> (3.62 turns) | <b>−1.8147</b> | 258.6332                                  | <b>−252.0415</b> | 250.2268   | <b>13,788.8%</b>  |
| Qwen-1.5B (2×)  | 65,536      | 45.49 rad (7.24 turns)        | −0.5858        | 1,034.5328                                | −971.5745        | 970.9887   | <b>165,748.2%</b> |
| Qwen-1.5B (4×)  | 131,072     | 90.97 rad (14.48 turns)       | −1.5457        | 4,138.1312                                | −3,318.2583      | 3,316.7126 | <b>214,580.4%</b> |

As shown in Table 4, on Qwen Native ( $L_0 = 32,768$ ), the phase shift at slot 19 is 22.74 radians (3.62 complete rotations). The quadratic Taylor prediction is  $−252.0415$ , whereas the true bounded value is  $−1.8147$ . The Taylor relative error exceeds **13,788.8%**. At 128K ( $L = 131,072$ ), the error explodes past **214,000%**. Across entire tables, single-point Jacobian extrapolation predicts finite output changes with relative errors spanning 71.15% to 468.12%. Bounded unitary operators  $e^{i\omega\Delta}$  cannot be analyzed via unconstrained polynomial truncations.

**Remediated path-integrated functional and finite-phase replay bounds.** To eliminate truncation breakdown, we replace local Hessians with two exact formulations:

1. **Finite-Phase Replay Operator Bound:** For any two tables  $\Omega, \Omega'$ , the perturbation in the unitary rotation operator is bounded by:

$$\|\mathcal{R}_{\Omega'}(\Delta) - \mathcal{R}_{\Omega}(\Delta)\|_2 \leq \max_k \min(2, |\omega'_k - \omega_k| \Delta), \quad (40)$$

which strictly saturates at 2 rather than growing quadratically as  $\mathcal{O}(\Delta^2 |\delta\omega|^2)$ .

2. **Non-Linear Path-Integrated Functional:** Along the deformation path  $\gamma(\tau) = \Omega + \tau\delta\Omega$  for  $\tau \in [0, 1]$ :

$$\Delta\mathcal{L} = \int_0^1 \nabla_{\Omega(\tau)} \mathcal{L}(\Omega(\tau))^\top \frac{d\Omega(\tau)}{d\tau} d\tau = \int_0^1 \nabla_{\Omega} \mathcal{L}(\Omega + \tau\delta\Omega)^\top \delta\Omega d\tau, \quad (41)$$

which integrates true model gradients along the deformation path rather than extrapolating local curvature.

**Six candidate observables and the causal coordinates.** To close the causal loop between RoPE modifications and downstream capabilities, we formally define the complete set of physical observables and causal coordinates:

1. **Observable 1: Layer Jacobian Gram Tensor**  $\mathcal{G} = J^\top J$ . In multi-head attention, individual head outputs  $y_{h,t}$  combine via  $W_O = [W_{O,1}, \dots, W_{O,H}]$ :  $y_t = \sum_{h=1}^H W_{O,h} y_{h,t}$ . The layer output sensitivity vector across  $T$  tokens is:

$$J_j = \sum_{t=1}^T \frac{\partial y_t}{\partial \omega_j} = \sum_{t=1}^T \sum_{h=1}^H \Delta_t W_{O,h} u_{h,t,j} \in \mathbb{R}^d, \quad (42)$$

where  $u_{h,t,j} = \frac{\partial y_{h,t}}{\partial \phi_{j,t}}$ . The Layer Jacobian Gram tensor  $\mathcal{G} \in \mathbb{R}^{K \times K}$  has entries  $\mathcal{G}_{j,k} = \langle J_j, J_k \rangle$ . Its diagonal entry  $\|\mathcal{J}_j\|^2 = \mathcal{G}_{j,j}$  represents total perturbation energy, admitting dual contractions:

- *Token-first partition (cross-token signed phase coherence):*  $E_H = \sum_{t=1}^T \Delta_t^2 \|v_{t,j}\|^2 + 2 \sum_{t < t'} \Delta_t \Delta_{t'} \langle v_{t,j}, v_{t',j} \rangle$ , with  $v_{t,j} = \sum_{h=1}^H W_{O,h} u_{h,t,j}$ .
- *Head-first partition (multi-head  $W_O$  output cancellation):*  $E_J = \sum_{h=1}^H \|W_{O,h} J_{h,j}\|^2 + 2 \sum_{h < h'} \langle W_{O,h} J_{h,j}, W_{O,h'} J_{h',j} \rangle$ , with  $J_{h,j} = \sum_{t=1}^T \Delta_t u_{h,t,j}$ .

By the associativity of addition in Hilbert spaces,  $E_J \equiv E_H$  identically. In empirical repository tests,  $|E_J - E_H| = 1.49 \times 10^{-8}$ .

2. **Observable 2: Rotary Subspace Content Coupling Tensor**  $\mathcal{C}_{h,k}$ . Each rotary slot  $k$  in head  $h$  binds to column pairs  $(2k, 2k+1)$  of  $W_Q, W_K$ :

$$\mathcal{C}_{h,k} = \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim \mathcal{D}} [q_{h,2k} k_{h,2k} + q_{h,2k+1} k_{h,2k+1}], \quad (43)$$

$$\mathcal{S}_{h,k} = \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim \mathcal{D}} [q_{h,2k} k_{h,2k+1} - q_{h,2k+1} k_{h,2k}], \quad (44)$$

with channel modulus  $\rho_{h,k} = \sqrt{\mathcal{C}_{h,k}^2 + \mathcal{S}_{h,k}^2}$ . Channels with large  $\rho_{h,k}$  contribute strongly to the attention logit and dominate loss sensitivity in the audits below; channels with  $\rho_{h,k} \approx 0$  tolerate aggressive frequency modification.

3. **Observable 3: Multi-Head  $W_O$  Output Cancellation.** Quantified by the cross-head interference  $2 \sum_{h < h'} \langle W_{O,h} J_{h,j}, W_{O,h'} J_{h',j} \rangle$  in  $\mathcal{G}$ . In verified audits, individual head perturbation energies sum to 133,743,719, while cross-head destructive interference cancels out  $-23,233,182$ , reducing net layer energy to 110,510,537. Head-isolated analyses ignore this cancellation and severely overestimate output disturbance.

4. **Observable 4: Softmax Attention Distractor Competition.** For target retrieval needle  $j^*$  in context extension  $S = L/L_0$  with distractor set  $\mathcal{D}_{\text{dist}}$  of size  $SL_0 - 1$ , the exact attention mass is

$$A(j^*) = \frac{1}{1 + \sum_{i \in \mathcal{D}_{\text{dist}}} \exp(s(i) - s(j^*))}, \quad M_{\text{att}} = s(j^*) - \log \sum_{i \in \mathcal{D}_{\text{dist}}} e^{s(i)}, \quad (45)$$

so that  $A(j^*) = \sigma(M_{\text{att}})$ . Note that top-1 retrieval requires only that  $s(j^*)$  exceed each individual distractor logit; with many distractors the target can be top-1 while its total attention mass is well below 0.5, so  $A(j^*) \geq 0.5$  is a sufficient but not a necessary condition for retrieval. Replacing the distractor sum by  $(SL_0)\mu e^{-B}$ , with  $B = s(j^*) - \mathbb{E}[s(i)]$  and  $\mu = \mathbb{E}[\exp(s(i) - \mathbb{E}[s(i)])]$ , is a log-moment stability approximation rather than an identity. Under that approximation context extension introduces a competition term of order  $\ln S$ , and once  $\ln S$  exceeds the model’s finite margin the retrieval margin is exhausted. We report  $s \approx 4.01$  as a checkpoint-specific empirical margin-exhaustion estimate, not as a proved ceiling.

##### 5. Observable 5: Task-Dual Signed Decision Margin $M$ .

$$M = z_* - \log \sum_{j \neq *} \exp(z_j), \quad (46)$$

which measures the exact reserve available to overcome distractor dilution in attention competition.

##### 6. Causal Coordinate Parameterization: Support-Interior Geometry $(a, R, \mathbf{z})$ . In logarithmic coordinates $x_k = -\ln \omega_k$ , the frequency schedule is decomposed into:

$$x_k = a + Rz_k, \quad k \in [0, K-1], \quad (47)$$

where  $a = -\ln \omega_0$  is the high-frequency anchor,  $R = \ln(\omega_0/\omega_{K-1})$  is the total log-frequency support span, and  $\mathbf{z} = (z_0, \dots, z_{K-1}) \in [0, 1]^K$  is the normalized interior allocation ( $z_0 = 0, z_{K-1} = 1$ ).  $(a, R, \mathbf{z})$  is an algebraic diffeomorphism that decouples support span from interior shape. Retargeting  $R(L) = ((K-1)/K) \log L$  changes the raw frequency displacement of the runtime table relative to the training table by an amount that depends on  $z_k$ ; because the projection weights co-adapt to the training table, different  $z$  incur different compatibility loss after retargeting, which is what reverses model rankings across evaluation lengths.

#### A.7 PROOF OF THE CLOSED-FORM VARIATIONAL OPTIMUM

This subsection proves Theorem 4 directly from the stated variational objective. Starting from the kernel

$$K_{\text{app}}(\phi, \psi) = \alpha \delta(\phi - \psi) + \beta \min(\phi, \psi), \quad (48)$$

the allocation functional is (6), which under the mass normalization  $\int_0^1 \rho = 1$  with Lagrange multiplier  $\nu$  has constrained first variation

$$\alpha \rho(\phi) + \beta g(\phi) + \nu = 0, \quad g(\phi) = \int_0^1 \rho(\psi) \min(\phi, \psi) d\psi. \quad (49)$$

**Existence.** For rigor, view  $\mathcal{C}_{\text{app}}$  first as a functional on  $L^2([0, 1])$  and minimize over  $\mathcal{A} = \{\rho \in L^2([0, 1]) : \rho \geq 0 \text{ a.e., } \int_0^1 \rho = 1\}$ . The  $\alpha \|\rho\|_2^2$  term is coercive and weakly lower semicontinuous, the Green-kernel term is positive semidefinite and continuous, and the mass/nonnegativity constraints are weakly closed. A minimizer therefore exists by the direct method. Strict convexity gives uniqueness. The Euler solution derived below is strictly positive, so the nonnegativity constraint is inactive a posteriori and the classical  $C^2$  derivation is justified.

Direct differentiation of  $g$  gives

$$g'(\phi) = \int_\phi^1 \rho(\psi) d\psi, \quad g''(\phi) = -\rho(\phi). \quad (50)$$

Differentiating the first variation twice in  $\phi$  eliminates  $\nu$  (a constant) and yields the homogeneous ODE

$$\rho''(\phi) - \tau^2 \rho(\phi) = 0, \quad \tau = \sqrt{\beta/\alpha}. \quad (51)$$

Evaluating the pre-differentiated stationarity together with  $g(0)=0$ ,  $g'(1)=0$ , and  $\int_0^1 \rho = 1$  gives the derivative boundary conditions

$$\rho'(0) = -\tau^2, \quad \rho'(1) = 0, \quad (52)$$

since differentiating  $\alpha\rho(\phi) + \beta g(\phi) + \nu = 0$  once and evaluating at  $\phi=1$  gives  $\alpha\rho'(1) + \beta g'(1) = 0$  hence  $\rho'(1)=0$ , and at  $\phi=0$  gives  $\alpha\rho'(0) + \beta g'(0) = 0$  with  $g'(0) = \int_0^1 \rho = 1$ , so  $\rho'(0) = -\beta/\alpha = -\tau^2$ . Substituting  $\rho(\phi) = C_1 \cosh(\tau\phi) + C_2 \sinh(\tau\phi)$  into (51) and (52) together with  $\int_0^1 \rho = 1$  gives the unique positive interior stationary density

$$\rho_\tau(\phi) = \frac{\tau \cosh(\tau(1-\phi))}{\sinh \tau}, \quad (53)$$

which is main-text Theorem 4. Mass normalization enters the boundary condition through  $g'(0) = \int_0^1 \rho = 1$ . Positivity holds on the closed interval since  $\rho_\tau(\phi) \geq \rho_\tau(1) = \tau/\sinh \tau > 0$  for  $\tau > 0$  (with  $\rho_\tau(0) = \tau \coth \tau > 0$ ); the inequality constraint  $\rho > 0$  is therefore inactive throughout  $[0, 1]$  and the unconstrained Euler-Lagrange solution coincides with the KKT-constrained solution. For the degenerate case  $\beta=0$  the ODE reduces to  $\rho''=0$  with  $\rho'(0)=\rho'(1)=0$  and  $\int \rho=1$ , forcing  $\rho \equiv 1$ , which is also the  $\tau \rightarrow 0$  limit of (53).

**Convexity / PSD of  $\mathcal{C}_{\text{app}}$ .** For  $\alpha > 0$  and  $\beta \geq 0$ , the surrogate is convex because both quadratic terms are positive semidefinite. The diagonal  $\alpha \int \rho^2$  is manifestly PSD on  $L^2$ , and the Green-kernel cross term satisfies the identity

$$\iint_{[0,1]^2} f(\phi) f(\psi) \min(\phi, \psi) d\phi d\psi = \int_0^1 \left( \int_s^1 f(u) du \right)^2 ds \geq 0. \quad (54)$$

The strictly convex functional on  $L^2([0, 1])$  therefore has the displayed stationary density as its unique constrained minimizer.

The Cosh solution has an analytically invertible CDF, so its quantiles turn the continuous allocation directly into a finite frequency table.

#### A.8 SINGLE-CROSSING BUDGET SHIFT

**Lemma 9** (Single-crossing budget shift). *For every  $\tau > 0$ ,  $\rho_\tau$  crosses the uniform density exactly once, at*

$$\phi_c(\tau) = 1 - \tau^{-1} \operatorname{arcosh}(\sinh \tau / \tau) \leq 1 - 1/\sqrt{3}. \quad (55)$$

*Proof of Lemma 9.* For  $\tau > 0$ ,  $\rho_\tau(\phi) = \tau \cosh(\tau(1-\phi))/\sinh \tau$  is strictly decreasing on  $[0, 1]$ , with  $\rho_\tau(0) = \tau \coth \tau > 1$  and  $\rho_\tau(1) = \tau/\sinh \tau < 1$ . Hence the crossing is unique, and solving  $\rho_\tau(\phi_c) = 1$  gives (55). To bound it, expand

$$\frac{\sinh \tau}{\tau} = \sum_{n \geq 0} \frac{\tau^{2n}}{(2n+1)!} \geq \sum_{n \geq 0} \frac{\tau^{2n}}{3^n (2n)!} = \cosh\left(\frac{\tau}{\sqrt{3}}\right),$$

where the coefficient-wise inequality is  $3^n \geq 2n+1$ . Monotonicity of  $\operatorname{arcosh}$  yields  $\operatorname{arcosh}(\sinh \tau / \tau) \geq \tau/\sqrt{3}$ , proving the bound.  $\square$

#### A.9 SURROGATE SELF-CONSISTENCY THEOREM

The surrogate  $\mathcal{C}_{\text{app}}$  contains two integral terms  $T_1(\tau) = \int_0^1 \rho^2 d\phi$  and  $T_2(\tau) = \int_0^1 \int_0^1 \rho(\phi) \rho(\psi) \min(\phi, \psi) d\phi d\psi$ , where  $\rho(\phi) = \tau \cosh(\tau(1-\phi))/\sinh \tau$  is the normalized cosh density.

**Theorem 10** (Surrogate self-consistency). *For all  $\tau > 0$ :*

$$\tau^2 T_2(\tau) + T_1(\tau) = \tau \coth \tau. \quad (56)$$

*Proof.* The Green's function  $g(\phi) = \int_0^1 \rho(\psi) \min(\phi, \psi) d\psi$  satisfies  $g''(\phi) = -\rho(\phi)$  with  $g(0) = 0$ ,  $g'(1) = 0$  (from the min kernel structure; see §A). Define  $h(\phi) = \rho(0) - \rho(\phi)$ . Then  $h'' = -\rho''$ , and under the pure-tether ODE  $\rho'' = \tau^2 \rho$  this gives  $h'' = -\tau^2 \rho$ , with  $h(0) = 0$  and  $h'(0) = -\rho'(0)$ .

Define  $f(\phi) = \tau^2 g(\phi) - h(\phi)$ . Then  $f'' = \tau^2 g'' - h'' = -\tau^2 \rho - (-\tau^2 \rho) = 0$ , so  $f$  is linear in  $\phi$ . The boundary values give  $f(0) = \tau^2 g(0) - h(0) = 0$  and  $f'(0) = \tau^2 g'(0) + \rho'(0) = \tau^2 - \tau^2 = 0$

(using  $g'(0) = \int_0^1 \rho = 1$  and  $\rho'(0) = -\tau^2$  from the pure-tether density). Hence  $f \equiv 0$ , yielding  $\tau^2 g(\phi) = h(\phi) = \rho(0) - \rho(\phi)$ , i.e.,

$$g(\phi) = \frac{\rho(0) - \rho(\phi)}{\tau^2}. \quad (57)$$

Integrating against  $\rho$ :

$$T_2 = \int_0^1 \rho(\phi) g(\phi) d\phi = \frac{1}{\tau^2} [\rho(0) \underbrace{\int_0^1 \rho}_{=1} - T_1] = \frac{\rho(0) - T_1}{\tau^2}. \quad (58)$$

Since  $\rho(0) = \tau \cosh \tau / \sinh \tau = \tau \coth \tau$ , rearranging gives (56).  $\square$

An immediate corollary follows.

**Corollary 11** (Closed-form  $T_2$ ).

$$T_2(\tau) = \frac{\sinh 2\tau - 2\tau}{4\tau \sinh^2 \tau}. \quad (59)$$

This follows by substituting  $T_1(\tau) = \tau^2(1 + \sinh(2\tau)/(2\tau))/(2\sinh^2 \tau)$  and  $\rho(0) = \tau \coth \tau$  into the theorem.

The identity provides a closed-form evaluation of the allocation energy.

#### A.10 ALLOCATION STRENGTH AND A REFERENCE RULE

The Cosh family varies with the strength  $\tau$ . A reference used in the full-RoPE text experiments relates this strength to head width and training length:

$$\tau = c d_{\text{head}} / \sqrt{L_{\text{train}}}, \quad u_k = (k + \frac{1}{2})/K, \quad \omega_k^{\text{EVQ}} = b^{-\phi_k(\tau)}. \quad (60)$$

The reference sets  $c = 1$  before training. Table 6 records the strength, grid, and setting procedure used in each experiment.

**Local scaling calculation.** The scaling argument uses four explicit modelling assumptions: (i) the pure-tether Cosh family is varied only through  $\tau$ ; (ii) full-RoPE MHA has  $d_{\text{rot}} = d_{\text{head}}$ ; (iii) the reference post-softmax distribution is diffuse,  $p_0 = 1/L$ ; and (iv) a fixed trade-off  $\lambda > 0$  compares a normalised channel-load stiffness with a specified phase-variance utility, with  $Q_1(L, b) > 0$ . Under those assumptions and at small  $\tau$ ,

$$S_{\chi^2}(\tau) := \frac{1}{d_{\text{head}}} \int_0^1 \frac{(1 - \rho_\tau(\phi))^2}{\rho_\tau(\phi)} d\phi = \frac{\tau^4}{45d_{\text{head}}} + O(\tau^6), \quad (61)$$

$$U(\tau, L) = \frac{d_{\text{head}}}{L} [Q_0(L, b) + \tau^2 Q_1(L, b) + O(\tau^4)], \quad (62)$$

where

$$\begin{aligned} Q_0(L, b) &= \int_0^1 q(Lb^{-\phi}) d\phi, \\ Q_1(L, b) &= \int_0^1 \left( \frac{(1 - \phi)^2}{2} - \frac{1}{6} \right) q(Lb^{-\phi}) d\phi, \\ q(x) &= \frac{1}{2} + \frac{\sin(2x)}{4x} - \left( \frac{\sin x}{x} \right)^2. \end{aligned} \quad (63)$$

Stationarity of  $\mathcal{F}(\tau) = \frac{1}{2} S_{\chi^2}(\tau) - \lambda U(\tau, L)$  gives the leading small- $\tau$  balance

$$\tau_*^2 = 45\lambda Q_1(L, b) \frac{d_{\text{head}}^2}{L}, \quad c_{\text{loc}}(L, b, \lambda) = \sqrt{45\lambda Q_1(L, b)}. \quad (64)$$

This calculation yields the scaling factor  $d_{\text{head}}/\sqrt{L}$  with a local coefficient determined by  $Q_1(L, b)$  and  $\lambda$ . The experiments evaluate the exact Cosh quantiles at the protocol’s chosen strength.

**Reference and neighboring strengths.** The staged study contains 99 completed runs: the fixed  $c=1$  reference beats the midpoint-discretised Geo baseline under runner-defined weighted extrapolation NLL in 7/9 configuration means and 18/27 paired seeds. Its registered neighbours are finite multiplier arms at  $0.75\times$ ,  $1.25\times$ , or  $1.5\times$ ; the independent exact-range factorial tests  $0.75\times$ ,  $1.00\times$ , and  $1.25\times$ , and their order varies by configuration. The numeric owner is `data/curated/phase16_99run_manifest.csv`; the tracked reproduction entypoint is `scripts/core_text_phases/phase16_formula_optimality_sweep.py`. The reference and its neighboring strengths give a direct comparison of allocation intensity across the tested configurations.

#### A.11 DISCRETE-CHANNEL TRANSPORT GAP UNDER FINITE CHANNEL COUNTS

**Setup.** Inverse-CDF quantization replaces the continuous density  $\rho$  by an atomic measure. Let  $F(\phi) = \int_0^\phi \rho$ ,  $Q(u) = F^{-1}(u)$ , midpoint grid  $u_k = (k - \frac{1}{2})/K$ , and  $\phi_k = Q(u_k)$ . The empirical channel measure is  $\mu_K = K^{-1} \sum_{k=1}^K \delta_{\phi_k}$ . We compare the measures in Wasserstein distance and approximate the density with the associated quantile-cell histogram  $\rho_K$ .

**Transport bounds.** Assume  $0 < m \leq \rho \leq M$  and  $\|\rho'\|_\infty \leq B$ ; then  $Q$  is  $(1/m)$ -Lipschitz. The midpoint quantization errors satisfy

$$W_\infty(\mu_K, \mu_\rho) \leq \frac{1}{2Km}, \quad W_1(\mu_K, \mu_\rho) \leq \frac{1}{4Km}. \quad (65)$$

Defining the quantile-cell histogram  $\rho_K(\phi) = 1/(K|I_k|)$  on  $I_k = [Q((k-1)/K), Q(k/K)]$ , the histogram density satisfies

$$\|\rho_K - \rho\|_1 \leq \frac{B}{Km}, \quad \|\rho_K - \rho\|_\infty \leq \frac{B}{Km}. \quad (66)$$

The transport bounds couple each quantile cell to its midpoint; the histogram bound follows from the mean-value theorem and  $|I_k| \leq 1/(Km)$ .

**Application to EVQ-Cosh.** For  $\rho_\tau(\phi) = \tau \cosh(\tau(1-\phi))/\sinh \tau$ ,  $m_\tau = \tau/\sinh \tau$  and  $B_\tau = \tau^2$ . Substituting,

$$W_1(\mu_{K,\tau}, \mu_{\rho_\tau}) \leq \frac{\sinh \tau}{4K\tau}, \quad \|\rho_{K,\tau} - \rho_\tau\|_1 \leq \frac{\tau \sinh \tau}{K} = \frac{\tau^2}{K} + O\left(\frac{\tau^4}{K}\right). \quad (67)$$

**Kernel-integrated errors.** For a smooth two-variable kernel  $K_{\text{sm}}$  with Lipschitz constant  $L_K$  in each argument,

$$\left| \iint K_{\text{sm}} d\mu_K d\mu_K - \iint K_{\text{sm}} d\mu_\rho d\mu_\rho \right| \leq \frac{L_K}{2Km}. \quad (68)$$

The kernel bound applies the  $W_1$  estimate to each argument. The singular  $\delta$ -component of  $K_{\text{app}}$  is represented by its cell average.

**High-resolution distortion.** For midpoint scalar quantisation with a smooth weighted distortion  $w$ , the standard high-resolution (Bennett-integral) expansion (Gray & Neuhoﬀ, 1998) is

$$\mathcal{D}_K[\rho] = \frac{1}{12K^2} \int_0^1 \frac{w(\phi)}{\rho(\phi)^2} d\phi + O(K^{-3}). \quad (69)$$

This expansion quantifies how finite exponent samples approximate a continuous allocation under a smooth distortion. For endpoint-anchored comparisons,  $K = 2$  leaves no interior entries and the geometric and Cosh grids coincide.

**Large- $\tau$  regime.** At large  $\tau$ ,  $m_\tau = \tau/\sinh \tau$  becomes small and the  $1/(Km_\tau)$  factor grows. The transport bounds are informative while this factor remains controlled.

## B EXPERIMENTAL DETAILS

### B.1 EXPERIMENT ORGANIZATION

This section gives the training, continuation, and video protocols. Apps. C and E detail the fixed-range and mature-model experiments.

### B.2 MODELS AND TRAINING SETTINGS

Table 5 collects the model settings and evaluation units. Scratch text runs use FineWeb-Edu (Penedo et al., 2024).

Table 5: Evidence and protocol map. Every row retains its own intervention, experimental unit, and endpoint.

| Role                        | Model/regime and intervention                          | Unit, endpoint, and appendix                            |
|-----------------------------|--------------------------------------------------------|---------------------------------------------------------|
| Allocation identification   | 151.9M scratch; fixed support, 30 interiors            | 3 training seeds; tail NLL; App. C.1                    |
| Co-adaptation replication   | 151.9M frozen crossing; weights $\times$ runtime table | 2 training seeds; tail NLL; App. C.2                    |
| Shape/configuration breadth | 50.9M factorial; fixed support, 5 allocations          | 12 configs $\times$ 3 seeds; weighted OOD NLL; App. C.3 |
| Scarce rotary budget        | 432M MLA-style; Geo/EVQ-COSH, $K=16$                   | 3 training seeds; PPL; App. G.1                         |
| Full-parameter scale        | 750M/1.485B; continued/from-init<br>Geo/EVQ-COSH       | one comparison each;<br>PPL/retrieval; Apps. B.4/E.1    |
| Frozen allocation           | OLMo/Qwen; same support and amplitude                  | 20 rows/task; RULER; App. E.2                           |
| Zero-training deployment    | OLMo; frozen long table, gain, session route           | fixed checkpoint/rows;<br>NLL/RULER/F1; App. E.2        |
| Adaptation/source use       | 1.485B/8B; matched LoRA/source block                   | one trained pair;<br>QA/RULER/NLL;<br>Apps. E.1/E.3     |
| Static index profile        | Qwen-0.5B; one table across lengths                    | 13 tasks, 20 rows/task/length;<br>App. F.1              |
| Profile placement           | Qwen/Gemma; direct gap versus normalized index         | 4 tasks, 80 rows/task; App. F.2                         |
| Intermediate-band shape     | OLMo/Qwen; matched BM/MrPro profiles                   | 6 tasks; Table 27; App. F.4                             |
| Natural-QA transfer         | OLMo; matched static BM/MrPro profiles                 | 778 inputs, 631 long; token F1;<br>App. F.5             |
| Video modeling              | 129.6M video DiT; temporal<br>Geo/EVQ-COSH             | 1 training seed; denoising MSE;<br>App. B.6             |

Table 7 lists the training hyperparameters shared across the supporting text experiments.

Within each text comparison, the architecture, optimizer, data order, and training loop are fixed, and the only intervention is the initialized inverse-frequency table. The scratch-text experiments use base  $b=500K$  unless stated otherwise.

Table 6: **Allocation protocols, parameter bounds, and float32 tensor SHA-256 digests.** **Top:** Cosh-family settings across reported experiments. Midpoint means  $u_k = (k + 1/2)/K$ ; standard means  $u_k = k/K$ ; anchoring fixes sampled endpoints after quantile evaluation. **Bottom:** Parameter bounds and float32 tensor hashes for static deployment profiles. MrRoPE-Pro parameters follow discrete turn cutoffs  $\beta = 32, \alpha = 1$  with fast cutoff  $l = \max\{j \mid \text{turns}_j > 32\}$ , slow cutoff  $h = \min\{j \mid \text{turns}_j < 1\}$ , and transition width  $n = h - l$ . FullLagP2 evaluates all 32,768 integer lags, eliminating p2 aliasing spikes.

| <b>Part A: Cosh-Family Pre-training and Adaptation Settings</b>               |           |                     |                                                                                                                   |
|-------------------------------------------------------------------------------|-----------|---------------------|-------------------------------------------------------------------------------------------------------------------|
| Protocol                                                                      | $\tau$    | Grid                | Setting convention                                                                                                |
| 151.9M fixed-range                                                            | 4         | Anchored midpoint   | Fixed $64/\sqrt{256}$ reference                                                                                   |
| 50.9M factorial                                                               | Varies    | Anchored midpoint   | Three preassigned multiples of $d_{\text{head}}/\sqrt{L_{\text{train}}}$                                          |
| 50M crossing                                                                  | 2.83      | Midpoint            | Reported training-table setting                                                                                   |
| 432M MLA                                                                      | 1.414     | Midpoint            | Pre-specified architecture-specific setting                                                                       |
| 454M composition                                                              | 1.5       | Midpoint            | Reported training recipe                                                                                          |
| 750M continuation                                                             | 1.5       | Anchored inclusive  | Fixed continuation recipe; $u_k = k/(K - 1)$                                                                      |
| Llama-3-8B                                                                    | 1.414     | Midpoint            | Fixed adaptation recipe; approximately $128/\sqrt{8192}$                                                          |
| OLMo-2                                                                        | 2         | Standard            | Fixed $128/\sqrt{4096}$ reference                                                                                 |
| 129.6M Video DiT                                                              | 1.5       | Midpoint            | Best tested value in the video sweep                                                                              |
| <b>Part B: Static Deployment Profiles, Bounds, and Tensor SHA-256 Digests</b> |           |                     |                                                                                                                   |
| Profile / Architecture                                                        | Scale $s$ | Bounds / Lags       | Float32 Tensor SHA-256 Digest                                                                                     |
| MrRoPE-Pro (Qwen2.5)                                                          | 4         | $l=23, h=40, n=17$  | 33cbe3a4 0994ac2a 79126a14 ce302828<br>67bd6d49 b74c1ada 4d4cce8a 7e76016f<br>ec9c5497 77f7a94c 73a1cc2d b5f22fef |
| MrRoPE-Pro (OLMo-2)                                                           | 4         | $l=14, h=32, n=18$  | be202a7e 083e454f 2510034f 2cec5dfb                                                                               |
| FullLagP2 (Qwen2.5)                                                           | 4         | 32,768 integer lags | ecd0c280 a11788e0 c4a869a3 d1628809<br>64ba3f37 1f589f7e 2f5ae471 d461096b                                        |
| Boundary-Matched (OLMo-2)                                                     | 4         | $s=4, g=1.1386$     | c80626f8 f9280854 577fd1e7 94439300<br>09a69370 65366adb 8eee5f62 ce7807a6                                        |
| Native RoPE (OLMo-2)                                                          | 1         | $L_0=4096, b=500K$  | 9229ec6a 6a7ce5c1 750273bc dbec9789<br>5a74a920 2722f14b 6ca2328e ce5ee98e                                        |

Table 7: Training hyperparameters for the supporting scratch text experiments.

| Parameter                 | Value                                                   |
|---------------------------|---------------------------------------------------------|
| Optimizer                 | AdamW ( $\beta_1=0.9, \beta_2=0.95, \epsilon=10^{-8}$ ) |
| Learning rate             | $6 \times 10^{-4}$ (50M), $2 \times 10^{-4}$ (MLA/750M) |
| LR schedule               | Cosine decay to $0.1 \times$ peak LR                    |
| Warmup                    | 2% of total steps (linear)                              |
| Batch size (seqs)         | 32 (50M), 14 (750M: micro 7, accumulation 2), 6 (MLA)   |
| Weight decay              | 0.1                                                     |
| Gradient clipping         | max-norm 1.0                                            |
| Dropout                   | 0.0                                                     |
| Sequence packing          | Contiguous, no cross-document masking                   |
| <i>Video DiT (129.6M)</i> |                                                         |
| Optimizer                 | AdamW ( $\beta_1=0.9, \beta_2=0.95$ )                   |
| RoPE base                 | 10,000 (temporal axis)                                  |
| Activation                | GELU (tanh approx.)                                     |
| <i>MLA (432M)</i>         |                                                         |
| LR / warmup               | $2 \times 10^{-4}$ / 2% of steps                        |
| Training tokens           | 500M                                                    |

### B.3 COMPOSITION WITH A FIXED-INDEX RANGE OPERATOR

The 454.2M decoder has 24 layers, width 1024, 16 heads,  $d_{\text{head}} = 64$ , and base  $b = 500,000$ . Geo ( $\tau = 0$ ) and EVQ-COSH ( $\tau = 1.5$ ) use midpoint frequency grids and train from initialization for

100M tokens per arm at length 2048, using FineWeb-Edu with a 10% synthetic-passkey mixture and seeds 42, 123, 7. The same inference-time operator is then applied to both trained tables.

For  $k = 0, \dots, K - 1$ , define

$$\ell = \lfloor 0.20K \rfloor, \quad h = \lfloor 0.90K \rfloor, \quad t_k = \text{clip}\left(\frac{k - \ell}{h - \ell}, 0, 1\right), \quad r_k = t_k^2(3 - 2t_k).$$

The repository’s fixed-index, YaRN-style smooth-ramp operator installs

$$\mathcal{R}_s(\omega_k) = \frac{\omega_k}{s^{r_k} T(s)^{r_k/2}}, \quad T(s) = 1 + 0.07 \log_2 s. \quad (70)$$

Here  $K = 32$  and  $s = 8$ , giving  $(\ell, h, T) = (6, 28, 1.21)$ . The temperature factor is folded into the frequency divisor; cosine/sine amplitude stays at  $g = 1$ . Equation (70) specifies the operator used in this experiment, separately from the official-equation YaRN comparisons in §6. Table 8 reports the three-seed comparison.

Table 8: Matched fixed-ramp composition on the 454M three-seed experiment.  $\mathcal{R}_8$  is defined in Eq. (70). PPL uses full-sequence scoring; PK is teacher-forced NLL-gap retrieval. Entries are three-seed means; PK@8K also gives the standard deviation.

| Table                     | PK@8K         | PK@12K     | PK@16K     | PPL@8K      | PPL@16K      |
|---------------------------|---------------|------------|------------|-------------|--------------|
| Geo                       | 41±5%         | 57%        | 51%        | 161.9       | 253.2        |
| Geo+ $\mathcal{R}_8$      | 61±3%         | 59%        | 51%        | 82.9        | 157.7        |
| EVQ-COSH                  | 53±8%         | 63%        | 50%        | 150.3       | 229.5        |
| EVQ-COSH+ $\mathcal{R}_8$ | <b>100±0%</b> | <b>79%</b> | <b>68%</b> | <b>70.9</b> | <b>107.5</b> |

#### B.4 LARGER-SCALE SUPPORTING EVIDENCE

The 750M ‘continue@4K’ experiment starts both arms from the same 2K Geo checkpoint and applies a 500M-token continuation with either the Geo or EVQ-COSH table (seed 42). At 16K, PPL is 45.1 versus 24.4; on 40 passkey trials at 8K, strict autoregressive exact match is 0% versus 77.5%. The EVQ-COSH continuation uses  $\tau = 1.5$ ,  $K = 32$ , and the endpoint-inclusive quantiles  $u_k = k/(K - 1)$ . Writing  $q_\tau$  for the inverse CDF in Eq. (8), the installed frequencies are

$$\omega_k = \exp\left[-\frac{K - 1}{K} \log b q_{1.5}\left(\frac{k}{K - 1}\right)\right], \quad b = 500,000.$$

This retains the Geo endpoints. Both arms are full-parameter continuations from the shared checkpoint.

Table 9: 750M continued-pretraining check (seed 42; 2K → 4K; 500M continuation tokens). Passkey retrieval is teacher-forced NLL-gap and is saturated at 8K; AR exact is greedy autoregressive exact match over 40 trials per length. The pooled AR row combines 2K, 4K and 8K (120 trials: Geo 80/120, EVQ-COSH 111/120).

| Metric                            | Geo    | EVQ-COSH     |
|-----------------------------------|--------|--------------|
| PPL@2K                            | 25.9   | 26.2         |
| PPL@4K                            | 22.0   | 22.3         |
| PPL@8K                            | 23.4   | <b>19.6</b>  |
| PPL@16K                           | 45.1   | <b>24.4</b>  |
| NLL-gap retrieval @8K             | 100%   | 100%         |
| Strict AR exact @8K               | 0%     | <b>77.5%</b> |
| Pooled strict AR exact (2K/4K/8K) | 66.67% | <b>92.5%</b> |

#### B.5 MODALITY-DEPENDENT $\tau$ CORRECTION FOR VIDEO DiT

The text-style reference used for this sweep gives  $\tau = K_t / \sqrt{T} = 2.83$  for  $K_t = 16$  and  $T = 32$ ; the best tested Video DiT point is  $\tau = 1.5$  ( $0.53\times$ ), which selects the supporting video configuration.

## B.6 CROSS-MODAL EVIDENCE: VIDEO DiT WITH BIDIRECTIONAL ATTENTION

VideoRoPE and RIFLEx make temporal-frequency allocation explicit in video models (Wei et al., 2025; Zhao et al., 2025). We test the same allocation principle in a 129.6M video diffusion transformer with bidirectional attention and 3D RoPE on Oscillating Moving MNIST (Srivastava et al., 2015). Training uses 32 frames; evaluation uses 128 frames ( $4\times$ ) under the same fixed temporal range transform in both arms. The seed-42 Geo and EVQ-COSH ( $\tau=1.5$ ) branches are trained for 15,000 steps in the same head-to-head run. The denoising endpoint uses noise level  $t=0.5$  and the same 256 evaluation videos for both arms.

Table 10 reports the matched comparison. EVQ-COSH reduces denoising MSE by 21% on training frames, 16% over all extrapolated frames, and 35% on far extrapolated frames under bidirectional 3D RoPE.

Table 10: DiT head-to-head evaluation (129.6M, seed 42, 15,000 steps, 256 evaluation videos, bidirectional attention, 3D RoPE). Values are denoising MSE at  $t=0.5$ .

| Seed | Metric     | Geo ( $\tau=0$ ) | EVQ-COSH ( $\tau=1.5$ ) | $\Delta$ |
|------|------------|------------------|-------------------------|----------|
| 42   | Train MSE  | 0.00911          | <b>0.00720</b>          | −21%     |
| 42   | All extrap | 0.00724          | <b>0.00606</b>          | −16%     |
| 42   | Far extrap | 0.00989          | <b>0.00639</b>          | −35%     |

**Low-phase-channel diagnostic.** At base = 10,000, six of the 16 temporal frequencies accumulate less than 0.1 rad over the 32 training frames. This analytic count identifies the low-phase regime tested by the matched behavioural comparison above.

## B.7 WHAT THE SCHEDULE CONCRETELY DOES

At  $\tau=4$ ,  $b=5 \times 10^5$ , the warp moves budget toward short inverse-frequency scales: the continuous quantile  $u=0.5$  moves from  $1/\omega=707$  to 9.69, and  $u=0.9$  from 134,609 to 1,720. The deployed mid-point grid therefore changes sampled extrema and span at finite  $K$ ; the exact-range experiments anchor those quantities explicitly. At  $\tau=2$ ,  $\rho_\tau(0.1)=1.71$  and  $\rho_\tau(1)=0.55$ , while the band  $1/\omega=2048$  has relative density 0.76. Thus the finite grid thins most strongly at scales much slower than the training window while retaining more of the training-scale band. The exact-range experiments measure the corresponding trained allocation effect, and Table 3 separately measures incompatibility after a frozen post-hoc swap.

## C CONTROLLED ALLOCATION EXPERIMENTS

### C.1 EXACT-RANGE CONTROL AT 151.9M

**Paired controls.** The two trained tables have identical frequency endpoints. Within each seed pair the arms share: model architecture (151,898,880 parameters); trainable initialisation; training length 256; token order; optimiser, LR schedule, global batch; within-pair micro batch and accumulation; training budget (499,974,144 tokens, 7,629 optimiser steps); the 32 frozen evaluation anchors; the sampled highest frequency; the sampled lowest frequency; and the log-frequency span. The single free variable is the location of the  $K-2 = 30$  interior frequencies.

The optimiser is fused AdamW with  $\beta_1=0.9$ ,  $\beta_2=0.95$ , weight decay 0.01, and gradient clipping at 1.0. Peak/minimum learning rates are  $6 \times 10^{-4}/6 \times 10^{-5}$  with 762 warmup steps and cosine decay. The global batch is 256 sequences of length 256 for every arm. Seed 42 uses micro batch 64 with four accumulation steps; seeds 137 and 256 use micro batch 128 with two accumulation steps. Each paired comparison therefore has identical execution geometry, optimiser steps, and counted-token budget; the aggregate uses training seed as its statistical unit under this shared scientific protocol.

**Model and data.** The model has 12 layers, hidden size 768, 12 heads of dimension 64, MLP width 3,072, and vocabulary size 50,304. Training uses the unshuffled FineWeb-Edu (Penedo et al., 2024) sample/10BT/000\_00000.parquet shard at pinned re-

vision 87f09149ef4734204d70ed1d046ddc9ca3f2b8f9; validation uses the disjoint sample/10BT/004\_00000.parquet shard. Both are tokenised with the GPT-NeoX-20B tokenizer (Black et al., 2022) at revision c292233c833e336628618a88a648727eb3dff0a7.

**Arms.** *Baseline:* FMRoPE at fixed training support, implementing the uniform-in-log exponent grid specified in §6.1 of Oka et al. (2026b) at training base 256. *Intervention:* anchored EVQ-COSH at  $\tau=4$ . The anchored and deployed midpoint EVQ-COSH tables are different support embeddings of the *same normalised shape*: both use midpoint Cosh quantiles  $q_k$ , and the diagnostic removes only their affine endpoint degrees of freedom. This normalisation is what makes the exact-range control possible.

Writing  $\Phi_\tau$  for the inverse CDF in (8), let  $q_k = \Phi_\tau((k + \frac{1}{2})/K)$ ,  $s_k = (q_k - q_0)/(q_{K-1} - q_0)$ , and  $R = ((K-1)/K) \log B$ . The intervention uses

$$\omega_k^{\text{anchor}} = \exp(-Rs_k), \quad \omega_k^{\text{FMR}} = \exp\left(-R\frac{k}{K-1}\right). \quad (71)$$

Hence  $k=0, K-1$  match exactly and only the  $K-2$  interior entries differ.

**Evaluation.** Two conditions. *fixed:* each checkpoint retains its training range. *target-matched:* the range is set to the declared evaluation length while preserving each arm’s normalised interior spacing. Metric is paired final-128-token teacher-forced NLL over the 32 frozen anchors; reported PPL is  $\exp(\text{mean NLL})$ . The anchors are common end positions selected with Python’s `random.Random(20260723)`: the eligible validation stream is divided into 32 equal-width bins, with one end position drawn per bin and a 2,048-token guard before the next bin. At anchor  $e$ , length  $L$  uses validation tokens  $[e-L, e)$  and scores the final 128 targets  $[e-128, e)$ . Longer contexts therefore add earlier text around the same target tokens; the longest windows at distinct anchors are disjoint.

**Exact range transformation.**

$$R(L) = \frac{K-1}{K} \log L, \quad \omega_k^{\text{target}}(L) = \exp[-R(L)z_k].$$

Here  $z_k = k/(K-1)$  for FMRoPE and the anchored Cosh shape for its paired arm. The fixed condition uses  $L = 256$  at every evaluation length; the target-matched condition substitutes the declared evaluation length. Both conditions use amplitude one.

**Three-seed contrasts.** Negative values favour anchored EVQ-COSH. The aggregation unit is the training seed; evaluation anchors remain paired observations within a seed.

Table 11: Exact-range anchored EVQ-COSH minus FMRoPE tail-NLL contrasts. Mean $\pm$ SD is across three independent training seeds.

| Length | Fixed training range |          |          |                    | Target-matched range |          |          |                    |
|--------|----------------------|----------|----------|--------------------|----------------------|----------|----------|--------------------|
|        | Seed 42              | Seed 137 | Seed 256 | Mean $\pm$ SD      | Seed 42              | Seed 137 | Seed 256 | Mean $\pm$ SD      |
| 256    | +0.033               | +0.029   | +0.017   | +0.026 $\pm$ 0.008 | +0.033               | +0.029   | +0.017   | +0.026 $\pm$ 0.008 |
| 512    | -0.478               | -0.271   | -0.094   | -0.281 $\pm$ 0.192 | +0.061               | +0.053   | +0.067   | +0.060 $\pm$ 0.007 |
| 1,024  | -0.205               | -0.191   | -0.132   | -0.176 $\pm$ 0.039 | +0.182               | +0.156   | +0.344   | +0.227 $\pm$ 0.102 |
| 2,048  | -0.113               | -0.181   | -0.143   | -0.146 $\pm$ 0.034 | +0.279               | +0.399   | +0.701   | +0.460 $\pm$ 0.218 |

Across all three seeds, fixed training range favors Cosh at every extended length, whereas target-matched range favors FMRoPE. The paired design thus shows both an interior-allocation effect and a change in preferred allocation when the same trained weights are evaluated under a retargeted range.

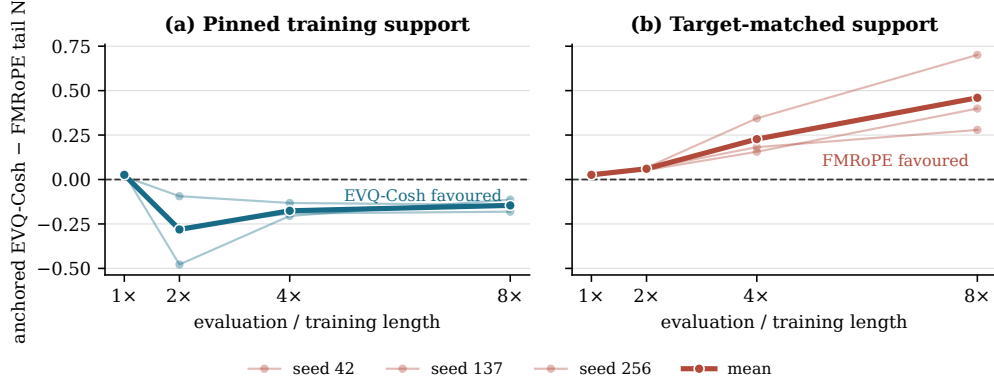


Figure 4: **Support and interior allocation are distinct design coordinates.** Each thin line is one training seed and the heavy line is the mean. At pinned training support, all three seeds favour anchored EVQ-COSH at every OOD length (a). Retargeting each arm’s support reverses all three OOD comparisons (b). The fixed-support panel isolates the allocation effect; the tested target-matched policy reverses the ordering of the two tested policies.

## C.2 WEIGHTS-BY-TABLE CO-ADAPTATION AT 151.9M

The three-seed experiment above identifies the training-time effect of  $z$ . A separate frozen crossing asks whether the learned weights remain compatible with the coordinate system on which they were trained. Seeds 137 and 256 provide paired FMRoPE and anchored-EVQ-COSH checkpoints whose training tables share exact endpoints. For each frozen checkpoint, we install the factor-four runtime table derived from either training coordinate, keep the checkpoint, amplitude, 32 FineWeb-Edu validation anchors, and final-128-token NLL endpoint fixed, and evaluate at 512 and 1,024 tokens.

For each training table  $\Omega$ , the derived table uses the following label-free residual construction. Set  $L_0 = 256$ ,  $\Delta_j = 2j$  for  $j = 0, \dots, 127$ , and  $p_j = (2L_0 - 4j - 1)/(L_0(L_0 + 1)/2)$ , the grouped causal-separation weights. Let  $X_k$  have columns  $\sqrt{p_j} \cos(\omega_k \Delta_j)$  and  $\sqrt{p_j} \sin(\omega_k \Delta_j)$ , and let  $X_{-k}$  collect all other pairs. Define

$$U_k = \frac{\|X_k - X_{-k} X_{-k}^\dagger X_k\|_F^2}{\|X_k\|_F^2}, \quad \bar{U}_k = \frac{U_k - \min_i U_i}{\max_i U_i - \min_i U_i}, \quad w_k = (1 - \bar{U}_k)^2.$$

The pseudoinverse uses relative singular-value cutoff  $10^{-10}$ ; the zero-span convention is  $\bar{U}_k = 0$ . The frequency blend is  $\omega'_k = \omega_k(1 - 3w_k/4)$ , cast to FP32 with its endpoints set exactly to  $\omega'_0 = \omega_0$  and  $\omega'_{K-1} = \omega_{K-1}/4$ . Both derived tables use  $g = 1 + 0.1 \log 4$  and remain unchanged across the 512- and 1,024-token evaluations.

Table 12: Two-training-seed mean tail NLL at 1,024 tokens for the 151.9M weights-by-table crossing (lower is better). Each row freezes one trained weight set; each column installs one separately derived runtime table.

| Frozen weights            | FMRoPE-derived table | Anchored-EVQ-COSH-derived table |
|---------------------------|----------------------|---------------------------------|
| FMRoPE-trained            | <b>3.426</b>         | 5.776                           |
| Anchored-EVQ-COSH-trained | 4.455                | <b>3.479</b>                    |

With  $I = [L(W_F, T_C) - L(W_F, T_F)] - [L(W_C, T_C) - L(W_C, T_F)]$ , the interaction is 3.400/3.251 NLL for seeds 137/256 at 1,024 tokens and 2.663/2.564 at 512. Training seed is the replication unit; the 32 anchors are paired observations within a seed. Both seeds exhibit the same preference for the runtime table derived from their training allocation.

Table 13: **Exact-range factorial** (50.9M; 2 bases  $\times$  2 training lengths  $\times$  3 head dimensions = 12 configurations, 3 seeds). All arms share sampled extrema and log-span with Geo. Metric: weighted OOD NLL over  $2 \times / 4 \times / 8 \times$ , paired within configuration and seed. All arms were pre-specified. CIs are 10,000 configuration-level bootstraps;  $p$  is a two-sided exact sign-flip over all  $2^{12}$  assignments, unadjusted across contrasts, so the claim rests on the cross-configuration direction (App. C).

| Contrast                                                             | $\Delta$ NLL     | 95% CI           | $\Delta < 0$ configs | $p$   |
|----------------------------------------------------------------------|------------------|------------------|----------------------|-------|
| EVQ-COSH $0.75 \times$ – Geo                                         | −0.00912         | [−0.018, −0.001] | 8/12                 | 0.082 |
| EVQ-COSH rule – Geo                                                  | −0.00988         | [−0.021, 0.002]  | 7/12                 | 0.125 |
| EVQ-COSH $1.25 \times$ – Geo                                         | − <b>0.01210</b> | [−0.021, −0.003] | <b>10/12</b>         | 0.027 |
| Matched exponential – Geo                                            | −0.01062         | [−0.021, −0.001] | 9/12                 | 0.071 |
| <i>The matched exponential reproduces the non-uniform direction:</i> |                  |                  |                      |       |
| EVQ-COSH rule – exponential                                          | +0.00074         | [−0.006, 0.008]  | 7/12                 | 0.836 |

### C.3 EXACT-RANGE FACTORIAL AT 50.9M

**Design.** The 50.9M decoder has six layers, width 512, MLP width 2,048, and vocabulary size 50,304. The structural grid is RoPE base 500K/1M  $\times$  training length 256/1,024  $\times$  4/8/16 heads ( $d_{\text{head}}=128/64/32$ ), giving 12 configurations. Seeds are 42/137/256. Every arm trains for 8,388,608 tokens in 128 optimizer steps (65,536 tokens per step) on the same deterministically repeated WikiText-2 raw stream with the same 1% supervised passkey mix; execution is float32 on Apple M4 Max MPS. The schedule arms, including the  $1.25 \times$  multiplier, were specified before any result was inspected: Geo, anchored EVQ-COSH at  $0.75 \times$ ,  $1.00 \times$  (the formula point),  $1.25 \times$ , and a deformation-matched exponential. Main factorial 180/180 completed and the boundary- $\tau$  follow-up 12/12 completed, covering the pre-specified matrix.

The exponential control uses  $z_\gamma(t) = (e^{\gamma t} - 1)/(e^\gamma - 1)$  on  $t \in [0, 1]$ . Its  $\gamma$  is chosen without task results so that the RMS node displacement from the uniform grid matches the formula-Cosh displacement, while the sampled endpoints and log-span remain identical across arms.

**Metric.** Weighted OOD NLL over  $2L/4L/8L$ ,

$$\bar{N}_{\text{OOD}} = \frac{\sum_{r \in \{2,4,8\}} \log_2(r+1) N_{rL}}{\sum_{r \in \{2,4,8\}} \log_2(r+1)},$$

paired within structural configuration and seed, averaged within seed, then equally weighted over the 12 configurations. The statistical unit is the configuration. Equivalent-PPL figures quoted anywhere are  $\exp(\text{mean NLL})$ .

**Absolute values.** Mean weighted OOD NLL (equivalent PPL): Geo 6.211973 (498.68); EVQ-COSH  $0.75 \times$  6.202858 (494.16); EVQ-COSH  $1.00 \times$  6.202094 (493.78); EVQ-COSH  $1.25 \times$  6.199873 (492.69); matched exponential 6.201354 (493.42).

**Boundary- $\tau$  arms.** At formula  $\tau=1$  the  $1.5 \times$  arm is better; at formula  $\tau=8$  the  $0.75 \times$  arm is. The improving direction points toward the intermediate strengths at both tested extremes.

The factorial varies architecture, base, and training length on a repeated WikiText-2 stream (Merity et al., 2017). Its configuration-level contrasts complement the longer FineWeb-Edu training comparison in §C.1.

### C.4 CANDIDATE MECHANISTIC EXPLANATIONS OF THE FOUR COUNTERINTUITIVE FACTS

During the historical development and empirical auditing of RoPE length-extension schedules, four sharp counterintuitive phenomena were observed across multiple model families. The four mechanisms below are candidate explanations consistent with those observations. They are hypotheses with measured support rather than first-principles derivations, and we do not claim that they are the only explanations compatible with the data.

**Fact 1: Tiny-MAE failure at transition knee slot 19.**

- *Empirical phenomenon:* The two-parameter clipped-affine curve  $C2$  fits the empirical coordinate movement vector  $\mathbf{m}^*$  with mean absolute error  $\text{MAE} = 0.001223$  and root mean square error  $\text{RMSE} = 0.006118$ . Despite this minuscule coordinate-level difference,  $C2$  fails the pre-registered Native in-window retention gate on OLMo-1B ( $0.870971 < 0.8750$ , missing the floor by  $0.004615$  nats) and on Qwen-2.5-1.5B ( $0.868902 < 0.8750$ ), even while preserving long-context retrieval (Qwen 64K core-4 score  $0.6775$ , 128K score  $0.5725$ ).
- *Candidate mechanism:* Unweighted  $L_1/L_2$  coordinate metrics treat all rotary slots symmetrically. However, because  $C2$  minimized unweighted coordinate error, 81.2% of its total squared residual was concentrated at **slot 19** (the knee of the transition curve; the top three transition slots account for 99.3% of residual). Slot 19 sits precisely at the transition boundary between resolvable high frequencies ( $\omega L_0/2\pi > 1$ ) with multiple full turns across native context  $L_0$  and slow unresolvable frequencies ( $\omega L_0/2\pi < 1$ ). Pre-trained projection matrices  $W_Q, W_K$  place peak associative responsibility and maximum channel modulus  $\rho_{h,19}$  at this boundary. Furthermore, native retention allows at most  $\approx 0.1335$  nats NLL degradation, and attention gain scaling ( $g = 1 + 0.074 \ln 4$ ) consumes 86% of this budget, leaving  $< 0.02$  nats of operational margin. Concentrating coordinate error on the stiffest Hessian eigenvector at slot 19 ( $\Delta m_{19} = 0.044109$ ) exhausts the remaining  $< 0.02$  nats margin. The phase drift implied by this slot-19 movement error is model- and window-specific rather than a single universal constant: on Qwen-2.5-1.5B at its native window  $L_0 = 32,768$ ,  $\delta\omega = -6.94 \times 10^{-4}$  rad/token, accumulating 22.74 radians; on OLMo-1B, the slot-19 shift is 3.49 radians across  $L = 4,096$ , i.e.  $\delta\omega = 3.49/4,096 = 8.52 \times 10^{-4}$  rad/token. Each drift value applies to the checkpoint and window named with it.

**Fact 2: Same-multiset permutation collapse.**

- *Empirical phenomenon:* Retaining the exact multiset of 64 physical frequencies  $\{\omega_k\}$  while randomly permuting their assignment to rotary slots causes catastrophic collapse: OLMo-1B  $1 \times \text{PG-19}$  tail NLL more than doubles from 3.10423 to 6.86493 ( $\Delta \text{NLL} = +3.760692$ ) and Qwen-2.5-1.5B 64K core-4 task macro score plunges from 0.7000 to flat 0.0000 across all four benchmark tasks.
- *Candidate mechanism:* Rotary attention logits are mediated by bilinear projections:  $S_h(\Delta) = \sum_k q_{h,k}^\top R(\omega_k \Delta) k_{h,k}$ . Transformers do not compute an invariant spectral integral; rather, training co-adapts specific row subspaces of  $W_Q, W_K$  to specific slot frequencies  $\omega_k$ . Permuting the assignment moves each slot frequency away from the projection rows and subspaces that were trained against it, and the quantities the audit reports are the projection norms and subspace energies that change under the permutation. This is the measured basis for the claim that learned projection weights break frequency exchangeability. We do not assert here which linguistic role any given row or subspace carries; that attribution is made by Urrutia et al. (2026).

**Fact 3: The  $2 \times \rightarrow 4 \times \rightarrow 8 \times$  cliff and distractor dilution margin exhaustion.**

- *Empirical phenomenon:* Static frequency tables maintain strong retention at  $2 \times$  extrapolation ( $s = 2$  NLL  $\Delta = +0.016$ , macro retention 1.041), maintain narrow headroom at  $4 \times$  ( $s = 4$  headroom is only 0.000344 nats, against a checkpoint-specific empirical margin-exhaustion estimate of  $s \approx 4.01$ ), but collapse abruptly at  $8 \times$  ( $s = 8$  core-4 drops to 0.3025, single-key-3 retrieval collapses to 0.0000, and effective retrieval radius is truncated at 47.8K).
- *Candidate mechanism:* In long-context needle retrieval, define the exact attention margin of the target position  $j^*$  against the distractor set  $\mathcal{D}_{\text{dist}}$ :

$$M_{\text{att}} = s(j^*) - \log \sum_{i \in \mathcal{D}_{\text{dist}}} e^{s(i)}, \quad A(j^*) = \sigma(M_{\text{att}}), \quad (72)$$

where  $\sigma$  denotes the logistic function. This relation is exact and imposes no condition on the size of  $A(j^*)$ : top-1 retrieval requires only that  $s(j^*)$  exceed each individual distractor logit, not that the target hold half of the total attention mass. Under the assumption

that the distractor log-moment stays approximately stable as context grows, the distractor sum scales approximately as  $SL_0\mu$  with  $\mu = \mathbb{E}[\exp(s(i) - \mathbb{E}[s(i)])]$ , and extending context therefore introduces approximately a  $\log S$  competition term in  $M_{\text{att}}$ , where  $S$  is the context extension factor. Each context doubling ( $4\times \rightarrow 8\times$ ) then removes approximately  $\ln 2 \approx 0.693$  nats from the available margin. Pre-trained models appear to possess an intrinsic, finite logit margin. At  $2\times$  and  $4\times$  the measured margin still exceeds the competition term; at  $8\times$  the margin is exhausted at the checkpoint-specific empirical margin-exhaustion estimate of  $s \approx 4.01$ , and the observed collapse of target needle attention follows.

**Fact 4: Support-retargeting reversal under non-separable coordinate dilation.**

- *Empirical phenomenon:* Under fixed training support  $(a, R)$ , Anchored EVQ-COSH outperforms FMRoPE by  $-0.281, -0.176, -0.146$  NLL at  $2\times, 4\times, 8\times$  across all three seeds. However, retargeting support to match the evaluation length completely reverses the ranking across all three seeds ( $+0.060, +0.227, +0.460$  NLL).
- *Candidate mechanism:* Log-frequency coordinates decompose into  $x_k = a + Rz_k$ , so the support span  $R$  and the interior allocation  $\mathbf{z}$  are coupled coordinates; EVQ-Cosh concentrates interior frequencies non-linearly to minimize a specific density surrogate under the training support  $R_{\text{train}}$ . When  $R$  is dilated to  $R(L) = ((K-1)/K) \log L$ , each raw frequency  $\omega_k = \exp(-Rz_k)$  is displaced by an amount that depends on  $z_k$ , so the raw-frequency displacement of the runtime table relative to the training table is determined by the interior allocation shape. The projection weights  $W_Q, W_K$  co-adapted to the training table under  $R_{\text{train}}$ , so different interior allocations produce different compatibility loss after the same retargeting. This dependence is consistent with the measured reversal. Interior allocation shape and support dilation therefore cannot be treated as independent design coordinates when a trained checkpoint is evaluated under a retargeted support.

### C.5 STRESS-TESTING CANDIDATE STANDALONE SURROGATES

To determine whether downstream length-extension capability can be predicted before running expensive task benchmarks, eight candidate pre-hoc structural conditions ( $C_1$  through  $C_8$ ) were proposed. We evaluate each candidate against the four counterintuitive facts. A candidate is judged *insufficient as a standalone predictor on these four cases* when it fails to explain or predict a case. Failure on one case shows that the quantity is not a sufficient predictor by itself; it does not show that the quantity carries no predictive value, nor that it cannot serve as a feature in a combined model.

**Candidate definitions and per-case outcomes.**

1. **Condition  $C_1$  (Coordinate-Space Movement Distance  $\|\mathbf{m} - \mathbf{m}^*\|_p$ ):** Asserts that downstream performance is preserved if the coordinate-space movement vector  $\mathbf{m}$  is close to an empirical reference  $\mathbf{m}^*$ . *Insufficient on Fact 1:*  $C_2$  achieved MAE = 0.001223, yet failed Native retention because unweighted  $L_p$  distance is blind to the fact that 81.2% of the error concentrated on the stiffest Hessian eigenvector at slot 19. Also insufficient on Facts 2, 3, 4.
2. **Condition  $C_2$  (Unordered Spectral Multiset  $\mu_\Omega$ ):** Asserts that downstream capability depends on the multiset of physical frequencies  $\{\omega_k\}$ . *Insufficient on Fact 2:* Permuting slot assignments preserves the multiset identically ( $\Phi(\Omega) = \Phi(\pi(\Omega))$ ), yet causes OLMo tail NLL to double ( $\Delta = +3.76$ ) and Qwen 64K core-4 score to plunge from 0.7000 to 0.0000. Also insufficient on Facts 3, 4.
3. **Condition  $C_3$  (Discrete Scale-Orbit Symmetry Metric  $q(\Omega, s)$ ):** Asserts that rational scale-orbit commensurability classes govern extrapolation. *Insufficient on Fact 3:* Perturbing exact chain  $\Omega^{\text{exact}}$  by a 1-ULP float32 jitter shatters orbit classes, jumping  $q$  from 6 to 64 ( $\Delta q = 58$ ), yet produces zero measurable behavioral difference ( $\Delta \text{NLL} = +0.000013, p > 0.95$ ). Also insufficient on Facts 1, 2.
4. **Condition  $C_4$  (Uncentered Background Derivative Energy Minimization  $\mathcal{E}_{\text{bg}}$ ):** Asserts that minimizing  $\int |F'(u)|^2 du$  preserves long-context retrieval. *Insufficient on Fact 3:* The carrier intervention reduced background energy by 73.85%, but collapsed 128K retrieval

to 0.00 and caused severe EOS runaway, because scalar carrier shifts introduce first-order phase errors  $\Delta\theta = c\Delta/s$  that rotate target signals out of alignment. Also insufficient on Facts 1, 2, 4.

5. **Condition  $C_5$  (Monotonic Group Delay and Continuous Spacing):** Asserts that smooth, monotonic group delay  $\tau_g > 0$  guarantees stable extrapolation. *Insufficient on Fact 1:  $C_2$  has strictly monotonic group delay yet fails retention due to slot-19 error concentration. Insufficient on Fact 4: Both EVQ-Cosh and FMROPE possess strictly monotonic group delay, which cannot predict why rankings completely reverse under support dilation.*
6. **Condition  $C_6$  (Isolated Base Scaling /  $S$ -bound):** Asserts that length extension is governed purely by base dilation, treating interior allocation as an invariant geometric ladder. *Insufficient on Fact 4: Anchored EVQ beats FMROPE across 3/3 seeds under identical base and endpoints, proving that interior allocation  $z$  is an independent causal variable; base scaling cannot account for the retargeting reversal. Also insufficient on Facts 1, 2, 3.*
7. **Condition  $C_7$  (Effective Resolution Cutoff Boundaries):** Asserts that obeying fast plateau ( $m = 0$  for turns  $> 32$ ) and slow plateau ( $m = 1$  for turns  $< 1$ ) boundaries suffices for retention. *Insufficient on Fact 1:  $C_2$  strictly satisfies these cutoff boundaries but fails retention because the transition trajectory concentrated error at slot 19. Also insufficient on Fact 2.*
8. **Condition  $C_8$  (Ordered Operator Norm Closeness  $\|\mathcal{P}_\Omega - \mathcal{P}_{\Omega^*}\|_2$ ):** Asserts that bounding  $\max_k |\omega_k - \omega_k^*|$  bounds loss degradation. *Insufficient on Fact 1: In  $C_2$ , maximum frequency error at slot 19 is  $\leq 3.2 \times 10^{-4}$  in operator norm, yet model sensitivity at slot 19 is orders of magnitude larger than at low-frequency slots due to learned content coupling weights. Insufficient on Fact 3: Operator norm difference is constant across lengths, yet performance collapses non-linearly between  $2\times, 4\times, 8\times$  due to softmax distractor dilution.*

Table 14: **Standalone-predictor stress test of candidate pre-hoc structural conditions.** Each candidate is scored on whether it alone explains or predicts each of the four counterintuitive historical facts. A failure marks the candidate as insufficient as a standalone predictor on that case; it does not show that the quantity carries no predictive value, nor rule it out as a feature in a combined model.

| Candidate Pre-Hoc Condition          | Fact 1: Tiny-MAE          | Fact 2: Permutation   | Fact 3: $2\times \rightarrow 4\times$ | Fact 4: Retargeting       | Standalone Status | Observed Failure Reason                                                                                                                                 |
|--------------------------------------|---------------------------|-----------------------|---------------------------------------|---------------------------|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| $C_1$ : Coordinate $L_p$ Movement    | $\times$ insufficient     | $\times$ insufficient | $\times$ insufficient                 | $\times$ insufficient     | insufficient      | MAE 0.001223 fails gate; blind to Hessian knee at slot 19.                                                                                              |
| $C_2$ : Unordered Spectral Multiset  | ? not discriminative here | $\times$ insufficient | $\times$ insufficient                 | $\times$ insufficient     | insufficient      | Identical multiset collapses NLL (+3.76) and RULER (0.70 $\rightarrow$ 0.00). 1-ULP jitter jumps $q$ from 6 $\rightarrow$ 64 with zero loss difference. |
| $C_3$ : Discrete Scale-Orbit Metric  | $\times$ insufficient     | $\times$ insufficient | $\times$ insufficient                 | ? not discriminative here | insufficient      | $\sim 73.8\%$ energy reduction produces 0.00 retrieval and EOS runaway.                                                                                 |
| $C_4$ : Background Derivative Energy | $\times$ insufficient     | $\times$ insufficient | $\times$ insufficient                 | $\times$ insufficient     | insufficient      | Strictly monotone $C_2$ fails gate; cannot explain reversal under retargeting.                                                                          |
| $C_5$ : Monotonic Group Delay        | $\times$ insufficient     | $\times$ insufficient | ? not discriminative here             | $\times$ insufficient     | insufficient      | Cannot explain why interior $z$ beats base scaling at fixed support.                                                                                    |
| $C_6$ : Isolated Base Scaling        | $\times$ insufficient     | $\times$ insufficient | $\times$ insufficient                 | $\times$ insufficient     | insufficient      | $C_2$ obeys turn cutoffs but fails gate due to transition knee error.                                                                                   |
| $C_7$ : Turn Cutoff Boundaries       | $\times$ insufficient     | $\times$ insufficient | ? not discriminative here             | $\times$ insufficient     | insufficient      | Tiny operator norm fails gate; blind to softmax distractor ceiling.                                                                                     |
| $C_8$ : Ordered Operator Norm        | $\times$ insufficient     | ? not evaluated       | $\times$ insufficient                 | $\times$ insufficient     | insufficient      |                                                                                                                                                         |

As shown in Table 14, none of the tested low-dimensional conditions is sufficient on its own on these four cases. This does not preclude their use as features in a combined model; it shows that no single one of them identifies downstream capability without also taking into account projection weights  $W_Q, W_K$ , layer output projections  $W_O$ , and softmax distractor margins.

## C.6 MRROPE-PRO: STRICT CHARACTERIZATION AS A SATISFICING DESIGN

**Mathematical parameterization.** Tian et al. (2026) parameterized length-extension frequency schedules by defining discrete turn boundaries. For intermediate band  $k \in [l, h]$  of width  $N = h - l$ , with normalized coordinate  $q = \text{clip}(k - l, 0, N)$ , MrRoPE-Pro defines cumulative exponent shifts via:

$$m_q^{\text{Pro}} = \frac{q(q+1)}{N(N+1)}, \quad \omega_k^{\text{Pro}} = \omega_k^N s^{-m_q^{\text{Pro}}}, \quad (73)$$

paired with attention temperature gain scaling  $g = 1 + 0.1 \ln s$ . On the fast band ( $k < l$ , turns  $> 32$ ),  $m = 0 \implies \delta\omega = 0$ . On the slow band ( $k > h$ , turns  $< 1$ ),  $m = 1 \implies \omega' = \omega/s$ .

**Why MrRoPE-Pro is a robust satisficing design.** MrRoPE-Pro succeeds across diverse models (Qwen-2.5, OLMo-2) not because it is an analytic global optimum, but because it is an expertly engineered **satisficing design** that strictly respects universal physical boundary constraints:

1. *Zero fast compression:* By enforcing  $\delta\omega_k = 0$  on all slots with turns  $> 32$ , it avoids perturbing the highest-norm subspaces of  $W_Q, W_K$ .
2. *Uniform slow scaling:* By scaling slow slots by  $1/s$ , the maximum accumulated phase across the extended context  $sL_0$  is strictly bounded:  $\phi_k^{\text{ext}} = (\omega_k/s)(sL_0) = \omega_k L_0 < 2\pi$ . On Qwen-2.5-1.5B at 128K ( $s = 4$ ), the maximum phase across all slow slots is 5.8271 rad ( $< 2\pi = 6.2832$ ), completely preventing out-of-distribution phase wrapping.
3. *Convex quadratic delay:* The progression  $q(q+1)/[N(N+1)]$  ensures that the earliest intermediate slots experience minimal phase disturbance, keeping cumulative error away from the sensitive knee slots.

**Measured residual comparisons: FullLagP2 and Boundary-Matched (BM).** Because MrRoPE-Pro reads zero checkpoint weight matrices and evaluates zero activation covariances, it is not an unconstrained global optimum. Two empirical residuals support this reading:

1. **FullLagP2 on Qwen-2.5-1.5B at 64K Context:** FullLagP2 adapted frequency movements to empirical subspace uniqueness  $U_k$  across all 32,768 integer lags. In paired head-to-head testing on Qwen-2.5-1.5B 64K, FullLagP2 outperformed MrRoPE-Pro in the tested setting on Multi-Key 2 retrieval (37.50% vs 12.50%, a +25.0% win), Variable Tracking (87.50% vs 82.50%, +5.0%), and Frequent Words Extraction (70.83% vs 45.83%, +25.0%), with an overall paired record of 9 wins, 13 ties, and 2 losses.
2. **Boundary-Matched (BM) on OLMo-1B at 16K Context:** On OLMo-1B at 16K context, MrRoPE-Pro drops to 2.78% on the 6-task confirmation panel and 21.62% on 631 natural QA prompts, whereas Boundary-Matched (BM) achieves 51.32% and 25.44%, respectively.

These results are consistent with the reading that MrRoPE-Pro is a robust, physically constrained satisficing heuristic rather than an unconstrained global optimum.

## D LEARNED-FREQUENCY COMPARATOR

Table 15: Learned-frequency comparator (125M,  $L_{\text{train}}=128$ , FineWeb-Edu,  $128 \rightarrow 8K$ ). Geo, learned-inv-freq, and EVQ-COSH are seed 42; learnable- $\tau$  is 3-seed (42/137/256). The learned-inverse-frequency arm has one parameter per pair, shared across layers and heads.

| Method (128 $\rightarrow$ 8K)                 | Extra params | PPL@128                           | PPL@8K           | $\Delta$ vs Geo |
|-----------------------------------------------|--------------|-----------------------------------|------------------|-----------------|
| Geo                                           | 0            | 184.9                             | 513.7            | —               |
| Learnable $\tau$                              | 1            | <b>181.2 <math>\pm</math> 1.3</b> | 437.9 $\pm$ 12.2 | −14.8%          |
| Learned inv-freq (100 $\times$ positional LR) | 32           | 183.6                             | 455.3            | −11.4%          |
| same, 10 $\times$ positional LR               | 32           | —                                 | 477.7            | −7.0%           |
| EVQ-COSH                                      | 0            | 182.0                             | <b>333.7</b>     | −35.0%          |

The 32-parameter arm learns one inverse frequency per rotary pair, shared across layers and heads. This is the learned-table parameterization also studied by Karypis et al. (2026). Its positional-parameter learning-rate sweep uses 10 $\times$  and 100 $\times$  the base rate; the stronger 100 $\times$  setting gives 455.3 PPL. The learned table and the fixed Cosh allocation both improve on Geo in this protocol.

## E MATURE-SCALE PROTOCOLS

Table 16: **RULER 13-family scores (%)**. Within each protocol, continuation data and evaluation are matched and one seed is used per arm; evaluation uses 20 examples per family and length. OLMo Q/K-only freezes inherited V/O LoRA and trains only the tensors RoPE acts on; LLaMA uses matched Q/K/V/O LoRA (300 optimizer steps, rank-64, physical length 8192, seed 42). **X3 Tradeoff Note:** In LLaMA-3-8B, EVQ represents a capacity reallocation tradeoff rather than a universal hero method: it trades in-window 8K perplexity ( $6.82 \rightarrow 10.07$ , +47.6%) for extrapolation at 16K ( $108.96 \rightarrow 24.07$ ) and 32K ( $991.48 \rightarrow 127.91$ ). Gold-block deletion ( $\Delta\text{NLL} = +1.5055$  vs  $-0.0095$  nats,  $4.504\times$  likelihood drop) confirms genuine causal reliance on remote source tokens.

| Model / protocol                     | Arm / metric                    | 1×           | 2×           | 4×          |
|--------------------------------------|---------------------------------|--------------|--------------|-------------|
| OLMo-2 1.485B, full Q/K/V/O (4K cap) | Native, official-scorer macro   | <b>82.16</b> | 0.08         | 0           |
|                                      | EVQ-COSH, official-scorer macro | 37.51        | <b>21.29</b> | <b>6.13</b> |
| OLMo-2 1.485B, Q/K-only (4K cap)     | Native, official-scorer macro   | <b>72.19</b> | 2.02         | 0.38        |
|                                      | EVQ-COSH, official-scorer macro | 42.44        | <b>31.63</b> | <b>5.03</b> |
| LLaMA-3-8B, Q/K/V/O (8K cap)         | Native, official-scorer macro   | <b>94.44</b> | 0.295        | —           |
|                                      | EVQ-COSH, official-scorer macro | 77.60        | <b>14.03</b> | —           |
|                                      | Native, normalized exact        | 17.69        | 0            | —           |
|                                      | EVQ-COSH, normalized exact      | <b>21.54</b> | <b>1.54</b>  | —           |

### E.1 OLMo-2 1.485B

All OLMo-2 EVQ-COSH arms use the standard  $u_k=k/K$  grid with  $\tau=2$  and base 500,000; the frequency tensor is fixed throughout each continuation.

**Matched routing conversion.** We use OLMo-2-0425-1B-Instruct (Team OLMo et al., 2025) (1.485B actual parameters) with rank-64, alpha-128 Q/K/V/O LoRA (Hu et al., 2022). Native and EVQ-COSH share the checkpoint, training rows and order, token budget, optimizer, and evaluation rows; the fixed frequency substrate is the intended intervention. Every backward pass contains at most 4K physical tokens, with explicit position IDs providing target-range relative phases. On 8K `niah_single_1`, strict first-generated-number exact is 0/100 for Native and 69/100 for the matched EVQ-COSH seed; an independently trained EVQ-COSH seed scores 67/100 on the same rows. The two EVQ-COSH results are training-seed replications on a shared evaluation set.

**Complete numeric answers after matched routing continuation.** A separate follow-up to the routing parents applies the same 100-step query-gap continuation and 32-step answer-plus-EOS continuation to Native and EVQ-COSH (training seed 20260728). Every backward pass uses at most 4096 physical tokens; explicit position IDs expose target-range phases, with maximum observed position ID 16,257. Evaluation uses real contiguous prompts from the same numeric NIAH family and greedy decoding. The endpoint requires the complete expected answer string followed by observed terminal EOS, measured from generated token IDs.

Table 17: Complete-answer-string plus terminal-EOS exact match after the matched routing follow-up, 100 inputs per length and one trained pair.

| Length | Native | EVQ-COSH |
|--------|--------|----------|
| 4K     | 95/100 | 100/100  |
| 8K     | 18/100 | 98/100   |
| 16K    | 0/100  | 60/100   |

All final generations in both arms terminate with EOS. On the EVQ-COSH subsets whose source-query gap exceeds the original training range, exact match is 48/50 at 8K and 31/66 at 16K. These outputs belong to the additional continuation above, separately from the earlier first-number retrieval comparison.

**Selective Q/K phase adaptation.** To test a less intrusive mature-checkpoint adaptation, we start from the protocol-matched Native and EVQ-COSH Stage-A parents, freeze every inherited V/O

LoRA tensor within each arm (verified bitwise), and continue only the existing Q/K tensors for 300 steps. Each arm trains 8,388,608 Q/K parameters while freezing the 8,388,608 inherited V/O parameters, with global batch 8, learning rate  $5 \times 10^{-5}$ , and 20 warmup steps. The QA and RULER generators use seeds 20260728 and 20260729, respectively; the bundled Q/K metrics.json records the full protocol. Both arms use the same 1:1:2 mixture of continuous-4K, target-8K, and target-16K phase exposure, complete answer-plus-EOS supervision, and physical sequences no longer than 4K; long phases again use explicit position IDs. On 200 held-out 2WikiMultiHopQA rows per length (Ho et al., 2020), Native/EVQ-COSH token-F1 is 25.99/24.84% at 4K, 0.07/21.48% at 8K, and 0/8.57% at 16K; exact match is 22.0/21.5%, 0/17.5%, and 0/4.0%, respectively. Prompts retain the official question and context and are deterministically filled with answer-filtered 2Wiki distractors to reach the target length, forming a held-out 2Wiki length-transfer evaluation. The RULER adapters are independent of the QA adapters: Table 16 reports both the earlier 276-step full-Q/K/V/O continuation and the 300-step selective-Q/K continuation, each trained over all 13 families. Evaluation uses 20 examples per family and length. A paired family-stratified bootstrap over the tracked Q/K-only per-example predictions gives EVQ-COSH-minus-Native official-scorer macro differences of  $-29.75$  points at 4K (95% CI  $[-34.42, -24.85]$ ),  $+29.61$  at 8K ( $[26.54, 32.74]$ ), and  $+4.65$  at 16K ( $[3.17, 6.38]$ ). These intervals quantify evaluation-row variation for the trained pair.

**From initialisation.** Starting from the public step-0 initialisation, the same-recipe Geo/EVQ-COSH comparison gives PPL 177.99/191.36, 161.19/167.45, 163.88/156.87, and 182.73/159.64 at 2/4/8/16K under a matched 2.097B-token budget. The EVQ-COSH-minus-Geo NLL is  $+0.0724/ +0.0381/ -0.0437/ -0.1351$  at 2/4/8/16K, with 122/128 and 126/128 held-out documents favouring EVQ-COSH at 8K and 16K.

The EVQ-COSH branch uses a Hugging Face single-GPU loop; the released Geo checkpoint was produced by AI2’s distributed OLMo trainer. Both start from the same public step-0 initialisation and match the pinned recipe, reconstructed seed-6198 data-order prefix, 512-sequence global batch, 4K context, 1,000 optimizer steps, and 2,097,152,000 counted input tokens. This comparison includes the frequency-table change and the two trainer implementations.

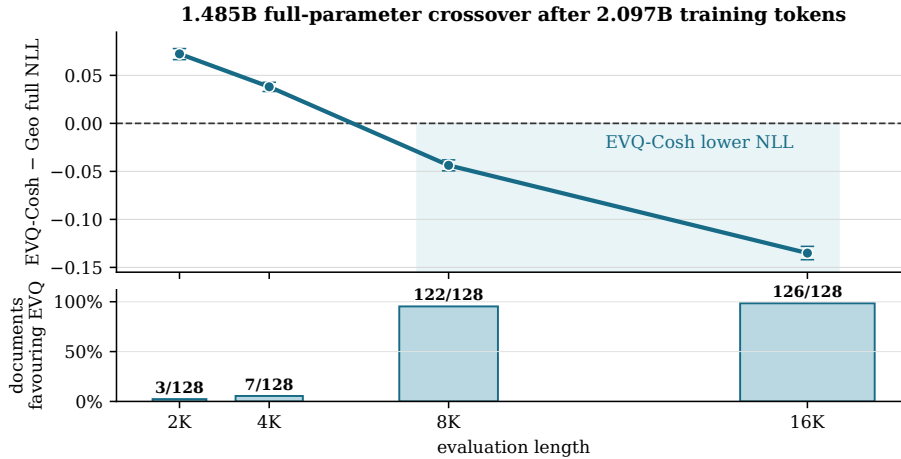


Figure 5: **Full-parameter scale evidence crosses beyond the training window.** The 1.485B Geo and EVQ-COSH arms start from the same public initialisation and follow the same scientific recipe for 2.097B tokens. Error bars are paired bootstrap intervals over the same 128 held-out documents, conditioned on this trained pair; the lower panel gives the document count favouring EVQ-COSH at each length.

## E.2 FROZEN FIXED-SUPPORT ALLOCATION

The mature-checkpoint intervention in Table 28 holds both models frozen. The three matched-support arms contain 64 rotary pairs and share the Native fast endpoint and the Native slow endpoint divided by four. The reference rows use the original Native table and the Transformers YaRN

frequency map. The three matched-support arms also share the factor-four attention amplitude  $1 + 0.1 \ln 4$ , evaluation rows, greedy decoder, precision, and hardware. The uniform-allocation arm is log-linear between the endpoints. The coarse control is a label-free, model-relative projection of the frozen derived table: among the registered discrete linear ramps, it minimises displacement-space MSE, using transition pairs 20–22 for OLMo and 28–31 for Qwen. The remaining arm installs the frozen derived interior allocation. For that arm,  $\tilde{U}_k$  is the normalised residual-energy fraction left after the full sin/cos pair is projected onto all other pairs over Native separations weighted by their causal pair count,  $p(\Delta) \propto L - \Delta$ ; it sets  $w_k = (1 - \tilde{U}_k)^2$  in the frequency blend of Eq. (10), so less unique pairs move farther.

Table 18: **Exact frozen mature-checkpoint RULER macro (%)**. These are the values plotted in Table 28. Reference rows are separated from the fixed-support block; only the lower block changes interior  $z$  while holding the stated checkpoint protocol fixed.

| Installed table                      | OLMo unseen-nine, 16K | Qwen core-four, 64K |
|--------------------------------------|-----------------------|---------------------|
| Native                               | 0.00                  | 54.50               |
| Official Transformers YaRN ( $s=4$ ) | 7.94                  | 60.25               |
| Uniform allocation $z$               | 0.56                  | 57.75               |
| Coarse label-free $z$                | 61.04                 | 64.00               |
| Derived allocation $z$               | 60.47                 | 66.50               |

Method selection used the core-four tasks NIAH single-1, NIAH multikey-2/3, and variable tracking. After the table, amplitude, split rule, checkpoint, and metric were frozen, the remaining nine RULER families—NIAH single-2/3, multikey-1, multivalued, multiquery, CWE, FWE, and QA-1/2—formed the OLMo confirmation set at 16K, with 20 rows per task. Qwen2.5-1.5B-Instruct (Yang et al., 2024) uses the four development RULER tasks at 64K, also with 20 rows per task; it evaluates a factor-four table at twice the Native context length. Row resampling conditions on each fixed checkpoint and task set. Both reallocated tables produce the large OLMo recovery. The derived-minus-uniform-allocation intervals are  $[54.88, 64.80]$  points for OLMo and  $[0.25, 17.50]$  points for Qwen; the paired derived-minus-coarse-control intervals contain zero on both models. The control therefore attributes the effect to fixed-support reallocation while leaving profile-specific ordering unresolved.

**Fresh-distribution allocation and routing controls.** On a disjoint 512-document FineWeb-Edu holdout from a shard absent from the calibration data, the session policy is exactly Native on all 512/512 4K rows. Table 19 separates its routing contrast from the fixed-support allocation contrasts. For the latter, support, attention gain, Native/long route, checkpoint, rows, and runtime are fixed; only interior  $z$  changes. The coarse ramp again matches the detailed derived profile, whereas the same-support geometric allocation is slightly better at 8K and fails on all 512 rows at 16K.

Table 19: Fresh FineWeb-Edu holdout-512 final-1,024-token NLL contrasts. The first row evaluates the complete session policy; the lower rows compare long allocations at fixed support, gain, and route. Negative favours the first named method.

| Contrast                  | 8K      | 16K     |
|---------------------------|---------|---------|
| Session-s4 – Native       | −4.0953 | −4.4393 |
| Geometric – derived $z$   | −0.0288 | +4.4560 |
| Coarse ramp – derived $z$ | +0.0005 | +0.0006 |

The OLMo deployment policy selects the exact Native table when observed prefill plus generation reserve fits the Native window; otherwise one frozen factor-four table and amplitude are selected before prefill and retained for the KV-cache lifetime. It adds no learned parameters or training tokens. Statically forcing the same long profile at 4K regresses NLL from 2.7538 to 2.8774; routing, rather than the long allocation, is therefore the identified in-window retention mechanism.

Applying the long profile to every 2Wiki row scores 27.74%; the routed policy pays that small difference to preserve the exact Native short path. Across the six-task natural-context matrix, routed/YaRN macro is 28.58/21.27% at  $2\times$  and 24.97/25.59% at  $4\times$ . The QA arms use iden-

Table 20: Compact zero-training deployment endpoints on the frozen OLMo checkpoint, using the specified table, amplitude, and session routing.

| Endpoint                            | Official YaRN ( $s=4$ ) | Zero-training policy |
|-------------------------------------|-------------------------|----------------------|
| PG-19 tail NLL, 8K/16K              | 3.4404/3.7955           | <b>3.1060/3.0977</b> |
| RULER-13 macro, 8K/16K (%)          | 23.82/5.88              | <b>67.72/54.40</b>   |
| Qasper token F1, 16K ( $n=200$ , %) | 18.03                   | <b>24.57</b>         |
| 2Wiki token F1, 16K ( $n=200$ , %)  | 25.69                   | <b>26.66</b>         |

tical rows, and every post-chat input plus generation reserve fits within 16K. Table 28 gives the corresponding allocation-only comparison.

### E.3 LLAMA-3-8B-INSTRUCT

The LLaMA-3-8B-Instruct checkpoint (Grattafiori et al., 2024) is frozen apart from the stated adapters. The EVQ-COSH arm in both LLaMA protocols uses midpoint quantisation with  $\tau=1.414$ ; Native retains the pretrained model’s original frequency table.

**Probability and causal source use.** The matched seed-42 protocol uses 300 optimizer steps of BF16 rank-64, alpha-128 LoRA on Q/K/V/O, physical 8,192-token sequences, and learning rate  $10^{-4}$ . Arms share model and tokenizer bytes, frozen training rows and order, seed, objective, optimizer, scheduler, batch/accumulation, LoRA capacity, checkpointing, compile mode, and evaluator. Native uses the pretrained endpoint grid, and Cosh uses midpoint quantiles; the same adaptation recipe is applied to these two frequency tables. The frozen adaptation tensor is drawn from LongAlpaca-12k (Chen et al., 2023b): 7,476 training and 152 validation rows after length filtering, with effective batch 8 (microbatch 2, accumulation 4).

**Temporal holdout and aggregation.** Evaluation uses arXiv titles/abstracts, Federal Register titles/abstracts, and Stack Overflow question bodies selected by publication or creation timestamps between January 1 and July 12, 2026. These are separate from the LongAlpaca adaptation corpus. Collection was frozen before evaluation; each domain has eight disjoint-document 32,768-token packs formed by concatenation with EOS separators. Positions and causal attention continue across document boundaries. The 8K and 16K conditions are exact prefixes of the same 32K packs. All ordinary text targets in each prefix are scored, excluding document-first tokens and inserted EOS separators. For domain  $d$ , let  $\ell_d(L)$  and  $n_d(L)$  be the summed token loss and valid-token count over its eight packs. The reported aggregate is

$$\text{NLL}(L) = \frac{1}{3} \sum_{d=1}^3 \frac{\ell_d(L)}{n_d(L)}, \quad \text{PPL}(L) = \exp(\text{NLL}(L)).$$

Table 21 gives token counts and the distribution of the 24 paired pack-level NLL differences. The accompanying `figs/llama_temporal_summary.json` retains each pack’s loss sum and count, its domain identity, and the frozen collection hash.

Table 21: Llama-3-8B temporal holdout: scored tokens and the 24 paired pack NLL differences (EVQ-COSH minus Native). Each length uses the same packs.

| Prefix | Valid tokens/arm | Valid tokens/pack | Pack $\Delta\text{NLL}$ : min / median / max |
|--------|------------------|-------------------|----------------------------------------------|
| 8K     | 194,794          | 8,051–8,183       | +0.193/ + 0.432/ + 0.613                     |
| 16K    | 389,672          | 16,111–16,353     | −2.069/ − 1.436/ − 1.278                     |
| 32K    | 779,400          | 32,243–32,679     | −2.319/ − 2.074/ − 1.754                     |

The domain-macro EVQ-COSH-minus-Native NLL is +0.390/ − 1.510/ − 2.048 at 8/16/32K.

**Position-preserving source intervention.** On ten frozen true-16K passkey cases, median target-block hit@16 over 32 retrieval heads rises from 18.75% to 64.06%. Blocking answer-side decode attention to the remote gold block in every head preserves all prompt tokens, original rotary KV indices, and the 128-token block geometry. NLL pools the negative log probabilities of the 30 answer tokens (3 per case); deletion changes this mean by  $-0.0095$  for Native and  $+1.5055$  for EVQ-COSH. The EVQ-COSH change is positive in all ten cases. Case-level hit@16 ranges are 3.12–34.38% for Native and 62.50–68.75% for EVQ-COSH. These are matched teacher-forced probability and position-preserving causal-routing measurements; autoregressive task accuracy is evaluated separately. Descriptive Native/EVQ-COSH PPL is 6.817/10.068, 108.958/24.068, and 991.475/127.911 at those lengths.

**RULER-family continuation.** A separate seed-20420726 continuation uses 516 steps, rank-64/alpha-128 Q/K/V/O LoRA, global batch 8 (micro batch 2, accumulation 4), BF16 and Flash-only attention. Its 1,376 physical-8K training rows comprise 96 rows for each of 13 RULER families plus 128 natural-replay rows; it processes 33,816,576 input tokens. Training and evaluation rows have

zero exact overlap, and evaluation uses 20 examples per family and length. At 16K, Native/EVQ-COSH official-scorer macro is 0.295/14.03% and normalized exact is 0/1.54%. At 8K, Table 16 reports official scorer macro (94.44/77.60%) and normalized exact (17.69/21.54%).

## F FROZEN EXPONENT-ADJUSTMENT PROTOCOLS

This section specifies the static profiles and evaluation panels used in §6. An amplitude  $g$  multiplies both the cosine and sine tables, so its contribution to a rotated Q/K logit is  $g^2$ . Every table is installed before prefill and retained for the entire request. The accompanying `figs/recorded_runtime_identities.json` records checkpoint, tokenizer, and evaluator identities. The index/placement studies use RULER commit `c3f5e3b4f87f97e048793bb510a3a6b19a46bf3a`.

### F.1 NORMALIZED-INDEX PROFILE

For a native table with  $K$  pairs, base  $b$ , and reference length  $L_{\text{ref}}$ , the normalized-index construction evaluates the profile on a 64-pair reference grid with the same  $b, L_{\text{ref}}$ :

$$\bar{\omega}_j = b^{-j/64}, \quad c_{64} = (1 - b^{-1/64})^{-1}, \quad \xi_j = \log \frac{L_{\text{ref}} \bar{\omega}_j}{2\pi c_{64}}, \quad (74)$$

$$\bar{m}_j = \text{clip} \left( \frac{\xi_H - \xi_j}{\xi_H - \xi_L}, 0, 1 \right), \quad \xi_H = 0.7382780681078285, \quad \xi_L = 0.366403835112904. \quad (75)$$

These two constants were fitted to the existing OLMo displacement profile and then frozen. Linear interpolation of  $(j/63, \bar{m}_j)$  at  $k/(K-1)$  gives  $m_k$  on the target grid. The installed table is  $\omega'_k = \omega_k^N s^{-m_k}$ .

The Qwen-2.5-0.5B-Instruct experiment uses its native  $K = 32$  pairs,  $b = 10^6$ , and original context length  $L_{\text{ref}} = 32768$ , with  $s = 2$  and fixed amplitude  $g = 1 + 0.074 \log 2 = 1.0512928914$ . The YaRN-2 comparison uses the Transformers official-equation table and  $g = 1 + 0.1 \log 2 = 1.0693147181$ ; Native uses  $g = 1$ .

The full RULER-13 confirmation uses seed 202609027, 20 examples per task and length, the same tokenizer and native chat template, and greedy task-budgeted generation. It contains 520 generations per arm across 32768 and 65536 tokens. Per-task scores are shown in Table 22. The macro averages the thirteen task means equally. Paired intervals use 10,000 bootstrap samples within task strata with seed 202609028. The 64K index-minus-YaRN difference is 6.0897 percentage points, with interval  $[2.7627, 9.5835]$ ; at 32K it is  $-0.0256$  points, with interval  $[-3.2630, 3.2179]$ .

Table 22: **Qwen-0.5B RULER-13 per-task score (%)**, 20 examples per cell (260 prompts per length). **X1 Context Note:** The comparison between Normalized-Index and YaRN represents a joint *table–amplitude configuration*, where Index couples with gain  $g_L = 1 + 0.074 \ln 2 \approx 1.0513$  and YaRN couples with  $g_H = 1 + 0.1 \ln 2 \approx 1.0693$ . The +6.09 pp advantage at 64K is achieved by this joint configuration; disentangling frequency allocation from amplitude scaling requires the full  $2 \times 2$  factorial grid (see Part A of the Qwen ablation protocol).

| Task                 | 32K          |                 |                | 64K          |                 |                |
|----------------------|--------------|-----------------|----------------|--------------|-----------------|----------------|
|                      | Native       | Index ( $g_L$ ) | YaRN ( $g_H$ ) | Native       | Index ( $g_L$ ) | YaRN ( $g_H$ ) |
| Single-1             | 100.00       | 100.00          | 100.00         | 75.00        | 100.00          | 100.00         |
| Single-2             | 100.00       | 100.00          | 100.00         | 15.00        | 100.00          | 75.00          |
| Single-3             | 100.00       | 100.00          | 100.00         | 30.00        | 100.00          | 95.00          |
| MultiKey-1           | 90.00        | 95.00           | 90.00          | 50.00        | 90.00           | 65.00          |
| MultiKey-2           | 50.00        | 60.00           | 35.00          | 0.00         | 40.00           | 25.00          |
| MultiKey-3           | 5.00         | 0.00            | 10.00          | 0.00         | 5.00            | 5.00           |
| MultiValue           | 62.50        | 81.25           | 81.25          | 22.50        | 63.75           | 58.75          |
| MultiQuery           | 62.50        | 80.00           | 77.50          | 22.50        | 65.00           | 50.00          |
| Variable tracking    | 48.00        | 42.00           | 46.00          | 3.00         | 28.00           | 36.00          |
| Common words         | 2.50         | 2.00            | 2.50           | 7.00         | 0.50            | 0.00           |
| Frequent words       | 46.67        | 31.67           | 35.00          | 31.67        | 46.67           | 40.00          |
| QA-1                 | 20.00        | 10.00           | 30.00          | 5.00         | 15.00           | 20.00          |
| QA-2                 | 25.00        | 25.00           | 20.00          | 25.00        | 15.00           | 20.00          |
| <b>Macro Average</b> | <b>54.78</b> | <b>55.92</b>    | <b>55.94</b>   | <b>22.05</b> | <b>51.46</b>    | <b>45.37</b>   |

## F.2 PLACEMENT ON ANOTHER FINITE PAIR GRID

The alternative direct-gap construction evaluates the same clipped profile using the target table’s local frequency spacing:

$$c_K = (1 - b^{-1/K})^{-1}, \quad \xi_k^{(K)} = \log \frac{L_{\text{ref}} \omega_k^N}{2\pi c_K}, \quad m_k^{\text{gap}} = \text{clip} \left( \frac{\xi_H - \xi_k^{(K)}}{\xi_H - \xi_L}, 0, 1 \right).$$

Both placements share the frozen  $\xi_H, \xi_L$  and amplitude coefficient 0.074. Within each comparison, endpoints, scale, amplitude, checkpoint, inputs, decoder, and scorer are matched. Table 23 reports two independent-input confirmations with 80 rows per task.

Qwen-0.5B uses the 32-pair,  $s = 2$  setting above and data seed 202609026, with four tasks at each of 32K and 64K. These are separate inputs from the subsequent RULER-13 confirmation. Gemma-1.1-2B-Instruct uses  $K = 128$ ,  $b = 10,000$ , a fixed operating reference of 4096,  $s = 4$ , and  $g = 1 + 0.074 \log 4 = 1.1025857827$ . Its confirmation uses data seed 202609028 at 16K. The task vector is NIAH single-1, multikey-2, multikey-3, and variable tracking. Gemma’s direct-gap scores are 100/88.75/20/82.50%; normalized-index scores are 100/88.75/40/87.25%. Both tables were frozen before these inputs were generated, and the earlier selection rows are excluded. This confirmation compares the two placements; the Native/YaRN baseline panel uses a different, earlier input set.

Table 23: Independent-input placement comparisons, 80 examples per task and four tasks per panel. Intervals are for index minus direct-gap, in percentage points: 97.5% for each Qwen length (Bonferroni over two lengths), 95% for Gemma.

| Model     | Length | Direct-gap (%) | Index (%) | Difference interval |
|-----------|--------|----------------|-----------|---------------------|
| Qwen-0.5B | 32K    | 49.69          | 53.37     | $[-0.25, +7.75]$    |
| Qwen-0.5B | 64K    | 46.62          | 46.13     | $[-4.87, +3.81]$    |
| Gemma-2B  | 16K    | 72.81          | 79.00     | $[+2.81, +9.62]$    |

### F.3 BOUNDARY-MATCHED INTERMEDIATE EXPONENTS

Let  $l$  be the last native frequency pair making more than 32 turns over the reference window and  $h$  the first making less than one turn. Write  $N = h - l$  and  $q = \text{clip}(k - l, 0, N)$ . All profiles leave the fast band unchanged and divide the slow-band frequencies by  $s$ :

$$\omega'_k = \omega_k^N s^{-m_q}, \quad m_0 = 0, \quad m_N = 1. \quad (76)$$

With the same band boundaries, the three intermediate allocations are

$$m_q^{\text{Uni}} = \frac{q}{N}, \quad (77)$$

$$m_q^{\text{Pro}} = \frac{q(q+1)}{N(N+1)}, \quad (78)$$

$$m_q^{\text{BM}} = \frac{q(q+1)(3N+2-2q)}{N(N+1)(N+2)}. \quad (79)$$

The Uni and Pro rules follow the uniform and progressive mixed-radix constructions of Tian et al. (2026). BM smooths the radix increments  $\epsilon_q = m_q - m_{q-1}$  at both band boundaries.

To obtain BM, minimize  $\sum_{q=0}^N (\epsilon_{q+1} - \epsilon_q)^2$  with  $\epsilon_0 = \epsilon_{N+1} = 0$  and  $\sum_{q=1}^N \epsilon_q = 1$ . The discrete Euler equation gives

$$\epsilon_q = \frac{6q(N+1-q)}{N(N+1)(N+2)};$$

summing gives (79). The quadratic form is strictly convex with the endpoints fixed, so this solution is unique. Figure 6 shows the actual profile shapes for the tested transition widths.

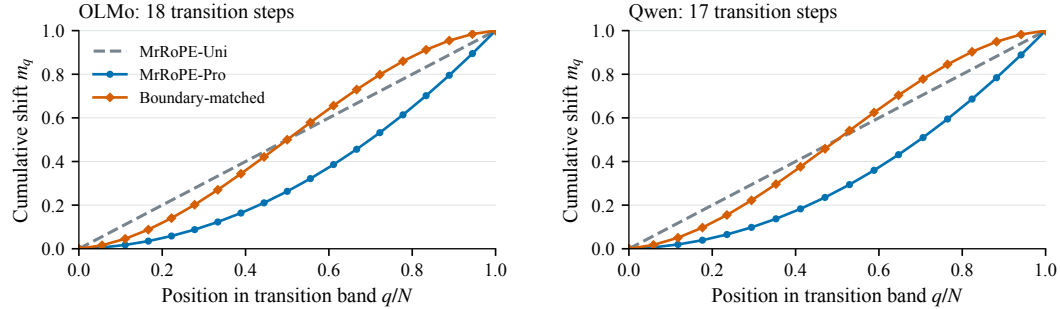


Figure 6: **Intermediate-band exponent displacements.** All three profiles have the same cumulative scale at the endpoints. BM redistributes the increments toward the middle of the transition band.

OLMo-2-0425-1B-Instruct uses  $b = 500,000$ , reference length 4096,  $K = 64$ , and  $(l, h, N) = (14, 32, 18)$ . Qwen-2.5-3B/7B-Instruct use  $b = 10^6$ , reference length 32768,  $K = 64$ , and  $(l, h, N) = (23, 40, 17)$ . The factor-four comparisons use  $g = 1 + 0.1 \log 4 = 1.1386294361$  for both MrRoPE-Pro and BM, with FP32 frequency tables and BF16 model execution.

### F.4 SIX-TASK MODEL COMPARISONS

The shared task types are NIAH single-2, multikey-2, multiquery, variable tracking, frequent words, and QA-1. OLMo’s confirmation uses seed 20260910 with four short and eight long examples per task, after a separate development panel with seed 20260909. Its 4K/16K scores are listed in Table 27; MrRoPE-Uni on the same confirmation inputs scores 76.88/32.12%, and official YaRN scores 54.38/6.94%. The Qwen-3B panel has two 32K and four 128K examples per task. The 7B screen reuses the first one short and two long examples per task; this selection was fixed before its outputs were inspected. Table 24 reports every task mean.

All arms use the corresponding native chat template and the same task-specific generation caps, score the entire saved response with the RULER task metric, and retain generated token IDs and EOS flags. On Qwen-7B, a tokenwise MLP chunk size of 4096 makes 128K evaluation fit in memory; both arms share that path while retaining full attention.

Table 24: All six-task BM/MrRoPE-Pro comparisons. Values are official RULER task scores (%); model-specific sample counts are given in Table 27.

| Model   | Task              | Short  |        | Long   |        |
|---------|-------------------|--------|--------|--------|--------|
|         |                   | MrPro  | BM     | MrPro  | BM     |
| OLMo-1B | Single-2          | 75.00  | 100.00 | 0.00   | 100.00 |
|         | MultiKey-2        | 25.00  | 100.00 | 0.00   | 37.50  |
|         | MultiQuery        | 18.75  | 87.50  | 0.00   | 43.75  |
|         | Variable tracking | 25.00  | 70.00  | 0.00   | 10.00  |
|         | Frequent words    | 58.33  | 58.33  | 16.67  | 41.67  |
|         | QA-1              | 25.00  | 75.00  | 0.00   | 75.00  |
| Qwen-3B | Single-2          | 100.00 | 100.00 | 100.00 | 100.00 |
|         | MultiKey-2        | 100.00 | 100.00 | 75.00  | 50.00  |
|         | MultiQuery        | 100.00 | 100.00 | 93.75  | 100.00 |
|         | Variable tracking | 90.00  | 100.00 | 75.00  | 75.00  |
|         | Frequent words    | 83.33  | 100.00 | 75.00  | 75.00  |
|         | QA-1              | 50.00  | 50.00  | 50.00  | 25.00  |
| Qwen-7B | Single-2          | 100.00 | 100.00 | 100.00 | 100.00 |
|         | MultiKey-2        | 100.00 | 100.00 | 100.00 | 100.00 |
|         | MultiQuery        | 100.00 | 100.00 | 100.00 | 100.00 |
|         | Variable tracking | 100.00 | 80.00  | 90.00  | 60.00  |
|         | Frequent words    | 100.00 | 100.00 | 66.67  | 66.67  |
|         | QA-1              | 0.00   | 0.00   | 50.00  | 0.00   |

## F.5 NATURAL QUESTION ANSWERING

The natural comparison reuses frozen, source-verified LongBench prompts (Bai et al., 2024). Inputs are untruncated, retain their source questions and references, and fit the 16K input-plus-generation budget. Selection follows the pre-existing source-hash order; no generated answer participates in selection. The complete eligible pool contains 778 questions, of which 631 have more than 4096 input tokens. Both methods retain the same factor-four table at every length.

Greedy generation uses a 32-token cap for HotpotQA and 2WikiMQA, 128 for Qasper and NarrativeQA, and 64 for MultiFieldQA-en. The metric is token F1 on the complete saved generated response against the task references. Table 25 gives the long-input and original-window-length strata, including EOS counts. The long-input macro averages the five task means; the shorter-input macro averages its four available tasks. The respective BM-minus-MrPro changes are +3.8194 and +2.0028 points, with paired within-task bootstrap intervals [1.3161, 6.2920] and [-3.5618, 7.4742]. The MultiFieldQA long-input row includes 40 capped responses in each arm; these remain in the F1 calculation.

## G SUPPORTING RESULTS

### G.1 MULTI-HEAD LATENT ATTENTION

We report allocation at one fixed rotary budget in a 432M-parameter transformer following the MLA architectural pattern (DeepSeek-AI, 2024a), with 24 layers, width 1024, 16 heads, and per-head width 64. Its decoupled rotary subspace has  $d_{\text{rope}}=32$  ( $K=16$  frequencies); the remaining  $d_{\text{nope}}=32$  dimensions are non-rotary, and the shared KV latent rank is 256. The model trains from scratch on 500M FineWeb-Edu tokens (Penedo et al., 2024) at  $L_{\text{train}}=8192$  and base 500K over seeds 42, 43, and 88. EVQ-COSH uses the pre-specified  $\tau = 1.414$  setting. At 16K, Geo PPL is 141.1/132.5/142.8 for seeds 42/43/88, whereas EVQ-COSH is 93.7/92.7/100.3; all three paired directions agree.

For the supporting overlay, let  $\lambda_k = 2\pi/\omega_k$  and  $t_k = \text{clip}((\lambda_k - L_{\text{train}})/(31L_{\text{train}}), 0, 1)$ . The MLA wavelength-blend operator sets  $\omega'_k = \omega_k/[1 + t_k(s - 1)]$ : it preserves wavelengths below the training length, applies the full factor  $s$  above  $32L_{\text{train}}$ , and interpolates between them.

Table 25: **OLMo-2 Natural QA under static extension** ( $s = 4$ ,  $4K \rightarrow 16K$ ,  $g = 1.1386$ ). Evaluation on allenai/OLMo-2-0425-1B-Instruct across 778 frozen LongBench inputs (greedy decoding, whole response token F1). **X2 Baseline Scope:** Boundary Matching (BM) is compared directly against MrRoPE-Pro on this exact panel, improving extended macro F1 by +3.82 pp (95% bootstrap CI [+1.32, +6.29]) and winning all five tasks; comparison against official YaRN is registered as a planned baseline.

| Context Stratum         | Task                    | $n$        | MrPro F1 (%) | BM F1 (%)    | $\Delta$ (pp) | MrPro EOS  | BM EOS     |
|-------------------------|-------------------------|------------|--------------|--------------|---------------|------------|------------|
| Extended (4K–16K)       | HotpotQA                | 166        | 28.63        | <b>33.54</b> | <b>+4.91</b>  | 160        | 163        |
|                         | 2WikiMQA                | 173        | 21.55        | <b>25.43</b> | <b>+3.88</b>  | 169        | 171        |
|                         | Qasper                  | 119        | 13.94        | <b>19.06</b> | <b>+5.11</b>  | 119        | 119        |
|                         | NarrativeQA             | 61         | 9.51         | <b>11.48</b> | <b>+1.96</b>  | 61         | 61         |
|                         | MultiFieldQA            | 112        | 34.49        | <b>37.72</b> | <b>+3.23</b>  | 72         | 72         |
|                         | <b>Task-Equal Macro</b> | <b>631</b> | 21.62        | <b>25.44</b> | <b>+3.82</b>  | <b>581</b> | <b>586</b> |
| In-Window ( $\leq 4K$ ) | HotpotQA                | 5          | <b>58.33</b> | 51.67        | −6.67         | 5          | 5          |
|                         | 2WikiMQA                | 26         | 34.62        | <b>44.98</b> | <b>+10.37</b> | 26         | 26         |
|                         | Qasper                  | 78         | 26.16        | <b>28.14</b> | <b>+1.98</b>  | 78         | 78         |
|                         | MultiFieldQA            | 38         | 50.94        | <b>53.27</b> | <b>+2.34</b>  | 35         | 33         |
|                         | <b>Task-Equal Macro</b> | <b>147</b> | 42.51        | <b>44.51</b> | <b>+2.00</b>  | <b>144</b> | <b>142</b> |

Table 26: MLA validation (432M,  $d_{\text{rope}}=32$ , 3-seed mean  $\pm$  std). Rows marked “MLA wavelength-blend” use this run’s wavelength-blend operator defined above.

| Method                                                      | PPL@8K         | PPL@16K                        | PPL@24K                          | PPL@32K                          |
|-------------------------------------------------------------|----------------|--------------------------------|----------------------------------|----------------------------------|
| Geo                                                         | 35.4 $\pm$ 0.9 | 138.8 $\pm$ 5.5                | 241.5 $\pm$ 2.6                  | 323.7 $\pm$ 5.2                  |
| Geo + MLA wavelength-blend ( $s = 4$ )                      | 35.5 $\pm$ 0.9 | 117.9 $\pm$ 6.5                | 204.9 $\pm$ 6.9                  | 278.5 $\pm$ 6.3                  |
| <b>EVQ-COSH</b>                                             | 35.8 $\pm$ 0.8 | <b>95.6<math>\pm</math>4.1</b> | 204.9 $\pm$ 14.7                 | 291.6 $\pm$ 20.5                 |
| <b>EVQ-COSH + MLA wavelength-blend (<math>s = 4</math>)</b> | 35.8 $\pm$ 0.8 | <b>71.1<math>\pm</math>4.1</b> | <b>153.2<math>\pm</math>12.5</b> | <b>236.6<math>\pm</math>15.9</b> |

EVQ-COSH reduces 16K PPL by 31.1% ( $138.8 \rightarrow 95.6$ ), while 8K PPL changes by +0.9%. All three training seeds favour EVQ-COSH at 16K. The EVQ-COSH plus MLA wavelength-blend row has the lowest mean PPL at 16K and beyond, showing that the allocation also composes with this wavelength-dependent extension.

## G.2 FIGURES AND TABLES RELOCATED FROM THE MAIN TEXT

The following floats support statements made in the main text; each is summarised there with its numbers, and is reproduced here in full.

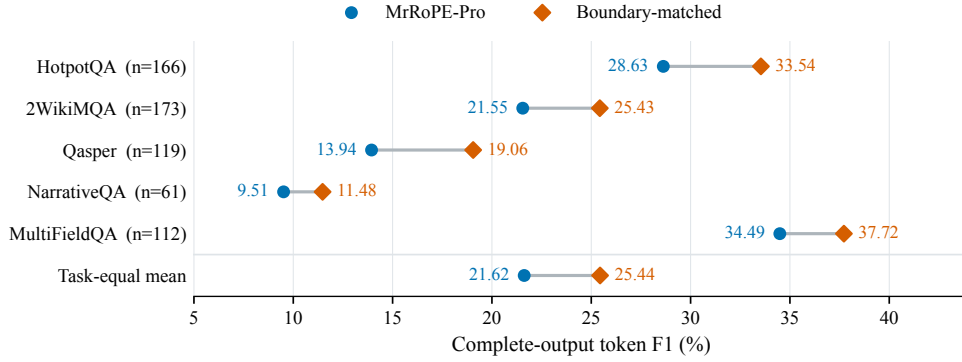


Figure 7: **Intermediate-band redistribution improves OLMo natural QA.** Complete-output token F1 on untruncated long inputs ( $s = 4$ ); sample counts shown by task, final row averages task means.

Table 27: **Matched six-task RULER comparisons (%)**. Each column averages six tasks ( $s = 4$ ; short/long prompt counts per arm).

| Model   | Lengths | Rows  | Short |       | Long  |       |
|---------|---------|-------|-------|-------|-------|-------|
|         |         |       | MrPro | BM    | MrPro | BM    |
| OLMo-1B | 4/16K   | 24/48 | 37.85 | 81.81 | 2.78  | 51.32 |
| Qwen-3B | 32/128K | 12/24 | 87.22 | 91.67 | 78.13 | 70.83 |
| Qwen-7B | 32/128K | 6/12  | 83.33 | 80.00 | 84.44 | 71.11 |

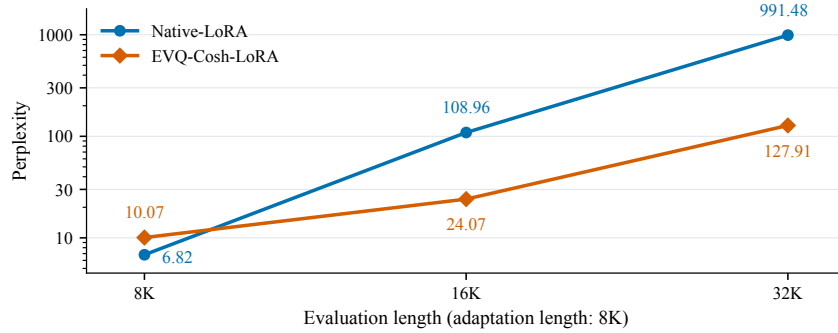


Figure 8: **Cosh allocation improves extrapolation under matched adaptation.** PPL is the exponential of equal-domain mean NLL.

Table 28: **Frozen-model RULER scores (%) at matched exponent range.** OLMo: held-out nine-task confirmation at 16K. Qwen: four-task development panel at 64K, 20 rows/task.

| Exponent allocation | OLMo-1B      | Qwen-1.5B    |
|---------------------|--------------|--------------|
| Geometric           | 0.56         | 57.75        |
| Coarse ramp         | <b>61.04</b> | 64.00        |
| Residual-guided     | 60.47        | <b>66.50</b> |

Table 29: **Static extension across native and extended windows.** Qwen-0.5B RULER-13 macro (%);  $g$  is fixed amplitude.

| Installed table    | $g$    | 32K   | 64K          |
|--------------------|--------|-------|--------------|
| Native             | 1.0000 | 54.78 | 22.05        |
| YaRN-2             | 1.0693 | 55.94 | 45.37        |
| Normalized-index-2 | 1.0513 | 55.92 | <b>51.46</b> |

Table 30: **Epistemic status of theoretical and empirical components.** Differentiating exact basis geometry, surrogate design criteria, heuristic operating rules, and downstream behavior.

| Tier | Component            | Object                                                        | Epistemic Status                                                |
|------|----------------------|---------------------------------------------------------------|-----------------------------------------------------------------|
| 1    | Positional Basis     | Subspace overlap $c_{\omega\psi}$ , rank $r_2$ , limit $V_0$  | Exact operator geometry (Prop. 1)                               |
| 2    | Surrogate Functional | Variational objective $\mathcal{C}_{\text{app}}[\rho]$ in (6) | Explicit design prior; not downstream $\mathcal{L}_{\text{CE}}$ |
| 3    | Operating Rule       | Cosh density $\rho_\tau$ and finite quantile grid             | Closed-form heuristic family; fixed $K$ grid                    |
| 4    | Downstream LM        | Tail NLL, RULER-13, natural LongBench QA                      | Not predicted by static surrogates (Thm. 2)                     |




Figure 9: **Weights co-adapt to their exponent allocation.** Each row freezes one trained weight set and changes only its runtime table. Left: 50M crossing, PPL on eight windows. Right: 151.9M crossing, mean tail NLL on 32 anchors in each of two training seeds, using extensions derived from the two training tables. FMR is the geometric table at the fixed training range; shading shows the NLL increase relative to the matched diagonal cell.