

BEYOND THE BASE: FREQUENCY ALLOCATION IN ROPE

Anonymous authors

Paper under double-blind review

ABSTRACT

Rotary position embeddings (RoPE) typically tie their frequency placement to a single base. We study frequency allocation as a design dimension for quality within and beyond the native context window. Controlled interventions show that internal placement matters even when the frequency range and total log-frequency displacement are fixed. Full sine–cosine geometry characterizes positional overlap, while coordinate interventions distinguish this structure from the model’s learned use of frequencies. Allocation also determines how each channel stretches its distance response. This view motivates TailSpline, a closed-form transition that smoothly joins the extended low-frequency tail. It retains standard rotary computation without weight updates or calibration. Matched comparisons show improvements in retrieval and language modeling across model families and scales, together with gains on natural question answering. Supporting studies extend this design freedom to native-window language modeling and learned extrapolation. These findings establish internal frequency placement as a practical means of improving how models use context.

1 INTRODUCTION

As native context windows grow, the largest extension ratio is only one aspect of positional-encoding quality. For a model with native length L deployed up to sL , we seek useful task quality across the intended range and ask whether performance within L can improve over its native encoding. We study frequency allocation z : the placement of rotary frequencies within a chosen log-frequency range.

RoPE’s base selects a geometric table. YaRN and MrRoPE redistribute scaling for extension, while LeRoPE learns frequencies (Peng et al., 2024; Tian et al., 2026; Karypis et al., 2026). What does internal placement contribute beyond choosing the range? Matched interventions show that it can improve quality even at fixed endpoints and equal total log-frequency displacement. Its value depends on the chosen range and the model’s learned use of frequencies.

Allocation changes each learned channel’s distance response. A longer input can move a query farther from a relevant record while preserving local relations within it. TailSpline distributes this distance stretch through a closed-form transition to the interpolated low-frequency tail. A phase bound explains its approach to that reference; matched controls and deployment comparisons test its practical value. Native-window interventions and Cosh transport provide complementary tests of allocation within the original window and during learning.

Our contributions are:

- **Allocation matters beyond frequency range.** Fixed-support and equal-displacement controls identify placement effects; changing the range can reverse the preferred allocation.
- **Learned frequency use affects allocation quality.** Complete sine–cosine pairs characterize positional overlap; reassigning the same spectrum across learned coordinates changes task performance.
- **Explicit allocations improve practical quality.** TailSpline improves long-context quality across model families and scales. Supporting native-window and learning studies establish further uses of allocation.

2 ALLOCATION AS A DESIGN VARIABLE

For $K = d_{\text{rot}}/2$ pairs, standard RoPE uses $\omega_k = b^{-k/K}$. Any strictly ordered positive table admits

$$x_k = -\log \omega_k = a + Rz_k, \quad z_k = \frac{x_k - x_0}{x_{K-1} - x_0}, \quad z_0 = 0, \quad z_{K-1} = 1. \quad (1)$$

The endpoints (a, R) fix the sampled log-frequency interval, called the support; $K - 2$ interior coordinates define allocation. Every geometric table has $z_k = k/(K - 1)$ irrespective of base. The full space $\mathcal{Z}_K = \{0 = z_0 < \dots < z_{K-1} = 1\}$ includes analytic subfamilies and positive-gap learned parameterizations.

The endpoints determine the covered range; z determines how the rotary pairs are placed inside it. We fix endpoints in controlled comparisons to isolate the effect of interior placement. For deployment relative to a native table, write

$$\omega'_k = \omega_k^N s^{-m_k}, \quad d_k = \log(\omega_k^N / \omega'_k) = m_k \log s. \quad (2)$$

For $m_0 = 0, m_{K-1} = 1$, the corresponding allocation is $z'_k = (R_N z_k^N + m_k \log s) / (R_N + \log s)$, where R_N is the native log-frequency span. Thus m describes the same frequency placement relative to the native grid. Changing frequencies preserves RoPE’s relative-position composition and norm preservation (Su et al., 2024). Multiplying cosine/sine amplitudes by g separately scales the rotary Q/K contribution by g^2 . Fixed-support experiments isolate z ; complete methods may jointly choose allocation, range and amplitude.

Geo denotes each protocol’s geometric reference and Native a released model’s inherited table. Standard and midpoint geometric grids share z but differ by an overall frequency offset. Exact grid and endpoint conventions are in Appendix B.3.

Separating the effects of allocation. At fixed support, total displacement $D = \sum_k d_k$ is a statistic of z . Matching D resolves finer allocation effects; permuting the same frequency multiset tests coordinate assignment. Appendix C specifies these controls.

3 FREQUENCY ALLOCATION MATTERS AT FIXED RANGE

We first isolate the effect of interior placement by fixing the frequency endpoints. We then examine how that effect depends on training, total frequency displacement, and the model’s learned use of its coordinates.




Figure 1: **Frequency range leaves consequential allocation choices open.** (a) Different interior placements share the same endpoints ($K = 8$ illustration). (b) Frozen OLMo/Qwen models attain different task scores within a fixed range. (c) In paired training, each weight set favors the runtime table derived from its own allocation: two-seed mean tail NLL at 1K.

3.1 DIFFERENT ALLOCATIONS WITH THE SAME FROZEN WEIGHTS

On OLMo-2-1B and Qwen-2.5-1.5B, we fix the checkpoint, sampled endpoints, 64 pairs, attention amplitude, input rows, and decoder. The range extends from the native fast frequency to one-quarter of the native slow frequency. On OLMo’s nine held-out RULER tasks (Hsieh et al., 2024), uniform

normalized exponents score 0.56%, a coarse ramp 61.04%, and a residual-guided profile 60.47%. Qwen’s four-task comparison gives 57.75/64.00/66.50% (Appendix C.5). Both reallocations yield large OLMo gains while retaining the same sampled range (Fig. 1b).

Internal shape beyond total displacement. A frozen OLMo comparison also matches total displacement $\sum_k d_k$, alongside endpoints, band and gain: the symmetric profile BM and uniform transition MrRoPE-Uni score 51.32 and 32.12% in a six-task, 48-input 16K follow-up comparison. Total displacement therefore does not explain this score difference (Appendix C.5). The clean TailSpline-C comparison confirms a shape effect at matched displacement on Llama (Section 6.1).

3.2 A PAIRED INTERVENTION AT FIXED SUPPORT

A paired 151.9M Cosh training experiment fixes both frequency endpoints to isolate interior placement, using $K = 32$ rotary pairs and training length 256.

Within each of three seeds, architecture, initialization, token order, optimizer, schedule, and the approximately 500M-token budget are identical. The matched geometric control (Geo) uses $b = L_{\text{train}}$; the EVQ-Cosh intervention anchors its quantiles with $\tau = 4$ to the same highest and lowest frequencies and log-span (§5.2). Only the 30 interior z coordinates differ.

With the training support retained at evaluation, mean Cosh-minus-Geo tail NLL is +0.026 at the training length and $-0.281 / -0.176 / -0.146$ at 512/1024/2048 tokens. All three paired seeds improve at every extended length (Table 5).

Retargeting both supports to $b = L_{\text{eval}}$ while retaining each trained checkpoint and normalized allocation reverses the ordering: Cosh-minus-Geo NLL becomes $+0.060 / +0.227 / +0.460$, with all three seeds favoring Geo (Table 5). Allocation is a useful design variable, and its preferred setting depends on the support policy.

The effect extends beyond one allocation family. A 50.9M short-training factorial (8.39M tokens per arm) with three paired seeds tests two bases, two training lengths and three head dimensions. Reference Cosh, its preassigned $1.25\times$ strength, and an RMS-movement-matched exponential improve OOD NLL in 7/12, 10/12, and 9/12 configurations. Cosh and the matched exponential perform similarly (mean difference +0.00074 NLL, interval $[-0.006, 0.008]$). These results support nonuniform allocation across families, with strength-dependent trade-offs; full settings and a separate standard-grid control are in Appendix C.4.

A learned model uses more than a frequency range. Crossing the runtime tables of the paired 151.9M models tests whether weights can use a different allocation interchangeably. Both seeds favor the table derived from their training allocation (Fig. 1c; Appendix C.3). A complementary frozen intervention permutes the same frequency multiset across coordinates and changes model performance (Appendix C.6). These findings connect the two operating settings: learning can adapt weights to an allocation, whereas frozen extension must work with the model’s existing frequency-coordinate association.

4 POSITIONAL STRUCTURE AND CONTENT USE

The controlled gains raise a structural question: what changes when the frequency range stays fixed? Complete rotary pairs supply positional directions whose finite-window overlap connects frequency spacing to redundancy. The model learns how to combine these directions with content.

4.1 COMPLETE ROTARY PAIRS AND POSITIONAL OVERLAP

For a fixed content pair, one rotary channel contributes $C \cos(\omega\Delta) + D \sin(\omega\Delta)$ to an attention logit, where Δ is relative token separation. Both coefficients depend on content; the positional object is

$$V_\omega = \text{span}\{\cos(\omega\Delta), \sin(\omega\Delta)\}. \quad (3)$$

To avoid privileging a content phase, let $x_\omega(\Delta) = [\cos(\omega\Delta), \sin(\omega\Delta)]$ and define $S_\omega = \mathbb{E}[x_\omega^\top x_\omega]$, $H_{\omega\nu} = \mathbb{E}[x_\omega^\top x_\nu]$ under a declared separation measure. For nonsingular self-Grams, the singular

values of $Q_{\omega\nu} = S_{\omega}^{-1/2} H_{\omega\nu} S_{\nu}^{-1/2}$ are the canonical correlations, or cosines of principal angles, of the two spaces (Björck & Golub, 1973). Their mean square,

$$c_{\omega\nu} = \frac{1}{2} \|Q_{\omega\nu}\|_F^2 \in [0, 1], \quad (4)$$

measures shared positional directions and is invariant to invertible basis changes within either pair. Let Γ be the Gram matrix of the concatenated block-whitened pairs, with diagonal blocks I_2 . The average pairwise overlap $\bar{c}$ gives the exact Rényi-2 effective-rank identity

$$r_2(\Gamma) = \frac{(\text{tr } \Gamma)^2}{\text{tr}(\Gamma^2)} = \frac{2K}{1 + (K-1)\bar{c}}. \quad (5)$$

Here $\text{tr}(\Gamma^2)$ is the frame potential of the normalized positional functions (Benedetto & Fickus, 2003). The identity specializes this classical measure to full rotary pairs: increasing pairwise overlap concentrates the same $2K$ coordinates into fewer effective positional directions. These are normalized directions in the supplied frequency basis; content-dependent coefficients determine their contribution to attention logits.

4.2 FINITE-WINDOW CONCENTRATION OF SLOW DIRECTIONS

Proposition 1 (Shared slow-frequency subspace). *Under uniform measure on $[0, L]$, with $t = \Delta/L$, as $x = \omega L \rightarrow 0$, $\cos(xt) \rightarrow 1$ and $\sin(xt)/x \rightarrow t$, so V_{ω} converges to $\text{span}\{1, \Delta\}$. For $\epsilon = \max\{|\omega L|, |\nu L|\} \rightarrow 0$, squared chordal separation satisfies $2 - \|Q_{\omega\nu}\|_F^2 = \mathcal{O}(\epsilon^4)$.*

Distinct slow frequencies therefore approach a shared two-dimensional positional space. On the identification grid $b = 256$, $K = 32$, $L = 256$, the eight slowest pairs span $\omega L \in [1.19, 4]$ and have effective rank $r_2 = 2.11$. Their overlap is substantial even beyond the small-phase limit (Fig. 2; Appendix A.2).




Figure 2: **Allocation determines the supplied positional directions.** Complete-pair overlap on the identification grid. The eight slowest pairs occupy 16 coordinates but have normalized positional effective rank 2.11 under uniform separations; the outlined block marks these pairs.

Why use complete rotary pairs? A cosine-only overlap proxy can reverse the allocation ordering: explicit four-pair tables with identical endpoints give $C_{\cos}(A) = 0 < C_{\cos}(B)$, yet $r_2(A) = 6.29 < 8.00 \approx r_2(B)$ because sine-cosine cross terms also carry positional overlap (Appendix A.3). The complete-pair measure captures this structure independently of phase.

4.3 SHARED POSITIONAL DEPENDENCE AND CONTENT USE

Each rotary block remains invertible even when its positional functions overlap. Shared dependence on distance can therefore coexist with distinct content comparisons. A slow-block construction makes this distinction explicit (Appendix A.5), complementing prior observations of positional and semantic frequency use (Barbero et al., 2025).

The complete frozen Llama tables provide a concrete example. At 16/32K, TailSpline has full-pair effective ranks 8.28/10.08, compared with MrPro’s 8.74/10.21 under uniform separations, while improving clean RULER by 3.39/11.72 points (Table 1). These gains establish the value of controlling distance dependence through allocation even when normalized positional diversity decreases. The model’s use of the supplied frequencies determines their practical value (Appendix A.6). Concretely, changing a frequency changes the distance response multiplying its content coefficients. A useful response improves target evidence relative to competing keys: a larger target logit alone does not determine its attention share. Appendix A.6 makes this comparison exact for fixed pre-RoPE states, linking positional structure to content competition.

Frequency changes and coordinate assignment. Changing frequencies also differs from reassigning a fixed spectrum. In the no-alias interval $(0, \pi)$, position-independent invertible Q/K maps preserve the full integer-position rotary kernel exactly if and only if the frequency multisets agree (Appendix A.4). This criterion complements the coordinate interventions and STRING’s structural characterization (Schenck et al., 2025). These distinctions motivate treating the native table as a design reference for changing the distance responses of learned content channels.

5 EXPLICIT FREQUENCY ALLOCATIONS

We derive TailSpline for frozen extension, then give a compact frequency transport for studying allocation during learning.




Figure 3: **TailSpline connects the transition to the extended tail.** Extra log-frequency gaps for $n = 17$. Both profiles retain the same outer bands; TailSpline allocates a larger entry gap and smoothly approaches the zero extra gap of the fully interpolated tail.

5.1 TAILSPINE: ALLOCATING EXTENSION AROUND A NATIVE TABLE

A frequency $\omega'_k = \omega_k^N s^{-m_k}$ evaluates the native rotary response at effective distance Δ/s^{m_k} . Hence $m_k = 0$ preserves native-distance responses, whereas $m_k = 1$ restores them at s times the distance. Longer inputs can contain both unchanged local dependencies and stretched query–evidence separations, motivating different amounts of scaling across channels (Wu et al., 2026).

We retain the high-frequency band and divide the low-frequency tail by s , then design the intervening allocation. Following MrRoPE’s 32/1-turn boundaries (Tian et al., 2026), let l be the last pair above 32 native turns and h the first below one. Put $n = h - l \geq 1$ and $q = \text{clip}(k - l, 0, n)$. The transition accumulates $\epsilon_q = m_q - m_{q-1}$, with $m_0 = 0, m_n = 1$.

The increments have a direct spectral meaning. For a geometric native grid with gap $c = \log b/K$,

$$x'_{l+q} - x'_{l+q-1} = c + \epsilon_q \log s. \quad (6)$$

Their unit sum distributes exactly $\log s$ of extra span across the band. At the tail, constant displacement $m = 1$ preserves native adjacent gaps, so the extra gap returns to zero.

MrRoPE-Pro uses increasing increments, $m_q = q(q+1)/[n(n+1)]$. We formulate a discrete tail-connection problem by penalizing variation in the extra log gaps and their mismatch with the fully interpolated tail:

$$\min_{\epsilon: \sum_{q=1}^n \epsilon_q = 1} \sum_{q=1}^{n-1} (\epsilon_{q+1} - \epsilon_q)^2 + \epsilon_n^2. \quad (7)$$

This one-sided minimum-bending objective smooths the interior log gaps and their connection to the stretched-distance reference. The high-frequency band stays fixed, with its entry gap free to accommodate the required span.

Theorem 2 (Finite-grid TailSpline). *The unique minimizer of Eq. (7) has positive increments and cumulative profile*

$$\epsilon_q = \frac{3(n+q)(n-q+1)}{n(n+1)(2n+1)}, \quad m_q^{\text{TS}} = \frac{q(3n^2 + 3n + 1 - q^2)}{n(n+1)(2n+1)}. \quad (8)$$

The solution assigns more extension throughout the transition: $m_q^{\text{TS}} \geq m_q^{\text{Pro}}$, with a larger entry gap and a smaller terminal gap (Fig. 3). The gap reduction has a direct phase consequence. At the last interior pair, for $n \geq 2$, $s > 1$ and any nonzero separation, residual phases relative to ω^N/s satisfy

$$\frac{\delta_{\text{TS}}}{\delta_{\text{Pro}}} = \frac{s^{\epsilon_n^{\text{TS}}} - 1}{s^{\epsilon_n^{\text{Pro}}} - 1} \leq \frac{3}{2n+1}. \quad (9)$$

Here $\delta_M = |\Delta| |\omega_{h-1}^M - \omega_{h-1}^N/s|$. The Llama $s = 4$ ratio is 0.080. More generally, Appendix B.5 relates phase mismatch to dependency stretch and gives a finite content-response bound. These properties explain the distance preference encoded by the construction.

Strict convexity gives the closed form (Appendix B.1). Penalizing both boundaries instead, by adding ϵ_1^2 , yields the symmetric construction BM with increments proportional to $q(n-q+1)$ (Appendix B.1).

What the comparison isolates. TailSpline and MrPro retain the same outer bands and endpoints, so their comparison tests complete internal allocations. The control C additionally matches total log-frequency displacement to test the remaining shape difference (Appendices B.2 and D.3).

Installation. (1) Read the public native grid, length and scale; (2) locate the 32/1-turn band; (3) compute m_k , $\omega'_k = \omega_k^N s^{-m_k}$ and shared gain $g = 1 + 0.1 \ln s$; (4) install before prefill across all layers and lengths. Construction costs $O(K)$, retains standard rotary computation and uses no checkpoint observations.

5.2 COSH FREQUENCY TRANSPORT

Cosh changes z through a monotone map of normalized exponents. Motivated by finite-window overlap among slow pairs, it balances log-frequency density concentration against the joint occupancy of the slow tail. The density objective and its unique minimizer are derived in Appendix B.3. Its inverse CDF is

$$Q_\tau(u) = 1 - \frac{\text{arcsinh}((1-u) \sinh \tau)}{\tau}, \quad u_k = \frac{k+1/2}{K}. \quad (10)$$

The uniform-density limit is recovered as $\tau \rightarrow 0$. Affine anchoring of $Q_\tau(u_k)$ to the chosen endpoints gives a convex allocation that moves interior frequencies toward the faster end; Appendix B.3 gives the installed grid and Eq. (32). This is the transport used in the fixed-range training intervention of §3.2. The strength τ determines the amount moved: the fixed-support reference is $d_{\text{head}}/\sqrt{L_{\text{train}}}$, and the MLA extrapolation recipe uses 1.414. The factorial compares neighboring strengths and a matched exponential transport (Appendix C.4). Every installation retains the standard rotary operator.

6 MODEL QUALITY WITHIN AND BEYOND THE NATIVE WINDOW

6.1 TASK QUALITY WITH A FROZEN TAILSPLINE TABLE

We compare TailSpline (T), MrRoPE-Pro (P) and static YaRN (Y) across Llama, OLMo, Qwen, GLM and Kanana, with a Llama-70B NF4 scale-transfer test. T/P share weights, outer bands, endpoints, gain, inputs and decoder; only internal allocation changes. Both use static $s = 4$ tables except Kanana’s 64K panel ($s = 2$). YaRN retains its public recipe, including Kanana’s publisher-recommended factor 4.4 (Peng et al., 2024; Kakao, 2025). All methods are installed without training. RULER (Hsieh et al., 2024) uses matched source-order prompts and official scoring.

Table 1: **Quality across lengths, families and scale.** Clean Full-13 scores (%). T/P use static $s = 4$ tables except Kanana ($s = 2$); its YaRN uses the official factor 4.4. T/P/Y: TailSpline/MrPro/static YaRN. Scores are paired within each row; (a)/(b) retain their stated sample sizes, with small Llama/OLMo panels drawn as subsets. Original RoPE scores 90.03% at Llama 8K. Intervals are paired 95% T-P contrasts (pp).

Model	Length	N/task	T	P	Y	T-P interval
<i>(a) Large-panel confirmation across the deployment window</i>						
Llama-8B	8K	50	85.16	82.39	–	[1.29, 4.25]
	16K	50	86.09	82.71	–	[1.53, 5.34]
	32K	200	68.27	56.54	–	[10.32, 13.11]
OLMo	16K	200	50.65	9.23	–	[39.99, 42.81]
<i>(b) Matched direct baselines and scale transfer</i>						
Llama	32K	10	70.38	58.35	55.88	[6.44, 17.68]
OLMo	16K	10	52.24	10.60	9.81	[35.38, 47.88]
Qwen	128K	10	64.40	59.87	62.04	[-0.41, 9.49]
GLM	128K	10	39.65	30.78	27.85	[4.04, 13.69]
Kanana	64K	10	72.65	70.33	65.64	[-3.13, 8.04]
Llama-70B NF4	32K	10	78.67	60.95	–	[11.65, 23.78]

Quality across the window. At Llama 32K, TailSpline improves Full-13 by 11.72 percentage points over MrPro on 2,600 paired prompts, with gains in all four task families and 12 of 13 tasks. At 16K, 650 pairs confirm the intermediate-window benefit; RULER QA improves by 8/9 points at 16/32K. The same table scores 85.16% on 650 native-length 8K prompts, with a slight reduction from original RoPE’s 90.03%. Figure 4 shows the task-level differences.




Figure 4: **One static allocation improves quality across context lengths.** Clean Llama Full-13 T-P differences. (a) Task-equal means with paired 95% intervals. (b) All task differences. The 8/16K panels use 50 inputs per task and 32K uses 200.

Direct baselines and transfer. On OLMo, the Full-13 gain reaches 41.42 points across 2,600 pairs, with all 13 tasks improving. The matched three-arm panels extend the comparison across model families, including GLM’s partial rotary subspace and Kanana’s 8-million base. TailSpline scores highest on each Full-13 panel; on Kanana it exceeds the publisher-recommended YaRN deployment by 7.01 points at 64K. Llama/OLMo NIAH-8 with 200 inputs per task and 46-document PPL also favor TailSpline over both baselines. Llama’s retrieval gains are concentrated in multikey tasks; OLMo improves across retrieval types (Appendix D.6). The allocation also transfers across model scale: directly reusing the Llama-8B frequency tables on the 70B NF4 checkpoint improves Full-13 by 17.72 points and PPL from 2.5075 to 2.4954.

Natural question answering. Question answering over complete passages and books tests the application-level value of improved retrieval (Yen et al., 2024). TailSpline improves OLMo’s five-task Natural-QA score, with gains on every task, and leads the Qwen book-QA panel and both GLM

book pools (Table 2). The 70B NF4 panel also scores 49.28/48.39% for T/P. At Kanana’s 128K deployment, TailSpline improves book-QA F1 over the recommended YaRN recipe by 1.74 points across 118 questions from 23 complete source contexts.

Natural-QA uses 631 questions from 524 source contexts. LongBench v2 (Bai et al., 2025) adds complete Llama inputs at 8K–32K. InfiniteBench En.QA (Zhang et al., 2024) uses complete books; the Qwen/GLM pool-I inputs span 100K–128K. To test transfer to new sources, GLM retains the same allocation on a second pool of 15 books disjoint from pool I’s seven, with complete inputs spanning 70K–95K.

Table 2: **Natural-task evaluation.** Scores (%): five-task macro F1 for Natural-QA, accuracy for LongBench v2 and official F1 for InfiniteBench book QA. I/II are separate book pools. Paired 95% T–P intervals use source contexts (pp); T–Y and simultaneous intervals are in Appendix D.7.

Model	Benchmark	N	T	P	Y	T–P interval
Llama-8B	Natural-QA	631	41.08	40.88	41.75	[−1.48, 1.86]
OLMo	Natural-QA	631	24.92	21.62	22.27	[0.39, 6.37]
Llama-70B NF4	Natural-QA	631	49.28	48.39	–	[−0.60, 2.44]
Llama-8B	LongBench v2	117	35.04	30.77	–	[−3.39, 11.97]
Qwen	Book QA (I)	35	19.14	15.90	18.45	[0.26, 7.51]
GLM	Book QA (I)	35	28.61	25.28	27.32	[−3.28, 8.98]
GLM	Book QA (II)	77	26.87	22.86	23.50	[−1.73, 9.68]
Kanana	Book QA	118	19.59	–	17.85	–

Internal shape at equal displacement. T–C matches endpoints, outer bands, gain and total log-frequency displacement. At 32K, T/C score 68.27/66.17%, a +2.10-point difference with interval [1.11, 3.08]; at 16K they score 86.09/86.35%. T assigns less stretch than C in the entry half of the transition and more in its tail half, at the same total displacement. This comparison links the target-length gain to how stretch is distributed (Appendix D.3). C has slightly higher positional effective rank at 32K (10.093/10.077), complementing Section 4. Higher-factor profiles at Llama 128K ($s = 16$) and Qwen 256K ($s = 8$) appear in Appendix D.8.

6.2 NATIVE-WINDOW ALLOCATION AT FIXED SUPPORT

The fixed-range findings also motivate an intervention within the native window. We use Native Contrastive Proximal allocation (NCP), a public-parameter construction that shifts interior frequencies while preserving both endpoints, frequency ordering, model weights and unit gain. Its reference target–distractor comparison favors a more stable aligned response while keeping shifts close to the native grid (Appendix E.1). One table is used across layers and inputs.

On OLMo’s native 4K Full-13 panel (780 pairs), the score rises from 69.37 to 70.78%, a +1.41-point gain [0.20, 2.58]; a new 130-pair panel adds +3.26 [0.81, 5.92]. Beyond these task diagnostics, we test context use by comparing full and recent histories while keeping the prediction targets fixed. Across 128 windows from 103 documents, scoring the final 256 tokens gives Native/NCP full-history NLL 2.91527/2.90248. NCP gains 0.00794 nat/token more from the additional history, interval [0.00553, 0.01035], connecting its language-modeling improvement to better use of earlier context. Appendix E provides complete task scores, natural-QA results and the same-target protocol.

6.3 ALLOCATION DURING TRAINING AND ADAPTATION

Training with Cosh lets weights adapt to a transported frequency allocation. A 432M MLA-style model with 16 rotary pairs trains for 500M tokens at 8K across three seeds (DeepSeek-AI, 2024a;b). Cosh improves all five extended lengths in every seed; 16K PPL falls from 138.8 to 95.6, and the benefit persists under a common wavelength blend (Fig. 5). In a matched 750M continuation, Cosh reduces 16K PPL from 45.1 to 24.4 and raises 8K answer-token exact match from 0/40 to 31/40.




Figure 5: **Cosh improves extrapolation across training seeds.** Three-seed mean and range for the 432M MLA model; dashed curves use the common wavelength blend ($s = 4$). Evaluation uses common text chunks.

The long-range gains extend to recipe-matched 1.485B OLMo training and matched Llama-3-8B LoRA, whose 16K PPL falls from 108.96 to 24.07. With phase-exposure and answer-plus-EOS supervision, matched OLMo adaptation also improves complete-answer retrieval: Native/Cosh score 18/98% at 8K and 0/60% at 16K. Learning thus makes use of the transported allocation in both language modeling and downstream retrieval. Appendix F gives the training and adaptation protocols and their native-window results.

7 RELATED WORK

Nonuniform scaling and frozen extension. YaRN blends interpolation and extrapolation across frequency bands (Peng et al., 2024). MrRoPE organizes scaling through mixed-radix conversion with MrRoPE-Pro progressively redistributing the intermediate increments (Tian et al., 2026). TailSpline retains MrPro’s outer bands and derives a different transition from a discrete tail-connection objective; matched support and total-displacement controls identify allocation and residual shape effects. LongRoPE and LongRoPE2 combine nonuniform scaling with search and learning (Ding et al., 2024; Shang et al., 2025); our deployment comparisons use frozen rules.

Learnable frequencies. LeRoPE learns per-frequency log scales jointly with model weights and analyzes their attention gradients (Karypis et al., 2026). Its fixed-frequency transfer experiment also separates the learned table from joint learning dynamics. Our work complements this direction with controlled interior-placement interventions and public-parameter constructions for frozen extension and native-window quality.

Positional structure and learned frequency use. Round and Round, HoPE and Massive Values investigate frequency roles and content use (Barbero et al., 2025; Chen et al., 2025; Jin et al., 2025); FoPE changes Fourier components (Hua et al., 2025). STRING characterizes relative-translation encodings, while GRAPE connects group generators, frequencies and norms (Schenck et al., 2025; Zhang et al., 2026). Our complete-pair analysis and integer-position criterion distinguish positional overlap, coordinate assignment and a changed spectrum. Selective RoPE and RePo instead learn input-dependent rotations or positions (Movahedi et al., 2026; Li et al., 2026).

8 DISCUSSION AND CONCLUSION

Choosing RoPE’s frequency range leaves an important design choice open: how its interior frequencies support the model’s use of context. Our controls isolate this choice, and the constructions turn it into improvements in frozen extension, native-window language modeling, and learned extrapolation. The resulting view of allocation connects positional structure with learned frequency use, making quality throughout the working context range a concrete design objective.

ETHICS STATEMENT

This work studies frequency allocation in existing language models using public research corpora and evaluation suites under their stated licenses. It introduces no human-subject study or deployed decision system. The proposed constructions retain the rotary operator and add no trainable parameters; adaptation experiments use the weight updates specified in their protocols.

REPRODUCIBILITY STATEMENT

The appendices specify the installed frequency tables, exponent grids, attention amplitudes, training and adaptation recipes, datasets, and evaluation procedures. The fixed-range experiment reports three paired training seeds; the larger-model and frozen-model tables identify their comparison units and sample counts. The source archive provides the active TeX, vector figures, numerical summaries and CPU checks of the Gram-matrix and allocation identities. It also includes the frozen two- and three-arm data-preparation, table-construction, generation, scoring and paired-report entrypoints with their local dependencies. The runtime guide specifies commands, public data sources and the pinned RULER revision.

The core fixed-support identification uses six runs of 499,974,144 tokens each; the 432M comparison uses six 500M-token runs, and the 750M continuation uses two 500M-token runs. These are protocol-specific training budgets, with the corresponding recipes and additional studies detailed in the appendices.

AI USE STATEMENT

We used generative AI tools to assist with mathematical discussion, code development, literature organization, and manuscript editing. Experimental results are drawn from recorded model runs and numerical analyses. The authors are responsible for the correctness and originality of the submitted text, mathematical arguments, experimental claims, and accompanying code.

REFERENCES

- Zhangir Azerbayev, Edward Ayers, and Bartosz Piotrowski. Proof-Pile. Dataset, 2022. URL <https://huggingface.co/datasets/hoskinson-center/proof-pile>.
- Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. LongBench: A bilingual, multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3119–3137, 2024. doi: 10.18653/v1/2024.acl-long.172. URL <https://aclanthology.org/2024.acl-long.172/>.
- Yushi Bai, Shangqing Tu, Jiajie Zhang, Hao Peng, Xiaozhi Wang, Xin Lv, Shulin Cao, Jiazheng Xu, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. Longbench v2: Towards deeper understanding and reasoning on realistic long-context multitasks. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3639–3664, 2025. doi: 10.18653/v1/2025.acl-long.183.
- Federico Barbero, Alex Vitvitskyi, Christos Perivolaropoulos, Razvan Pascanu, and Petar Veličković. Round and round we go! what makes rotary positional encodings useful? In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/e6d58fc68c0f3c36ae6e0e64478a69c0-Abstract-Conference.html.
- John J. Benedetto and Matthew Fickus. Finite normalized tight frames. *Advances in Computational Mathematics*, 18:357–385, 2003. doi: 10.1023/A:1021323312367.
- Åke Björck and Gene H. Golub. Numerical methods for computing angles between linear subspaces. *Mathematics of Computation*, 27(123):579–594, 1973. doi: 10.1090/S0025-5718-1973-0348991-3.
- Sidney Black, Stella Biderman, Eric Hallahan, Quentin Anthony, Leo Gao, Laurence Golding, Horace He, Connor Leahy, Kyle McDonell, Jason Phang, Michael Pieler, Usvsn Sai Prashanth, Shivanshu Purohit, Laria Reynolds, Jonathan Tow, Ben Wang, and Samuel Weinbach. GPT-NeoX-20B: An open-source autoregressive language model. In *Proceedings of BigScience Episode #5 – Workshop on Challenges & Perspectives in Creating Large Language Models*, pp. 95–136, 2022. doi: 10.18653/v1/2022.bigscience-1.9. URL <https://aclanthology.org/2022.bigscience-1.9/>.
- Yuhan Chen, Ang Lv, Jian Luan, Bin Wang, and Wei Liu. HoPE: A novel positional encoding without long-term decay for enhanced context awareness and extrapolation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 23044–23056, 2025. doi: 10.18653/v1/2025.acl-long.1123. URL <https://aclanthology.org/2025.acl-long.1123/>.
- DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024a. doi: 10.48550/arXiv.2405.04434. URL <https://arxiv.org/abs/2405.04434>.
- DeepSeek-AI. DeepSeek-V3 technical report. *arXiv preprint arXiv:2412.19437*, 2024b.
- Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. LongRoPE: Extending LLM context window beyond 2 million tokens. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pp. 11091–11104, 2024. URL <https://proceedings.mlr.press/v235/ding24i.html>.
- Cheng-Ping Hsieh, Simeng Sun, Samuel Krman, Shantanu Acharya, Dima Rekesh, Fei Jia, Yang Zhang, and Boris Ginsburg. RULER: What’s the real context size of your long-context language models? In *First Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=kIoBbc76Sy>.

- Ermo Hua, Che Jiang, Xingtai Lv, Kaiyan Zhang, Youbang Sun, Yuchen Fan, Xuekai Zhu, Biqing Qi, Ning Ding, and Bowen Zhou. Fourier position embedding: Enhancing attention’s periodic extension for length generalization. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 24932–24949, 2025. URL <https://proceedings.mlr.press/v267/hua25b.html>.
- Mingyu Jin, Kai Mei, Wujiang Xu, Mingjie Sun, Ruixiang Tang, Mengnan Du, Zirui Liu, and Yongfeng Zhang. Massive values in self-attention modules are the key to contextual knowledge understanding. *arXiv preprint arXiv:2502.01563*, 2025. URL <https://arxiv.org/abs/2502.01563>.
- Kakao. Kanana-1.5-8B-Instruct-2505 model card. <https://huggingface.co/kakaocorp/kanana-1.5-8b-instruct-2505>, 2025. Official runtime YaRN configuration.
- Petros Karypis, Sean O’Brien, Shreyas Kadekodi, Rui Zhu, and Julian McAuley. LeRoPE: Learnable RoPE frequencies improve language modeling. *arXiv preprint arXiv:2607.10134*, 2026. doi: 10.48550/arXiv.2607.10134. URL <https://arxiv.org/abs/2607.10134>.
- Huayang Li, Tianyu Zhao, Deng Cai, and Richard Sproat. RePo: Language models with context re-positioning. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026. URL <https://icml.cc/virtual/2026/poster/64002>. arXiv:2512.14391.
- Sajad Movahedi, Timur Carstensen, Arshia Afzal, Frank Hutter, Antonio Orvieto, and Volkan Cevher. Selective rotary position embedding. In *International Conference on Learning Representations (ICLR)*, 2026. URL <https://openreview.net/forum?id=AQo1SEElNb>.
- Guilherme Penedo, Hynek Kydlíček, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro von Werra, and Thomas Wolf. The FineWeb datasets: Decanting the web for the finest text data at scale. *arXiv preprint arXiv:2406.17557*, 2024. doi: 10.48550/arXiv.2406.17557. URL <https://arxiv.org/abs/2406.17557>.
- Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=wHBfxhZulu>.
- Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, and Timothy P. Lillicrap. Compressive transformers for long-range sequence modelling. *arXiv preprint arXiv:1911.05507*, 2019. URL <https://arxiv.org/abs/1911.05507>.
- Connor Schenck, Isaac Reid, Mithun George Jacob, et al. Learning the ropes: Better 2d and 3d position encodings with string. *arXiv preprint arXiv:2502.02562*, 2025. URL <https://arxiv.org/abs/2502.02562>.
- Ning Shang, Li Lyna Zhang, Siyuan Wang, Gaokai Zhang, Gilsinia Lopez, Fan Yang, Weizhu Chen, and Mao Yang. LongRoPE2: Near-lossless LLM context window scaling. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *PMLR*, pp. 54203–54218, 2025. URL <https://proceedings.mlr.press/v267/shang25a.html>.
- Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- Qingyuan Tian, Wenhong Zhu, Xiaoran Liu, Xiaofeng Wang, and Rui Wang. MrRoPE: Mixed-radix rotary position embedding. In *International Conference on Learning Representations (ICLR)*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/69413f87e5a34897cd010ca698097d0a-Abstract-Conference.html.
- Xinyi Wu, Siyuan Liu, and Ali Jadbabaie. How data shapes RoPE frequency usage: From positional scale matching to length generalization. *arXiv preprint arXiv:2607.07678*, 2026. doi: 10.48550/arXiv.2607.07678. URL <https://arxiv.org/abs/2607.07678>.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. doi: 10.48550/arXiv.2412.15115. URL <https://arxiv.org/abs/2412.15115>.

- Howard Yen, Tianyu Gao, Minmin Hou, Ke Ding, Daniel Fleischer, Peter Izsak, Moshe Wasserblat, and Danqi Chen. Helmet: How to evaluate long-context language models effectively and thoroughly. *arXiv preprint arXiv:2410.02694*, 2024. URL <https://arxiv.org/abs/2410.02694>.
- Xinrong Zhang, Yingfa Chen, Shengding Hu, Zihang Xu, Junhao Chen, Moo Hao, Xu Han, Zhen Thai, Shuo Wang, Zhiyuan Liu, and Maosong Sun. ∞ Bench: Extending long context evaluation beyond 100k tokens. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pp. 15262–15277, 2024. doi: 10.18653/v1/2024.acl-long.814. URL <https://aclanthology.org/2024.acl-long.814/>.
- Yifan Zhang, Zixiang Chen, Yifeng Liu, Zhen Qin, Huizhuo Yuan, Kangping Xu, Yang Yuan, Quanquan Gu, and Andrew Chi-Chih Yao. Group representational position encoding. In *International Conference on Learning Representations (ICLR)*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/5cb58625f49ddf70fe2d527e9e4bbae5-Paper-Conference.pdf.

SUPPLEMENTARY MATERIAL

The supplement follows the main argument. Appendix A proves the positional statements; Appendix B derives the allocations; Appendix C specifies the controlled interventions. Complete frozen, native-window and learned results are in Appendices D, E and F, respectively. The source archive includes numerical inputs, construction/scoring code and CPU verification scripts.

A POSITIONAL GEOMETRY AND LEARNED USE

A.1 CROSS-GRAM FORMULA, INVARIANCE AND EFFECTIVE RANK

Closed-form cross-Gram. For $\Delta \sim \text{Unif}[0, L]$ and $x_\omega(\Delta) = [\cos(\omega\Delta) \ \sin(\omega\Delta)]$, write $d = (\omega - \nu)L$, $s = (\omega + \nu)L$, $a(t) = \sin t/t$ and $b_\star(t) = (1 - \cos t)/t$. Expanding the products of sinusoids and integrating termwise gives

$$H_{\omega\nu} = \mathbb{E}[x_\omega^\top x_\nu] = \frac{1}{2} \begin{bmatrix} a(d) + a(s) & b_\star(s) - b_\star(d) \\ b_\star(s) + b_\star(d) & a(d) - a(s) \end{bmatrix}, \quad (11)$$

and $S_\omega = H_{\omega\omega}$. Since $Q_{\omega\nu} = S_\omega^{-1/2} H_{\omega\nu} S_\nu^{-1/2}$ is the whitened cross-Gram of two two-dimensional subspaces, its singular values are the canonical correlations, and $c_{\omega\nu} = \frac{1}{2} \|Q_{\omega\nu}\|_F^2$ is invariant under $x_\omega \mapsto x_\omega O$ for any invertible O (in particular under rotation of the phase inside a pair), because S_ω transforms congruently and the whitening cancels O .

Basis invariance. The cancellation can be stated exactly. Under independent invertible basis changes $x_\omega \mapsto x_\omega A_\omega$ and $x_\nu \mapsto x_\nu A_\nu$, let $U_j = S_j^{1/2} A_j (A_j^\top S_j A_j)^{-1/2}$ for $j \in \{\omega, \nu\}$. Then $U_j^\top U_j = I$, and the transformed whitened cross-Gram is $\tilde{Q}_{\omega\nu} = U_\omega^\top Q_{\omega\nu} U_\nu$. Left and right orthogonal factors preserve the two canonical correlations and therefore $c_{\omega\nu}$. Thus $c_{\omega\nu}$ is an intrinsic measure of overlap between the two positional subspaces.

Theorem 3 (Spectral budget identity). *For the block-whitened Gram matrix Γ with diagonal blocks I_2 , $\text{tr } \Gamma = 2K$ and $\text{tr}(\Gamma^2) = 2K[1 + (K - 1)\bar{c}]$. Hence its Rényi-2 effective rank is given by (5).*

Proof of Theorem 3. Γ is the block matrix with blocks $\Gamma_{ij} = Q_{\omega_i \omega_j}$ and $\Gamma_{ii} = I_2$. Hence $\text{tr } \Gamma = \sum_{i=1}^K \text{tr } I_2 = 2K$. For the second moment, $\text{tr}(\Gamma^2) = \sum_{i,j} \text{tr}(\Gamma_{ij} \Gamma_{ji}) = \sum_i \text{tr}(I_2) + \sum_{i \neq j} \|Q_{\omega_i \omega_j}\|_F^2 = 2K + 2K(K - 1)\bar{c}$, using $\Gamma_{ji} = \Gamma_{ij}^\top$ and the definition $\bar{c} = \frac{1}{K(K-1)} \sum_{i \neq j} c_{\omega_i \omega_j}$ with $\|Q_{ij}\|_F^2 = 2c_{ij}$. Dividing gives (5). $\square$

The identity determines Rényi-2 effective rank exactly from the mean pairwise redundancy. Shannon effective rank and log-determinant retain higher-order multi-subspace dependence on the full spectrum.

Proof of Proposition 1. Put $t = \Delta/L$, $x = \omega L$, and use $\theta = x^2$ as the local parameter. The rescaled pair has the analytic basis

$$u_1(\theta, t) = \cos(\sqrt{\theta} t) = 1 - \frac{1}{2}\theta t^2 + O(\theta^2), \quad (12)$$

$$u_2(\theta, t) = \frac{\sin(\sqrt{\theta} t)}{\sqrt{\theta}} = t - \frac{1}{6}\theta t^3 + O(\theta^2). \quad (13)$$

Thus $V_\omega \rightarrow V_0 = \text{span}\{1, \Delta\}$. The fourth-order coefficient follows by retaining the part of the first perturbation transverse to V_0 . Under the uniform measure on $t \in [0, 1]$,

$$\Pi_0^\perp t^2 = p_2(t) := t^2 - t + \frac{1}{6}, \quad (14)$$

$$\Pi_0^\perp t^3 = p_3(t) := t^3 - \frac{9}{10}t + \frac{1}{5}, \quad (15)$$

where $\Pi_0^\perp$ is the $L_2[0, 1]$ projection off $\text{span}\{1, t\}$. In the base $[1, t]$, its Gram and the residual perturbation Gram are

$$G_0 = \begin{bmatrix} 1 & 1/2 \\ 1/2 & 1/3 \end{bmatrix}, \quad M = \begin{bmatrix} \langle -p_2/2, -p_2/2 \rangle & \langle -p_2/2, -p_3/6 \rangle \\ \langle -p_3/6, -p_2/2 \rangle & \langle -p_3/6, -p_3/6 \rangle \end{bmatrix} = \begin{bmatrix} 1/720 & 1/1440 \\ 1/1440 & 1/2800 \end{bmatrix}. \quad (16)$$

The squared chordal distance between the two nearby planes is therefore

$$2 - \|Q_{x,y}\|_F^2 = (x^2 - y^2)^2 \text{tr}(G_0^{-1}M) + O(\epsilon^6) = \frac{19}{12600}(x^2 - y^2)^2 + O(\epsilon^6), \quad (17)$$

with $\epsilon = \max\{|x|, |y|\} \rightarrow 0$. At $x = 0.05, y = 0.10$, the analytic Gram gives exact/leading ratio 1.00058.

For the softmax metric with fixed p , $F = \text{diag}(p) - pp^\top$ satisfies $F\mathbf{1} = 0$, so the constant direction is annihilated before whitening. Writing $\bar{h} = h - \mathbb{E}_p h$, the same expansion gives

$$\frac{\sin(\omega\Delta)}{\omega} \rightarrow \Delta - \mathbb{E}_p \Delta, \quad -\frac{2\cos(\omega\Delta)}{\omega^2} \rightarrow \Delta^2 - \mathbb{E}_p \Delta^2. \quad (18)$$

These functions span the centred limit whenever p has nondegenerate support on at least three distances. The separation weights p therefore determine which polynomial directions survive centering. $\square$

A.2 FINITE-WINDOW OVERLAP IN THE IDENTIFICATION EXPERIMENT

The slow-limit proposition has a finite-window counterpart that can be calculated directly on the identification grid. Set $b = 256$, $K = 32$, $L = 256$, and retain the n slowest frequencies of $\omega_j = b^{-j/K}$. Their dimensionless frequencies are $\omega_j L = 1.1892, 1.4142, 1.6818, 2, \dots$ in slow-to-fast order. The smallest is already greater than one, so we evaluate the full cross-Gram in (11) without a small-phase approximation.

Table 3: **Finite-window positional concentration on the paired-training grid.** The separation measure is uniform on $[0, L]$. The nominal coordinate count is $2n$; r_2 is the effective rank of the block-whitened Gram.

Slowest pairs n	Largest ωL	Coordinates	r_2
2	1.41	4	2.00030
4	2.00	8	2.00361
6	2.83	12	2.02263
8	4.00	16	2.11474
12	8.00	24	2.89256
16	16.00	32	4.43339

For $n = 8$, the largest two eigenvalues are 7.9681 and 7.5752, accounting for 97.1% of the trace. This is the finite-window geometry in Fig. 2; the full-pair calculation uses no small-phase approximation.

A.3 TWO STATIC ORDERING COUNTEREXAMPLES

Two explicit examples show how phase coupling and interval length change the ordering of frequency tables.

Cosine-only collision can choose the lower-rank table. Take $K = 4$, $L = 8\pi$, and the two descending frequency tables

$$\Omega_A = (1, 5/8, 3/8, 1/4), \quad \Omega_B = (1, 6.01/8, 4.01/8, 1/4).$$

They share both endpoints. Define $C_{\cos}$ as the mean squared normalized cosine-cosine overlap over distinct pairs. For Ω_A , every cosine is an integer harmonic on $[0, L]$, so $C_{\cos}(A) = 0$ exactly. Mixed harmonic parities nevertheless produce nonzero sine-cosine cross terms. For Ω_B , the small perturbation of an even-harmonic table gives

$$C_{\cos}(A) = 0 < 1.38725 \times 10^{-5} = C_{\cos}(B), \quad (19)$$

$$r_2(A) = 6.29478 < 7.99971 = r_2(B). \quad (20)$$

Thus a strict cosine-only preference can have the opposite full-subspace rank ordering. Both values follow directly from Eq. (11).

Full-subspace collision can reverse across lengths. Now use $K = 4$, $L = 2\pi$, and the fixed tables

$$\Omega_A = (1, 0.99, 0.02, 0.01), \quad \Omega_B = (1, 2/3, 1/3, 0.01).$$

Again the endpoints agree. Writing C_M for mean phase-invariant collision under the uniform measure on $[0, M]$, Eq. (11) gives

$$C_L(A) = 0.538996 < 0.631721 = C_L(B), \quad (21)$$

$$C_{2L}(A) = 0.384445 > 0.204995 = C_{2L}(B), \quad (22)$$

$$C_{4L}(A) = 0.345188 > 0.033506 = C_{4L}(B). \quad (23)$$

The close pairs in A retain high overlap as the interval grows, while the more evenly spread frequencies in B become easier to distinguish. The example makes the interval dependence of the geometry explicit.

Closed-form Gram evaluations agree with independent numerical quadrature. These examples make the roles of complete pairs and evaluation length explicit.

A.4 INTEGER-POSITION KERNEL EQUIVALENCE AND ITS BOUNDARIES

Corollary 4 (Integer-position equivalence). *Let Ω, Λ contain K frequencies in $(0, \pi)$ and $\mathcal{R}_\Omega(d) = \text{diag}_k R(\omega_k d)$. Position-independent invertible A, B satisfy $A^\top \mathcal{R}_\Omega(d) B = \mathcal{R}_\Lambda(d)$ for $d = 0, 1$ if and only if the frequency multisets agree.*

Corollary 4 applies to the full bilinear kernel on the $2K$ -dimensional rotary space. Equality for all q, k is matrix equality. At $d = 0$, $A^\top B = I$; at $d = 1$ the two rotation matrices are similar. Their complex eigenvalues are $e^{i\omega_k}, e^{-i\omega_k}$ including multiplicities. On $(0, \pi)$ the positive-angle eigenvalue identifies each frequency uniquely, proving necessity even with repeated frequencies. For sufficiency choose a block permutation P with $\mathcal{R}_\Lambda(d) = P \mathcal{R}_\Omega(d) P^\top$, and set $A = B = P^\top$. This holds for every integer d , including negative distances.

The interval $(0, \pi)$ selects unique positive rotation angles: outside it, ω and $\omega + 2\pi$ are indistinguishable at integer positions. The statement concerns equality of the full bilinear kernel for every content pair. These conditions are verified by the bundled discrete-kernel checker.

A.5 POSITIONAL OVERLAP PRESERVES CONTENT COORDINATES

For a block $\mathcal{B}$ of r rotary pairs, write $\mathcal{R}_\mathcal{B}(d) = \text{diag}_{k \in \mathcal{B}} R(\omega_k d)$. Each block is orthogonal, so $\mathcal{R}_\mathcal{B}(d)$ has rank $2r$ at every distance. This is the rank of an operator on content coordinates; the positional Gram instead measures shared dependence on distance.

Proposition 5 (Content matching in a slow block). *Let $H > 0$ and $\eta = H \max_{k \in \mathcal{B}} |\omega_k|$. For $|d| \leq H$,*

$$|q_\mathcal{B}^\top \mathcal{R}_\mathcal{B}(d) k_\mathcal{B} - q_\mathcal{B}^\top k_\mathcal{B}| \leq \eta \|q_\mathcal{B}\|_2 \|k_\mathcal{B}\|_2. \quad (24)$$

For $2r$ content symbols represented by distinct coordinate unit vectors, let the query and correct key share a vector and all incorrect keys use other vectors. Their unscaled score margin is at least $1 - 2\eta$ across arbitrary key distances in $[-H, H]$.

Proof. For each pair, $\|R(\omega d) - I\|_2 = 2|\sin(\omega d/2)| \leq |\omega d|$. The block-diagonal norm is the maximum pair norm. The bilinear norm bound gives Eq. (24). Unrotated matching scores are one for the correct symbol and zero for an incorrect symbol. Each score changes by at most η , giving the margin bound. In particular, $\eta < 1/2$ preserves the correct ranking. $\square$

As distinct frequencies approach zero, their positional subspaces approach $\text{span}\{1, d\}$ while the content construction remains distinct in all $2r$ coordinates. This provides a concrete distinction between repeated positional dependence and repeated content information, consistent with prior analyses of frequency use (Barbero et al., 2025). The construction concerns specified content vectors and the small-phase regime; the next subsection evaluates the complete finite frequency tables.

A.6 TASK GAINS WITH LOWER FULL-PAIR EFFECTIVE RANK

We use the Llama mathematical native grid $\omega_k^N = 500000^{-k/64}$, native length 8192, $s = 4$ and band [18, 35]. The TailSpline and MrPro formulas determine all frequencies without model observations. Effective rank uses the full-pair Gram from Eq. (11), with uniform separations on $[0, H]$.

Table 4: **A successful allocation can have lower positional effective rank.** Ranks describe the full mathematical frequency tables; task scores reuse the clean paired evaluations in Table 1. The 16K and 32K task panels contain 650 and 2600 paired prompts.

Window	r_2 TailSpline	r_2 MrPro	RULER TailSpline (%)	RULER MrPro (%)
16K	8.283391	8.739887	86.0949	82.7051
32K	10.077270	10.214888	68.2660	56.5436

Continuous and integer-position Grams, FP32 rounding and causal-separation weights retain this rank ordering at 8/16/32K. The table separates positional structure from task quality; a lower rank is not used as a selection rule.

From distance response to content competition. Fix the pre-RoPE query and keys for one attention read. With content coefficients C_{jk} and separations d_j , its logits can be written $\ell_j(\nu) = \text{Re} \sum_k C_{jk} e^{id_j \nu_k}$, absorbing the attention normalization into C_{jk} . Changing the table gives the finite shift $\eta_j = \text{Re} \sum_k C_{jk} (e^{id_j \nu'_k} - e^{id_j \nu_k})$. For nonempty target and competing groups G, B partitioning the valid keys, let p_G, p_B be their separately normalized original softmax weights. The log-odds $\mathcal{O} = \log \sum_G e^{\ell_j} - \log \sum_B e^{\ell_j}$ obey

$$\mathcal{O}' - \mathcal{O} = \log \mathbb{E}_{p_G} e^{\eta_j} - \log \mathbb{E}_{p_B} e^{\eta_j}. \quad (25)$$

Indeed, factor each original group partition function from its updated sum. If $\min_G \eta_j > \max_B \eta_j$, the target share strictly increases. This sufficient condition compares target evidence with its competitors, rather than using positional rank alone. The identity concerns one fixed-state read; a full forward pass also changes states in subsequent layers. It is an explanatory softmax identity, not a task-optimality criterion.

B EXPLICIT CONSTRUCTIONS AND DISTANCE SCALING

B.1 TAILSPINE AND THE SYMMETRIC CONSTRUCTION

Let $D\epsilon = (\epsilon_{q+1} - \epsilon_q)_{q=1}^{n-1}$ and $H = D^\top D + e_n e_n^\top$. A zero value of $\epsilon^\top H \epsilon$ requires constant increments and a zero terminal increment, hence $\epsilon = 0$. Thus H is positive definite. For $n > 1$, the proposed $\epsilon_q = c(n+q)(n-q+1)$, with $c = 3/[n(n+1)(2n+1)]$, satisfies

$$\epsilon_1 - \epsilon_2 = 2c, \quad 2\epsilon_q - \epsilon_{q-1} - \epsilon_{q+1} = 2c, \quad 2\epsilon_n - \epsilon_{n-1} = 2c.$$

Therefore $H\epsilon = 2c\mathbf{1}$, the equality-constrained stationarity equation. The increments are positive and sum to one; summing them yields Eq. (8). The objective minimum is $2c = 6/[n(n+1)(2n+1)]$. For $n = 1$, the constraint gives $\epsilon_1 = 1$, agreeing with the formulas. Positivity makes additional nonnegativity constraints inactive.

Adding ϵ_1^2 yields the symmetric objective, with unique solution

$$\epsilon_q^{\text{BM}} = \frac{6q(n-q+1)}{n(n+1)(n+2)}, \quad m_q^{\text{BM}} = \frac{q(q+1)(3n+2-2q)}{n(n+1)(n+2)}. \quad (26)$$

Both profiles are monotone with $m_0 = 0, m_n = 1$.

What the objective smooths. On the geometric native grid, Eq. (6) makes each change between neighboring log gaps equal to $(\epsilon_{q+1} - \epsilon_q) \log s$. The terminal change is $-\epsilon_n \log s$ because the fully interpolated tail has constant displacement. Thus $(\log s)^2$ times Eq. (7) is exactly the squared log-gap variation within the transition and at its low-frequency junction. The high-frequency junction is deliberately unpenalized.

At the terminal gap,

$$\epsilon_n^{\text{Pro}} = \frac{2}{n+1}, \quad \epsilon_n^{\text{TS}} = \frac{6}{(n+1)(2n+1)}, \quad \frac{\epsilon_n^{\text{TS}}}{\epsilon_n^{\text{Pro}}} = \frac{3}{2n+1}.$$

The opposite entry behavior is also explicit: $\epsilon_1^{\text{Pro}} = 2/[n(n+1)]$, whereas $\epsilon_1^{\text{TS}} = 3/(2n+1)$. TailSpline trades a larger entry jump for a smaller terminal jump. A useful boundary prior need not minimize both. In the continuum limit, $m^{\text{TS}}(u) = (3u - u^3)/2$ solves $\min \int_0^1 (m''(u))^2 du$ under $m(0) = 0, m(1) = 1, m'(1) = 0$; the natural entry condition is $m''(0) = 0$. The finite-grid formula is used in every model experiment, not this limiting polynomial.

B.2 AN EXACT EQUAL-DISPLACEMENT CONTROL

Use local indices $q = 0, \dots, n$. Write $T_q = m_q^{\text{TS}}, B_q = m_q^{\text{BM}}, U_q = q/n$, and $P_q = q(q+1)/[n(n+1)]$. Define the front-loaded profile $F_q = 2U_q - P_q$ and the control

$$C_q = (1 - w_n)U_q + w_n F_q, \quad w_n = \frac{3n}{2(2n+1)}. \quad (27)$$

Substitution of the finite-grid formulas gives $T = (1 - w_n)B + w_n F$. Subtracting C gives

$$m^{\text{TS}} - m^C = (1 - w_n)(m^{\text{BM}} - m^{\text{Uni}}). \quad (28)$$

In local coordinates, the explicit difference is

$$T_q - C_q = (1 - w_n)(B_q - U_q) = \frac{q(n-q)(2q-n)}{2n(n+1)(2n+1)}. \quad (29)$$

This cumulative-profile difference is antisymmetric under $q \mapsto n - q$, vanishes at both endpoints, and sums to zero. Since U and F have positive increments and $0 < w_n < 1$, C is monotone with $C_0 = 0, C_n = 1$. Extending both T and C by the same constant outer bands gives identical endpoints and total log displacement at a common scale. Their equal increment centroids follow from Eq. (42), rather than imposing a further constraint. For $n = 1, 2$, $B = U$ and $T = C$; for wider bands the contrast tests higher-order shape.

For $q < n/2$, $T_q < C_q$; for $q > n/2$, $T_q > C_q$. TailSpline thus allocates less stretch in the entry half and more in the tail half, at the same total displacement. The clean T-C results are in Appendix D.3.

B.3 COSH DENSITY AND FINITE INSTALLATION

Let $S_\rho(t) = \int_t^1 \rho(\phi) d\phi$. For $\alpha > 0, \beta \geq 0$, the Cosh transport minimizes

$$\mathcal{C}_{\text{app}}[\rho] = \frac{\alpha}{2} \int_0^1 \rho^2 + \frac{\beta}{2} \int_0^1 S_\rho^2, \quad \rho \geq 0, \quad \int_0^1 \rho = 1. \quad (30)$$

The first term penalizes concentrated log-frequency density. By Fubini, $\int_0^1 S_\rho(t)^2 dt = \mathbb{E}[\min(\Phi_1, \Phi_2)]$ for independent $\Phi_i \sim \rho$, so the second penalizes joint occupancy toward the slow end. These declared structural preferences define the transport objective. Coercivity and weak lower semicontinuity on the closed convex constraint set give existence; the first term makes the objective strictly convex. For $g(\phi) = \int_0^1 \min(\phi, \psi) \rho(\psi) d\psi$, stationarity is $\alpha\rho + \beta g + \lambda = 0$. Since $g'' = -\rho$, $g'(1) = 0$ and $g'(0) = 1$, this gives $\rho'' = \tau^2 \rho$, $\rho'(0) = -\tau^2$ and $\rho'(1) = 0$, where $\tau^2 = \beta/\alpha$. Thus the unique solution is

$$\rho_\tau(\phi) = \frac{\tau \cosh[\tau(1 - \phi)]}{\sinh \tau}. \quad (31)$$

It is positive and has unit mass, so it satisfies the inequality constraint. At $\beta = 0$ the solution is uniform. Integrating and inverting its CDF yields Eq. (10). Equivalently, for $h = Q'$, $\mathcal{C}_{\text{app}} = \frac{1}{2} \int_0^1 [\alpha/h + \beta(1 - u)^2 h] du$.

For midpoint quantiles $u_k = (k + 1/2)/K$, anchored installation uses

$$z_k = \frac{Q_\tau(u_k) - Q_\tau(u_0)}{Q_\tau(u_{K-1}) - Q_\tau(u_0)} \leq \frac{k}{K-1}, \quad \omega_k = \exp(-a - Rz_k). \quad (32)$$

The inequality follows from convexity of Q_τ and its endpoint chord. Unanchored midpoint installation instead uses $\omega_k = b^{-Q_\tau(u_k)}$. Its uniform limit is midpoint Geo; standard Geo uses k/K , differing by the common frequency factor $b^{-1/(2K)}$. Each experimental protocol below specifies its actual grid and support convention. Native denotes the inherited frequency table of a released model.

B.4 SCALE RESPONSE AND FINITE PHASE CHANGES

Write $\nu_k(s) = \omega_k^N s^{-m_k(s)}$ and let $A_k(s) = \omega_k^N / \nu_k(s)$ be the wavelength multiplier. For the index ramp in the released YaRN implementation (Peng et al., 2024), with fixed interior coordinate $0 < t < 1$,

$$A_Y(s, t) = \frac{1}{1 - t + t/s}, \quad \lim_{s \rightarrow \infty} A_Y(s, t) = \frac{1}{1 - t}. \quad (33)$$

A fixed frequency blend $\gamma + (1 - \gamma)/s$ similarly has limit $1/\gamma$ for $\gamma > 0$; the fully interpolated tail $\gamma = 0$ grows with s . MrPro and TailSpline instead have $A_P = s^{P_q}$ and $A_T = s^{T_q}$ for their fixed finite-grid exponents. Their log-scale derivatives are

$$\frac{\partial \log A_Y}{\partial \log s} = \frac{t}{s(1 - t) + t}, \quad \frac{\partial \log A_P}{\partial \log s} = P_q, \quad \frac{\partial \log A_T}{\partial \log s} = T_q. \quad (34)$$

These identities follow by differentiation. They describe constructor responses to s , rather than task-performance guarantees.

Index and turn-count ramps use different blend weights. All numerical comparisons here use one shared band and the index-ramp equation of the released YaRN implementation (revision 995db5b). They are frequency-table comparisons, not a re-evaluation of the YaRN baseline runs in the MrRoPE paper.

Magnitude of the TailSpline intervention. Direct subtraction of the finite-grid formulas gives

$$T_q - P_q = \frac{q(n - q)(3n + q + 1)}{n(n + 1)(2n + 1)} \geq 0, \quad \frac{A_T}{A_P} = s^{T_q - P_q}. \quad (35)$$

For the Llama $n = 17$ grid, the maximum wavelength ratio is 1.328763, 1.765612, 2.346080 and 3.117385 at $s = 2, 4, 8, 16$. At pair $k = 26$ ($q = 8$), the multipliers are:

Scale	YaRN index blend	MrPro	TailSpline
4	1.545455	1.385674	2.423868
16	1.789474	1.920093	5.875134

The YaRN/MrPro ordering at an interior pair can therefore change with scale. TailSpline redistributes substantial wavelength growth throughout the transition, in addition to its smaller terminal extra-gap jump.

B.5 TWO REFERENCE RESPONSES AND FINITE PHASE INTERVALS

The native reference is ω and the fully interpolated reference is ω/s . For a declared phase tolerance $0 < \theta \leq \pi$, define

$$D_N(\nu; \theta) = \frac{\theta}{|\omega - \nu|}, \quad D_D(\nu; \theta) = \frac{\theta}{|\nu - \omega/s|}, \quad (36)$$

with infinity at a zero denominator. These are the maximal intervals starting at zero on which the respective unwrapped phase deviation is at most θ . Since $\omega/s \leq \nu_T \leq \nu_P \leq \omega$,

$$D_D(\nu_T; \theta) \geq D_D(\nu_P; \theta), \quad D_N(\nu_T; \theta) \leq D_N(\nu_P; \theta). \quad (37)$$

The exact operator difference is $2|\sin(d(\nu - \nu_{\text{ref}})/2)|$. Thus the distances quantify two specified phase references; they are not effective context lengths of a model.

Tail-junction phase residual. At $q = n - 1$, $m_q = 1 - \epsilon_n$ and hence $\nu - \omega/s = (\omega/s)(s^{\epsilon_n} - 1)$. For $n \geq 2$ and $s > 1$, $0 < \epsilon_n^{\text{TS}} \leq \epsilon_n^{\text{Pro}}$. Since $(e^x - 1)/x$ is increasing for $x > 0$, dividing the two residuals gives Eq. (9). It quantifies the approach to the stretched distance reference at the final interior pair.

Dependency stretch and response. Let a reference dependency have separation $d_0 > 0$ and its extended separation be αd_0 . The residual phase in channel q is $\delta_M = \omega d_0 (\alpha s^{-m_q^M} - 1)$. For an interior pair with $T_q > P_q$ and $s > 1$, the smaller absolute residual switches from P to T at

$$\alpha_q^* = \frac{2}{s^{-T_q} + s^{-P_q}}. \quad (38)$$

Indeed, with $v_M = s^{-m_q^M}$, $\delta_T^2 - \delta_P^2 = (\omega d_0)^2 \alpha (v_T - v_P) [\alpha (v_T + v_P) - 2]$. Thus P has smaller residual for $0 < \alpha < \alpha_q^*$, and T for $\alpha > \alpha_q^*$. The relevant scale is the dependency’s stretch, not the total input length.

To allow a nonzero content phase, write a reference-aligned complex coefficient as $a + ib$, with $a > 0$ and $|b| \leq \kappa a$. Its response at residual δ is $f(\delta) = a \cos \delta - b \sin \delta$. If two residuals have the same sign and $0 \leq r = |\delta_A| < t = |\delta_B| \leq \pi$, then

$$f(\delta_A) - f(\delta_B) \geq 2a \sin \frac{t-r}{2} \left(\sin \frac{t+r}{2} - \kappa \left| \cos \frac{t+r}{2} \right| \right). \quad (39)$$

This follows from the cosine and sine difference identities and $|b| \leq \kappa a$; an allowed endpoint value of b attains the bound. A positive bracket guarantees a larger response. At $\kappa = 0$ it recovers the aligned cosine comparison; beyond that case the allowable phase offset is explicit. These are fixed-content responses. Their contribution to attention depends on competing keys, as specified in Appendix A.6. For example, $a = 1$, $b = -1$ and positive residuals 0.1/0.3 give a response difference of -0.156 , illustrating why the phase condition matters.

Equal displacement does not imply a small phase intervention. On the same Llama $s = 4$ mathematical grid, the largest T/C unwrapped phase difference over integer distances $0, \dots, 32767$ is 10.672514 radians. The corresponding T/P value is 108.197815. For finite comparisons one can use

$$\rho_H(\delta) = \frac{4}{H} \sum_{d=0}^{H-1} \sin^2(d\delta/2) = 2 - 2 \frac{\sin(H\delta/2)}{H \sin(\delta/2)} \cos((H-1)\delta/2), \quad (40)$$

with removable singularities evaluated by continuity. The fixed-slot rotation RMS is $[K^{-1} \sum_k \rho_H(\nu_k - \tilde{\nu}_k)]^{1/2}$. At $H = 32768$, it is 0.448329 for T/C and 0.687543 for T/P. These values compare unscaled rotary kernels, before the shared cosine/sine gain and the attention normalization. For $H|\delta| \ll 1$, $\rho_H(\delta) = ((H-1)(2H-1)/6)\delta^2 + O(H^4\delta^4)$. The small parameter is therefore distance times frequency difference; matching total displacement alone does not supply it.

Turn boundaries across native bases. The 32/1-turn rule places the transition in a comparable native phase range. For a geometric grid with $c = \log b/K$, assume both boundaries lie inside the sampled table without clipping. Let l be the last slot above 32 turns, h the first below one, $n = h - l$ and $u = q/n$. Define $\eta_+ = \log(L\omega_l/(64\pi))$ and $\eta_- = \log(2\pi/(L\omega_h))$. Both belong to $(0, c]$, and

$$nc = \log 32 + \eta_+ + \eta_-, \quad L\omega_{l+q} = 64\pi 32^{-u} \exp((1-u)\eta_+ - u\eta_-). \quad (41)$$

These identities follow from $\omega_{l+q} = \omega_l e^{-cq}$. Together with $d\nu = (d/L)(L\omega)s^{-m}$, they explain why a larger base alone does not remove a turn-selected allocation change. The remaining geometric differences include finite sampling, rounding and transition width; model task behavior additionally depends on learned content use. The fine-grid limit is $c \rightarrow 0$, not increasing the base at fixed K .

C CONTROLLED ALLOCATION EXPERIMENTS

C.1 INTERVENTION IDENTITIES

With the native fast endpoint fixed and its slow endpoint divided by s , $z'_k = (Rz_k + \log s m_k)/(R + \log s)$. Fixing the cumulative increment mass $\sum_q \epsilon_q = 1$ fixes the added log span; fixing total displacement $D = \log s \sum_k m_k$ further constrains the interior allocation. In particular,

$$\sum_{q=1}^n q\epsilon_q = n - \sum_{q=1}^{n-1} m_q, \quad D = Ka' + R' \sum_k z'_k - \sum_k x_k^N. \quad (42)$$

Thus equal total displacement already fixes the increment centroid. A frequency-multiset permutation instead keeps the spectrum and changes its coordinate assignment. Gain is controlled separately in every matched comparison.

C.2 FIXED-SUPPORT TRAINING AND RANGE INTERACTION

The paired 151,898,880-parameter decoder has 12 layers, width 768, 12 heads of dimension 64, MLP width 3072 and vocabulary 50304. Within each seed (42, 137, 256), Geo and anchored Cosh share initialization, token order, optimizer, batch and training length 256. Each arm sees 499,974,144 tokens in 7629 updates of 256 sequences. Seed 42 uses microbatch 64 with accumulation 4; the other seeds use 128×2 . Fused AdamW uses $(\beta_1, \beta_2) = (0.9, 0.95)$, weight decay 0.01, clipping 1, peak/minimum LR $6 \times 10^{-4}/6 \times 10^{-5}$ and 762 warmup steps followed by cosine decay.

FineWeb-Edu training and validation use disjoint unshuffled shards `sample/10BT/000_00000.parquet` and `sample/10BT/004_00000.parquet`, revision `87f09149ef47`, with GPT-NeoX-20B tokenizer revision `c292233c833e` (Penedo et al., 2024; Black et al., 2022). Geo uses $b = 256$ and $\omega_k = \exp[-Rk/(K-1)]$, $R = (K-1) \log 256/K$, $K = 32$. Cosh uses anchored midpoint quantiles with $\tau = 4$ and exactly the same endpoints. Only the 30 interior coordinates differ.

Evaluation scores the final 128 target tokens at 32 fixed validation anchors. Seed 20260723 selects one end position from each of 32 equal-width stream bins with a 2048-token guard; the longest windows are disjoint. At anchor e , length H uses tokens $[e - H, e)$ and scores $[e - 128, e)$. The retained-range condition keeps R fixed; the target-matched condition uses $R(H) = (K-1) \log H/K$ while retaining each trained weight set and normalized z . Both use unit gain. Training seed is the replication unit.

Table 5: Exact-range anchored EVQ-COSH minus the matched geometric control tail-NLL contrasts. Mean $\pm$ SD is across three independent training seeds.

Length	Fixed training range				Target-matched range			
	Seed 42	Seed 137	Seed 256	Mean $\pm$ SD	Seed 42	Seed 137	Seed 256	Mean $\pm$ SD
256	+0.033	+0.029	+0.017	+0.026 $\pm$ 0.008	+0.033	+0.029	+0.017	+0.026 $\pm$ 0.008
512	-0.478	-0.271	-0.094	-0.281 $\pm$ 0.192	+0.061	+0.053	+0.067	+0.060 $\pm$ 0.007
1,024	-0.205	-0.191	-0.132	-0.176 $\pm$ 0.039	+0.182	+0.156	+0.344	+0.227 $\pm$ 0.102
2,048	-0.113	-0.181	-0.143	-0.146 $\pm$ 0.034	+0.279	+0.399	+0.701	+0.460 $\pm$ 0.218

C.3 CROSSED WEIGHTS AND RUNTIME TABLES

For seeds 137, 256, each Geo- or Cosh-trained checkpoint installs a factor-four table derived from either training allocation. The derivation uses $L_0 = 256$, $\Delta_j = 2j$ for $j = 0, \dots, 127$, and grouped causal weights $p_j = (2L_0 - 4j - 1)/(L_0(L_0 + 1)/2)$. Let X_k contain the weighted sine/cosine pair and X_{-k} the other pairs. Define

$$U_k = \frac{\|X_k - X_{-k} X_{-k}^\dagger X_k\|_F^2}{\|X_k\|_F^2}, \quad \bar{U}_k = \frac{U_k - \min_i U_i}{\max_i U_i - \min_i U_i}, \quad \omega'_k = \omega_k [1 - 3(1 - \bar{U}_k)^2/4].$$

The pseudoinverse cutoff is 10^{-10} ; zero normalization span gives $\bar{U}_k = 0$. Tables are FP32 with endpoints overwritten to ω_0 and $\omega_{K-1}/4$, and gain $1 + 0.1 \ln 4$. Evaluation uses the same anchors and tail-NLL metric at 512/1024 tokens. The 1024-token mean crossing is: Both seeds favor the

Table 6: Two-training-seed mean tail NLL at 1,024 tokens for the 151.9M weights-by-table crossing (lower is better). Each row freezes one trained weight set; each column installs one separately derived runtime table.

Frozen weights	Geo-derived table	Anchored-EVQ-COSH-derived table
Geo-trained	3.426	5.776
Anchored-EVQ-COSH-trained	4.455	3.479

table derived from their training allocation. The seed-specific interaction contrasts are 3.400/3.251 NLL at 1024 tokens and 2.663/2.564 at 512.

Table 7: **Exact-range factorial** (50.9M; 2 bases $\times$ 2 training lengths $\times$ 3 head dimensions = 12 configurations, 3 seeds). All arms share sampled extrema and log-span with Geo. Metric: weighted OOD NLL over $2 \times /4 \times /8 \times$, paired within configuration and seed. All arms were pre-specified. CIs are 10,000 configuration-level bootstraps; p is a two-sided exact sign-flip over all 2^{12} assignments, unadjusted across contrasts, so the claim rests on the cross-configuration direction (App. C).

Contrast	Δ NLL	95% CI	$\Delta < 0$ configs	p
EVQ-COSH $0.75 \times$ – Geo	−0.00912	[−0.018, −0.001]	8/12	0.082
EVQ-COSH rule – Geo	−0.00988	[−0.021, 0.002]	7/12	0.125
EVQ-COSH $1.25 \times$ – Geo	−0.01210	[−0.021, −0.003]	10/12	0.027
Matched exponential – Geo	−0.01062	[−0.021, −0.001]	9/12	0.071
<i>The matched exponential reproduces the non-uniform direction:</i>				
EVQ-COSH rule – exponential	+0.00074	[−0.006, 0.008]	7/12	0.836

C.4 MULTIPLE SHAPES AT FIXED SUPPORT

The 50.9M decoder has 6 layers, width 512, MLP width 2048 and vocabulary 50304. The 12 configurations cross bases 500K/1M, training lengths 256/1024 and 4/8/16 heads (head dimensions 128/64/32). Each of three seeds 42/137/256 trains each arm for 8,388,608 tokens, 128 updates, on a shared repeated WikiText-2 raw stream with a 1% supervised passkey mix. Execution is float32 on MPS. The preassigned arms are Geo, anchored Cosh at $0.75/1/1.25$ times $d_{\text{head}}/\sqrt{L}$, and an exponential $z_\gamma(t) = (e^{\gamma t} - 1)/(e^\gamma - 1)$ whose RMS node movement matches reference Cosh without task fitting. Every arm shares sampled endpoints. The OOD mean weights ratios $r \in \{2, 4, 8\}$ by $\log_2(r + 1)$; the bootstrap unit is configuration after averaging paired seeds. Training and evaluation use separate WikiText-2 raw train/validation splits, tokenized with GPT-NeoX-20B. Nonempty stripped rows are concatenated without extra special tokens and independently repeated/truncated to 8,388,608 training and 1,048,576 evaluation tokens. NumPy RandomState(9999) selects four distinct starts at each of $L, 2L, 4L, 8L$, in that order, identically across arms. Full next-token cross-entropy scores $4(rL - 1)$ targets and averages the four chunk NLLs before exponentiation. Structural configurations, rather than repeated text windows, are the bootstrap units.

Standard-grid replication. A separate short-training study uses the 151.9M architecture, $b = 500000$, head dimension 64, length 128 and 14,974,976 tokens. Standard-grid RoPE and span-matched native-grid Cosh are evaluated on 32 test anchors separate from 16 selection anchors. Three-seed mean tail NLL for Native/Cosh is 5.5046/5.4645 at 128, 5.6563/5.5434 at 512, 5.8456/5.6966 at 1K and 6.3928/6.2935 at 8K. All three seeds favor Cosh at 512/1K/2K/4K/8K; complete recorded curves are in the bundled numerical inputs. This protocol is separate from the longer fixed-support training comparison.

C.5 FROZEN FIXED-SUPPORT AND EQUAL-DISPLACEMENT INTERVENTIONS

Frozen OLMo-2-1B and Qwen-2.5-1.5B use 64 pairs, their native fast endpoint and the native slow endpoint divided by four. Matched-support arms share gain $1 + 0.1 \ln 4$, inputs, greedy decoder and precision. Uniform z is log-linear between these endpoints. The derived arm computes residual fractions U_k as above over native integer separations weighted by $L - d$, grouped in pairs at rounded weighted centroids, then uses $\omega'_k = \omega_k[1 - 3(1 - \bar{U}_k)^2/4]$ and exact endpoint overwrite. The coarse arm enumerates integer linear-ramp bounds, minimizing MSE to $(1 - \bar{U}_k)^2$ with lexicographic tie-breaking: bounds are (20, 22) for OLMo and (28, 31) for Qwen. These are geometric, label-free controls.

On OLMo’s nine held-out tasks (20/task, 16K), uniform/coarse/derived score 0.56/61.04/60.47%; Qwen’s four development tasks (20/task, 64K) score 57.75/64.00/66.50%. The sampled support is identical within each three-arm block. A separate OLMo six-task follow-up uses 24/48 inputs at 4/16K, $s = 4$, band $[14, 32]$ and gain $1 + 0.1 \ln 4$:

	MrPro	Official YaRN	MrUni	BM
4K	37.85	54.38	76.88	81.81
16K	2.78	6.94	32.12	51.32

BM and Uni have the same interior sum $(n-1)/2$ and full-table sum 40.5; Pro’s full-table sum is $37\frac{2}{3}$. Thus BM–Uni identifies a shape contrast at equal total displacement. The newer clean T–C comparison is reported separately in Appendix D.3.

C.6 SPECTRUM-PRESERVING COORDINATE REASSIGNMENT

A block permutation obeys $\mathcal{R}_{\pi\Omega}(d) = P\mathcal{R}_{\Omega}(d)P^{\top}$, so jointly permuting Q/K and frequencies preserves the bilinear kernel exactly. With weights fixed, permuting only interior frequency assignments changes model performance. The historical OLMo log-law table uses $L_0 = 4096$, $b = 500000$, $K = 64$, $s = 4$, gain $1 + 0.074 \ln 4$, and fixes slots 0, 63. Its 20-document PG-19 tail NLL changes from 3.10423 to 6.86493. Qwen-2.5-1.5B recomputes the same profile at $L_0 = 32768$, $b = 10^6$; on 20 inputs per task at 64K (single-key, multikey-2, multikey-3, tracking), macro changes from 70 to 0%. These fixed-spectrum interventions isolate coordinate assignment; their table is distinct from TailSpline. Source identities remain in the numerical supplement.

D FROZEN EVALUATION: PROTOCOLS AND COMPLETE RESULTS

D.1 SHARED MODELS, TABLES, SCORING AND UNCERTAINTY

Llama-3-8B-Instruct uses native length 8192, $K = 64$, base 500000 and band [18, 35]; OLMo-2-0425-1B-Instruct uses 4096, $K = 64$, base 500000 and band [14, 32]. T/P share frozen weights, sampled endpoints, outer bands and gain $1 + 0.1 \ln 4$, while using their respective static $s = 4$ interior tables. Tables are installed in FP32 after loading BF16 weights and are fixed across layers and input lengths. Evaluation uses Flash-SDPA and greedy decoding. No table is fitted using model outputs. The gain multiplies both cosine and sine, hence contributes g^2 to the rotated Q/K logit. Clean panels reuse the tables fixed for the classic runs; changing input selection does not refit the allocation.

Clean RULER pins revision c3f5e3b4f87f, selects source records in order, and uses neither content padding nor depth selection. Both arms use identical prompt tokens, references and generation caps: 128 for NIAH, 30 for tracking, 120 for common words, 50 for frequent words and 32 for QA. The endpoint is the official within-task score averaged equally across the 13 tasks. All outputs, including empty and capped answers, remain in the score. Intervals use 20,000 paired resamples within task and length, with tasks fixed; all reported intervals are 95% percentile intervals. RULER resamples source records, while natural QA resamples source-context clusters.

D.2 CLEAN RESULTS ACROSS LENGTH AND MODEL

Llama uses 650/650/2600 pairs at 8/16/32K; OLMo uses 2600 at 16K. All clean runs use batch 1 and 8192-token prefill chunks. Llama 32K inputs span 28270–32606 tokens; the fixed QA-1 supplement completes 200 source-order records without output-based selection. OLMo’s clean panel uses generation seed 20261101 and QA source offset 5600.

Llama 32K T–P is +11.72 points, interval [10.32, 13.11]. Family differences (retrieval/tracking/aggregation/QA) are +10.86/ + 24.00/ + 11.76/ + 9.00 points, with intervals [9.11, 12.61]/[20.60, 27.40]/[8.80, 14.77]/[4.25, 13.75]. Leaving out any one task keeps the macro difference between 9.53 and 13.07 points. At 16K, the Full-13 difference is +3.39 [1.53, 5.34]; family differences are +4.50/ − 2.40/ − 2.77/ + 8.00. The latter QA interval is [1, 15] and retrieval interval [2.50, 6.69].

OLMo clean 16K T–P is +41.42 points, interval [39.99, 42.81]. Its family differences are +60.19/ + 9.70/ + 16.13/ + 7.50 points, with intervals [58.30, 62.03]/[7.00, 12.60]/[13.86, 18.42]/[2.75, 12.25]. All thirteen task means increase. T/P end with EOS on 1251/1392 inputs and hit the cap on 1349/1208, with no empty outputs.

Table 8: **Complete clean task scores (%)**. N is original RoPE; T/P are TailSpline/MrPro. Llama 8/16/32K use 50/50/200 inputs per task; OLMo 16K uses 200.

Task	Llama 8K			Llama 16K		Llama 32K		OLMo 16K	
	N	T	P	T	P	T	P	T	P
cwe	97.60	95.40	89.60	83.80	80.00	1.65	0.30	7.55	3.30
fwe	87.33	9.33	4.00	25.33	34.67	61.83	39.67	51.83	23.83
MK_1	100.00	100.00	98.00	100.00	88.00	84.50	62.50	65.00	6.50
MK_2	96.00	100.00	94.00	98.00	86.00	47.00	9.00	43.00	0.50
MK_3	98.00	98.00	94.00	92.00	84.00	24.00	2.50	4.50	0.00
MQ	99.50	96.50	97.00	98.00	97.00	95.25	94.88	73.25	4.12
MV	97.50	97.00	92.50	98.50	97.50	93.12	97.62	77.12	2.25
S_1	100.00	100.00	100.00	100.00	100.00	100.00	97.00	97.50	11.50
S_2	100.00	100.00	100.00	100.00	98.00	99.50	99.00	100.00	18.00
S_3	100.00	100.00	100.00	100.00	100.00	100.00	94.00	64.00	0.00
qa_1	70.00	70.00	72.00	76.00	70.00	64.00	55.50	39.50	31.50
qa_2	32.00	46.00	44.00	60.00	50.00	37.00	27.50	25.50	18.50
vt	92.40	94.80	86.00	87.60	90.00	79.60	55.60	9.70	0.00

At native 8K, original RoPE uses the native table and unit gain. T/P/original RoPE score 85.1564/82.3923/90.0256%. T-P is +2.7641 [1.2923, 4.2538] points; T-original RoPE is -4.8692 [-6.5718, -3.1462], a 5.41% relative reduction. The difference concentrates in FWE: T/P have 45/48 empty responses, all on that task, while original RoPE has none. Across all 650 inputs, T/P/original RoPE EOS counts are 648/644/645 and cap counts 2/6/5. The other twelve tasks have a descriptive T-original RoPE mean of +1.225 points; the primary Full-13 endpoint retains FWE.

D.3 CLEAN EQUAL-DISPLACEMENT COMPARISON

C is the analytic profile in Appendix B.2. T/C share outer bands, endpoints, gain and total log-frequency displacement; the deployed FP32 $\sum m_k$ discrepancy is below 10^{-8} . Completed clean T/P generations are reused; C adds the same 650/2600 inputs at 16/32K with batch 1 and 8192-token prefill chunks. T/C score 86.0949/86.3462% at 16K, difference -0.2513 [-1.4051, 0.8949] points; at 32K they score 68.2660/66.1699%, difference +2.0962 [1.1134, 3.0769] points.

Table 9: **Complete clean shape comparison**. Scores (%) at 16K (50/task) and 32K (200/task).

Task	T16	C16	P16	T32	C32	P32
cwe	83.80	83.20	80.00	1.65	0.65	0.30
fwe	25.33	30.00	34.67	61.83	54.83	39.67
multikey_1	100.00	98.00	88.00	84.50	88.00	62.50
multikey_2	98.00	98.00	86.00	47.00	36.50	9.00
multikey_3	92.00	92.00	84.00	24.00	9.50	2.50
multiquery	98.00	98.50	97.00	95.25	95.25	94.88
multivalue	98.50	98.00	97.50	93.12	94.88	97.62
single_1	100.00	100.00	100.00	100.00	99.50	97.00
single_2	100.00	100.00	98.00	99.50	98.50	99.00
single_3	100.00	100.00	100.00	100.00	100.00	94.00
qa_1	76.00	76.00	70.00	64.00	69.50	55.50
qa_2	60.00	60.00	50.00	37.00	37.00	27.50
vt	87.60	88.80	90.00	79.60	76.10	55.60

At 32K, T and C both have EOS/cap rates 98.96/1.04%. Their full-pair ranks are 10.0773/10.0926, compared with P’s 10.2149; the task contrast therefore also distinguishes positional rank from learned use.

D.4 NATURAL QA AND LONGBENCH V2

The frozen Natural-QA pool has 631 questions from 524 source-context clusters across HotpotQA, 2WikiMQA, Qasper, NarrativeQA and MultiFieldQA-en (Bai et al., 2024). Membership was defined

using the OLMo > 4096 -token threshold and retained under Llama retokenization. Actual inputs are 4110–16319 for OLMo and 3694–16304 for Llama. Llama has 316 questions within its native window and 315 above it. Generation caps are 32 for HotpotQA/2Wiki, 128 for Qasper/NarrativeQA and 64 for MultiFieldQA. Batch is 1; Llama uses 8192-token chunked prefill and OLMo direct prefill.

The primary score includes all responses, averaging whole-response normalized token F1 within each task, then equally across tasks. For paired intervals, 20,000 bootstrap draws resample source-context clusters within task. Each draw divides sampled score sums by sampled question counts and retains equal task weights. There are no shared clusters across tasks.

Table 10: Complete Natural-QA631 task scores (F1, %). Each model uses the same question pool across T/P/Y.

Task	Llama-8B			OLMo		
	T	P	Y	T	P	Y
2wikimqa	36.96	35.19	35.61	26.29	21.55	24.72
hotpotqa	52.98	53.09	56.07	34.65	28.63	28.27
multifieldqa_en	51.91	51.33	50.65	35.80	34.49	32.29
narrativeqa	22.82	23.77	26.09	9.85	9.51	11.52
qasper	40.73	41.04	40.33	18.03	13.94	14.56

The three-arm report gives Llama T-P $+0.1957$ $[-1.4826, 1.8590]$ points; YaRN scores 41.7473%. Llama’s within/above-native T-P contrasts are $+0.6998$ $[-1.2478, 2.8073]$ and -0.9383 $[-4.1494, 1.9539]$. Document-equal sensitivity gives T/P 40.8141/40.8375%. OLMo’s overall difference is $+3.2994$ $[0.3912, 6.3651]$ points, with every T-P task mean positive; YaRN scores 22.2719%. Question-equal sensitivity gives $+3.9194$ $[1.1977, 6.6794]$. OLMo’s 4K–8K and 8K–16K strata contain 316/315 questions and give T-P differences of $+6.3989$ $[1.8747, 10.8706]$ and $+1.5095$ $[-4.9432, 7.6943]$ points. Llama T/P EOS/cap/empty counts are 625/624, 6/7, 0/0; generated-token totals are 6972/7019. OLMo has 567/581, 64/50, 0/0.

LongBench v2 uses 117 complete Llama inputs between 8K and 32K, with its direct-answer multiple-choice adapter, batch 1, matched caps, static $s = 4$ tables and 8192-token prefill chunks. Correct parsed options score one; incorrect/unparseable answers score zero. T/P answer 41/36 correctly, giving 35.0427/30.7692% and difference $+4.2735$ $[-3.3906, 11.9658]$ points under source-context cluster resampling. T alone succeeds on 13 inputs, P on 8, with 96 ties.

Table 11: **LongBench v2 complete-input subset**. Accuracy (%) with actual Llama-tokenized input lengths between 8K and 32K.

Domain	Inputs	TailSpline	MrPro
Code repository	3	33.33	0.00
In-context learning	7	14.29	14.29
Structured data	1	0.00	0.00
Dialogue history	12	41.67	16.67
Multi-document QA	33	39.39	30.30
Single-document QA	61	34.43	37.70

The 28 inputs at 8K–16K give T/P 39.29/35.71%; the 89 above 16K give 33.71/29.21%. Overall EOS/cap/empty counts are 88/91, 29/26, 0/0. Accuracy and token F1 remain separate metrics.

D.5 CLASSIC LENGTH CURVES AND SUPPLEMENTARY CONTROLS

Classic Llama and OLMo use 10 RULER inputs/task at $L, 2L, 4L$ (390/arm), with output-independent depth selection and prefix padding. Single-evidence prompts have two examples at each 10/30/50/70/90% depth; multi-evidence prompts alternate dispersed and middle-cluster targets. A repeated irrelevant-record prefix fills the input to cap minus 128 tokens. The shared task-specific decoding budgets above are retained. Preparation seeds are 20260924/20260925 (Llama/OLMo), with fixed length/task offsets; the bundled entrypoints specify the deterministic construction.

PPL uses 32 Proof-Pile arXiv test documents and 14 PG-19 test-book prefixes (Azerbayev et al., 2022; Rae et al., 2019). At each length it is the exponential of summed whole-prefix loss divided by scored tokens. Normalized trapezoidal log-length AUC has weights (1/4, 1/2, 1/4); PPL AUC integrates PPL itself. The document bootstrap resamples document IDs jointly across lengths; task bootstrap resamples within task-length cells.




Figure 6: **Classic task curves.** Each model retains its own recorded depth-selected, prefix-padded protocol.

Llama Full-13 AUC T/P is 78.8013/75.6026%, difference +3.1987 [0.6538, 5.7853] points; OLMo is 66.60/17.36%, difference +49.23 [45.34, 53.04]. Llama PPL at 8/16/32K is T 5.295283/4.530318/4.110622, P 5.286429/4.518111/4.161831. Its NIAH AUC improves by 3.75 points; QA AUC changes by -1.25. Llama EOS/cap/empty counts are T/P 388/384, 2/6, 30/37.

The original classic T-C diagnostic gives AUC 78.8013/79.2115%, difference -0.4103 [-2.6282, 1.8205] points, with batch 1/2. The clean comparison uses matched batch-one evaluation. The classic 130-input native reference scores 91.8846%, versus T 89.7436%, difference -2.1410 [-6.1410, 1.9231] points. Whole-prefix 8K PPL on 46 documents is original RoPE/T/P 5.275919/5.295283/5.286429; T-original RoPE is +0.019364 with paired document interval [0.006602, 0.034861].

D.6 DIRECT BASELINES ACROSS MODEL FAMILIES AND SCALE

Qwen-2.5-3B-Instruct (Yang et al., 2024) and GLM-4-9B-0414 have native reference length 32768, with bases $10^6/10^4$ and 64/32 rotary pairs. GLM uses GQA with 32 query and 2 key/value heads; 64 of each head’s 128 coordinates rotate. We retain its adjacent-pair convention and nonrotary subspace. T/P use bands [23, 40] for Qwen and [17, 30] for GLM.

Official static YaRN uses its public 32/1 correction and amplitude rule, without fine-tuning; its own boundary rounding and blend are retained. All methods share each panel’s prompts, scoring and greedy decoder. Llama/OLMo NIAH uses 200 source-order inputs per task (1600 pairs); PPL uses all 46 matched Proof-Pile/PG-19 documents. Their Full-13 three-arm panels use the same first 10 inputs per task from the large T/P panels, adding only the missing YaRN arm. Qwen/GLM use 10 inputs/task at 128K (130 pairs); their earlier 5/task subsets are included, not additional replications. Their PPL uses five matched InfiniteBench LongBook documents. Table 12 reports the retrieval-family means and PPL; Tables 14–15 report complete Full-13 task scores and contrasts.

PPL is the exponential of token-pooled NLL; intervals concern the mean paired document NLL difference. No scores are pooled across tokenizers or corpora. Weights are BF16, batch one and Flash-SDPA. Llama/OLMo use 8192-token NIAH prefill chunks, GLM 65536, and Qwen direct prefill; LM uses direct prefill. NIAH intervals resample within task; document intervals resample documents, using 10000 draws. Full-13 uses 20000 paired within-task draws. Llama’s single-needle tasks are near ceiling; its multikey T/P/Y means are 51.83/24.67/27.00%. Multivalue/multiquery means are 94.19/96.25/97.13%. OLMo improves in all three groups: single-needle 87.17/9.83/9.83%, multikey 37.50/2.33/3.17%, and multivalue/multiquery 75.19/3.19/5.56%. These task profiles describe where the gains occur, separately from the structural analysis.

Table 12: **Retrieval and language modeling with static $s = 4$ tables.** NIAH uses 200/task for Llama/OLMo and the 10/task Full-13 retrieval subset for Qwen/GLM. PPL uses 46 matched documents for Llama/OLMo and 5 for Qwen/GLM; corpora and tokenizers remain model-specific. Below, NIAH intervals resample within task and NLL intervals resample documents; NLL differences are document means. Full-13 intervals accompany their task tables. T/P/Y are defined in Table 1.

Model	Length	NIAH $\uparrow$ (%)			PPL $\downarrow$		
		T	P	Y	T	P	Y
Llama	32K	80.42	69.56	70.97	4.1106	4.1618	4.1495
OLMo	16K	65.55	5.36	6.27	8.5362	12.8270	12.1410
Qwen	128K	79.69	76.88	76.56	2.3272	2.3293	2.3278
GLM	128K	52.81	38.44	35.00	27.3001	29.9507	30.0744

Paired differences and marginal 95% intervals

Model	Contrast	NIAH (pp)	Interval	Δ NLL	Interval
Llama	T-P	+10.86	[9.06, 12.62]	-0.01238	[-0.01408, -0.01078]
Llama	T-Y	+9.45	[7.78, 11.11]	-0.00941	[-0.01088, -0.00803]
OLMo	T-P	+60.19	[58.28, 62.08]	-0.40724	[-0.43628, -0.38052]
OLMo	T-Y	+59.28	[57.36, 61.16]	-0.35227	[-0.38037, -0.32606]
Qwen	T-P	–	–	-0.00093	[-0.00198, 0.00044]
Qwen	T-Y	–	–	-0.00026	[-0.00111, 0.00110]
GLM	T-P	–	–	-0.09266	[-0.13809, -0.02102]
GLM	T-Y	–	–	-0.09678	[-0.14876, -0.02092]

Table 13: Complete large-panel NIAH scores (% , 200/task). Passkey is single_1, already included in NIAH-8.

Task	Llama-8B			OLMo		
	T	P	Y	T	P	Y
multikey_1	84.50	62.50	63.00	65.00	6.50	9.00
multikey_2	47.00	9.00	17.00	43.00	0.50	0.50
multikey_3	24.00	2.50	1.00	4.50	0.00	0.00
multiquery	95.25	94.88	96.88	73.25	4.13	5.50
multivalue	93.13	97.63	97.38	77.13	2.25	5.63
single_1	100.00	97.00	100.00	97.50	11.50	11.00
single_2	99.50	99.00	98.50	100.00	18.00	18.50
single_3	100.00	94.00	94.00	64.00	0.00	0.00

Frozen transfer to 70B. The prequantized Llama-3-70B-Instruct checkpoint uses NF4 double quantization and BF16 computation. Its tokenizer, vocabulary, chat template and special tokens match the 8B evaluation. T/P directly reuse the 8B tables, band [18, 35] and common gain, with no allocation search on the 70B model. Batch-one 32K evaluation uses direct prefill. Full-13 uses the same 10/task source subset, Natural-QA the same 631 questions, and PPL the first five matched documents. T/P PPL is 2.4954/2.5075; the pooled NLL contrast is -0.00484 . This is a within-checkpoint NF4 comparison. T wins 10 Full-13 task means, ties one and loses two; EOS/cap rates are 99.23/91.54% and 0.77/8.46%. Natural-QA F1 is 49.2846/48.3924%, with T-P interval $[-0.5991, 2.4379]$ points under source-context resampling.

Table 14: Complete paired Full-13 scores (% , 10/task). These are the direct-baseline panels in Table 1; paired task-stratified 95% intervals appear below each model.

Task	Llama-8B			OLMo		
	T	P	Y	T	P	Y
cwe	6.00	0.00	0.00	7.00	2.00	3.00
fwe	70.00	50.00	40.00	66.67	13.33	20.00
multikey_1	80.00	70.00	60.00	70.00	0.00	10.00
multikey_2	80.00	10.00	10.00	30.00	0.00	0.00
multikey_3	30.00	0.00	0.00	10.00	0.00	0.00
multiquery	92.50	92.50	97.50	72.50	2.50	2.50
multivalue	92.50	100.00	95.00	65.00	0.00	0.00
single_1	100.00	100.00	100.00	100.00	20.00	40.00
single_2	100.00	100.00	90.00	100.00	30.00	10.00
single_3	100.00	100.00	100.00	50.00	0.00	0.00
qa_1	50.00	60.00	50.00	50.00	30.00	10.00
qa_2	30.00	10.00	10.00	50.00	40.00	30.00
vt	84.00	66.00	74.00	8.00	0.00	2.00
Macro	70.38	58.35	55.88	52.24	10.60	9.81
T-P 95% (pp)	[6.44, 17.68]			[35.38, 47.88]		
T-Y 95% (pp)	[8.77, 20.37]			[36.17, 48.71]		

Table 15: Complete paired Full-13 scores (% , 10/task). These are the direct-baseline panels in Table 1; paired task-stratified 95% intervals appear below each model.

Task	Qwen			GLM			70B NF4	
	T	P	Y	T	P	Y	T	P
cwe	3.00	4.00	6.00	1.00	6.00	2.00	70.00	58.00
fwe	56.67	43.33	50.00	60.00	56.67	60.00	86.67	83.33
multikey_1	100.00	90.00	90.00	60.00	50.00	50.00	80.00	90.00
multikey_2	50.00	40.00	50.00	0.00	0.00	10.00	90.00	20.00
multikey_3	10.00	10.00	0.00	0.00	0.00	0.00	30.00	20.00
multiquery	95.00	92.50	87.50	50.00	30.00	45.00	92.50	87.50
multivalue	82.50	82.50	85.00	62.50	47.50	45.00	97.50	97.50
single_1	100.00	100.00	100.00	80.00	10.00	0.00	100.00	70.00
single_2	100.00	100.00	100.00	70.00	80.00	50.00	100.00	90.00
single_3	100.00	100.00	100.00	100.00	90.00	80.00	100.00	90.00
qa_1	20.00	30.00	30.00	20.00	30.00	20.00	60.00	70.00
qa_2	40.00	30.00	50.00	10.00	0.00	0.00	20.00	10.00
vt	80.00	56.00	58.00	2.00	0.00	0.00	96.00	6.00
Macro	64.40	59.87	62.04	39.65	30.78	27.85	78.67	60.95
T-P 95% (pp)	[-0.41, 9.49]			[4.04, 13.69]			[11.65, 23.78]	
T-Y 95% (pp)	[-2.78, 7.49]			[6.37, 17.32]			-	

Publisher-recommended deployment on Kanana. Kanana-1.5-8B-Instruct-2505 has native length 32768, 64 rotary pairs and base 8,000,000 (Kakao, 2025). At 64K, T/P use $s = 2$, band [20, 35] and gain $1 + 0.1 \ln 2$. Official runtime YaRN uses factor 4.4, 64/2-turn corrections and gain $1 + 0.1 \ln 4.4$, as recommended by its model card. T/P isolates allocation; T/Y compares complete deployment recipes on the same checkpoint and inputs. All arms use BF16, batch one, Flash-SDPA and greedy decoding. Full-13 comprises 130 paired prompts. Table 16 reports every task; EOS rates for T/P/Y are 70.77/69.23/65.38%, with no empty outputs.

At 128K, TailSpline uses $s = 4$ and the same native band; official YaRN keeps its factor-4.4 recipe. InfiniteBench English-QA contains all 118 complete prompts fitting the tokenizer and input budget, from 23 book contexts. Inputs span 69680–130852 tokens; the output cap is 40 tokens. Question-mean official F1 is 19.59/17.85% for T/Y; source-context-equal F1 is 19.63/18.04%. T wins 31 questions, Y wins 30, and 57 tie; EOS/cap counts are 80/38 versus 76/42, with no empty outputs. These two target lengths use separately constructed TailSpline tables.

Table 16: Kanana at 64K: complete Full-13 scores (% , 10/task) and paired within-task intervals. T/P use $s = 2$ and Y the publisher-recommended factor 4.4.

Task	T	P	Y
single_1	100.00	100.00	100.00
single_2	100.00	100.00	100.00
single_3	100.00	100.00	100.00
multikey_1	80.00	80.00	80.00
multikey_2	60.00	50.00	50.00
multikey_3	70.00	50.00	40.00
multivalue	65.00	57.50	60.00
multiquery	22.50	32.50	25.00
vt	64.00	70.00	82.00
cwe	43.00	51.00	13.00
fwe	90.00	93.33	93.33
qa_1	80.00	70.00	70.00
qa_2	70.00	60.00	40.00
Macro	72.65	70.33	65.64
T-P: +2.32 pp, 95% interval [-3.13, 8.04]			
T-Y: +7.01 pp, 95% interval [0.21, 13.81]			

D.7 COMPLETE-BOOK QA AND SOURCE-CONTEXT UNCERTAINTY

Qwen and GLM use InfiniteBench English book QA (Zhang et al., 2024), with complete prompts between 100K and 128K tokens, no input truncation, and a 40-token answer cap. Each panel contains 35 questions from seven source-book contexts, five questions per context. Actual input ranges are 100907–127946 for Qwen and 100844–127811 for GLM. All three arms retain their static $s = 4$ table, batch one, BF16, Flash-SDPA and 65536-token chunked prefill.

The official English-QA adapter computes normalized token F1 against the accepted references; it is not RULER’s answer-containment score. All generated answers remain in the mean. The paired bootstrap keeps questions from a sampled source context together, using 20,000 draws and seeds 20261401/20261402 for T-P/T-Y. Table 17 reports marginal and simultaneous intervals for each three-arm panel, including the five-task Natural-QA comparisons. Task-macro Natural-QA resamples source clusters within task while retaining question-weighted means; complete-book QA uses its single task and retains all questions per book. GLM’s second source pool contains 77 questions from 15 books, none in pool I. Its complete inputs span 69660–95486 tokens; the same $s = 4$ table and decoder are retained. T/P/Y F1 is 26.8672/22.8626/23.5015%. The two pools retain separate means and source-context intervals.

Table 17: Natural-QA and book-QA contrasts (pp). Paired source-context bootstrap uses 20,000 draws. The last column gives Bonferroni simultaneous intervals for the two comparisons within each three-arm panel; the 70B panel has one comparison. II identifies GLM’s second book pool.

Model	Contrast	Difference	Marginal 95%	Simultaneous 95%
Llama-8B	T-P	+0.20	[-1.48, 1.86]	[-1.74, 2.12]
Llama-8B	T-Y	-0.67	[-2.87, 1.38]	[-3.22, 1.63]
OLMo	T-P	+3.30	[0.39, 6.37]	[-0.02, 6.83]
OLMo	T-Y	+2.65	[0.04, 5.49]	[-0.35, 5.92]
Llama-70B NF4	T-P	+0.89	[-0.60, 2.44]	[-0.60, 2.44]
Qwen	T-P	+3.24	[0.26, 7.51]	[-0.07, 7.89]
Qwen	T-Y	+0.69	[-3.66, 5.21]	[-3.99, 5.86]
GLM	T-P	+3.33	[-3.28, 8.98]	[-4.17, 9.70]
GLM	T-Y	+1.29	[-6.53, 8.61]	[-8.05, 9.29]
GLM II	T-P	+4.00	[-1.73, 9.68]	[-2.57, 10.61]
GLM II	T-Y	+3.37	[-1.58, 8.42]	[-2.38, 9.17]

Qwen T/P/Y EOS counts are 29/28/30, with cap counts 6/7/5; GLM has EOS counts 31/29/26 and cap counts 4/6/9. Neither panel has empty outputs. GLM pool II has 70 EOS and 7 capped outputs in each arm, with no empty responses. Relative to P, Qwen T has higher per-question F1 on 13 questions, lower on 6, and equal on 16.

D.8 HIGHER EXTENSION FACTORS: COMPLETE TASK PROFILES

Llama uses $s = 16$ at 128K and Qwen uses $s = 8$ at 256K. T/P share each condition’s frozen weights, public band, amplitude and prompts. All use BF16, batch one and Flash-SDPA. Llama uses 8192-token chunked prefill; Qwen uses direct prefill. These are separate length/scale conditions, not extra samples from the $s = 4$ comparisons.

Table 18: Complete higher-factor results. Generation scores are percentages; PPL uses its own document set. Intervals are paired 95% intervals.

Model	Endpoint	N	T	P	T-P interval
Llama 128K	Full-13	130	22.05	1.28	[15.96, 25.64] pp
	Proof-Pile PPL	10	2.5870	2.6714	[−0.03505, −0.02918] NLL
	En.QA F1	41	13.21	20.17	[−11.87, −2.07] pp
	En.Dia accuracy	100	12.00	8.00	[−1.00, 10.00] pp
Llama-70B NF4 128K	Proof-Pile PPL	10	2.2449	2.3129	–
Qwen 256K	Single-NIAH	15	93.33	40.00	[33.33, 66.67] pp
	LongBook PPL	5	30.3437	23.9283	[−1.0561, 1.7769] NLL

Llama Full-13 uses ten inputs per task. Qwen uses five inputs for each of single-1/2/3, whose T/P scores are 100/100, 100/20 and 80/0%; this three-task endpoint is distinct from NIAH-8. Llama’s natural panels retain complete 100K–131K contexts and the official InfiniteBench metrics; En.QA contains eight book contexts. Generation scores use paired task or source-context resampling; PPL uses paired documents. The 70B NF4 $s = 16$ PPL contrast is -0.02981 pooled NLL on ten documents; no document interval is supplied. It uses 32768-token prefill chunks and the shared gain $1 + 0.1 \ln 16$. Llama’s Full-13 and PPL improve while its book-QA F1 decreases; Qwen’s single-needle retrieval improves while its five-book PPL is higher. The results distinguish retrieval, language modeling and natural-task behavior at these extension factors.

E NATIVE-WINDOW ALLOCATION: CONSTRUCTION AND COMPLETE RESULTS

E.1 REFERENCE OBJECTIVE AND NUMERICAL CONSTRUCTION

NCP improves task performance within the native context by reallocating the interior frequencies while fixing both sampled endpoints. Let $\phi_k = (L - 1)\omega_k$ and $\omega'_k = \omega_k e^{-u_k}$. The reference problem is a binary target–distractor comparison with unit amplitudes, aligned target phase and uniformly random distractor phase:

$$r(\theta) = \frac{1}{2\pi} \int_0^{2\pi} \log(1 + \exp(\cos \psi - \cos \theta)) d\psi, \quad (43)$$

$$\mathcal{R}(\phi) = 2 \int_0^1 (1 - t) r(\phi t) dt. \quad (44)$$

The triangular distance weight is the distribution of absolute separation between two uniform positions on a continuous unit interval. The aligned reference comparison supplies a public design objective. The Transformer task scores below evaluate the resulting fixed construction.

NCP minimizes

$$\min_u \sum_k \left[\mathcal{R}(\phi_k e^{-u_k}) + \frac{9}{8} u_k^2 \right], \quad \omega'_k = \omega_k e^{-u_k}, \quad (45)$$

with fixed $\lambda = 9/4$ (quadratic coefficient $\lambda/2 = 9/8$), subject to

$$u_0 = u_{K-1} = 0, \quad 0 \leq u_k \leq 2/9, \quad \log \frac{\omega'_k}{\omega'_{k+1}} \geq \frac{1}{2} \log \frac{\omega_k}{\omega_{k+1}}. \quad (46)$$

The endpoint constraints keep support fixed; the gap condition preserves ordering and at least half each original adjacent log gap. Writing $r(\theta) = a_0 + \sum_{j \geq 1} a_j \cos(j\theta)$ gives

$$\mathcal{R}(\phi) = a_0 + \sum_{j \geq 1} a_j \operatorname{sinc}^2(j\phi/2), \quad \operatorname{sinc}(x) = \frac{\sin x}{x}, \quad \operatorname{sinc}(0) = 1. \quad (47)$$

This follows by integrating each cosine against $2(1-t)$. With $q(\phi) = d\mathcal{R}/d\log \phi$, an interior scalar solution satisfies $\lambda u = q(\phi e^{-u})$. We compute r on a 1024-point angular grid, retain 32 Fourier modes and solve the scalar stationary equations by bisection, using zero shift when the native gradient is nonpositive. The two endpoint shifts are reset to zero. The implemented 32-mode scalar objective is strictly convex. For $f(x) = \operatorname{sinc}^2(x/2)$, its log-curvature is $h(x) = 2 \cos x - 6 \sin x/x + 8(1 - \cos x)/x^2$ and $|h(x)| \leq 4$: for $0 \leq x \leq 4$, the integral representation gives $|h(x)| \leq x/3 + x^2/6$; for $x \geq 4$, it gives the bound $\sqrt{4 + 36/x^2} + 16/x^2 \leq 3.5$. The retained coefficients satisfy $\sum_{j=1}^{32} |a_j| < 0.554403$, so scalar curvature is at least $9/4 - 4 \sum_j |a_j| > 0.03238$. Thus the stationary root, or the zero-shift boundary when its derivative is nonnegative, is the unique box minimizer. On the deployed OLMo grid the log-gap constraints are inactive, so these scalar minimizers also solve the constrained finite-series problem.

The deployed grid has 64 pairs, base 500000 and native length 4096. It changes 46 FP32 entries, reaches maximum relative slowdown 16.7164% and retains a minimum log-gap ratio of 0.7681. The maximum scalar stationarity residual is below 10^{-11} . The mean reference risk decreases from 0.614082 to 0.609641; the penalized objective is 0.611853. These are reference-model quantities, separate from model task scores. The table and unit gain are fixed across layers and inputs; construction reads no checkpoint weights, activations or evaluation outputs. The source archive includes the original constructor and an exact FP32 table reproduction check.

E.2 NATIVE-CONTEXT TASK RESULTS AND COMPANION CONTROLS

The OLMo native 4K development panel contains 60 source-order inputs per RULER task, 780 total, without content padding. It is the same frozen panel used for the companion allocation studies. NCP and original RoPE share the checkpoint, unit gain, inputs, BF16 precision, greedy decoder, batch 1 and 4096-token prefill chunks. Their scores are 70.7821/69.3739%; the paired task-equal difference is +1.4081 points, interval $[0.2009, 2.5833]$. NCP has higher per-input scores on 62 inputs, original RoPE on 31, with 687 ties. Every task remains in the aggregate.

Table 19: Complete native-window Full-13 scores (%). Original/new OLMo panels use 60/10 inputs per task at 4K. Every task stays in its primary mean.

Task	Original (60/task)		New (10/task)	
	Native	NCP	Native	NCP
cwe	7.83	4.67	2.00	1.00
fwe	57.78	60.00	53.33	56.67
multikey_1	81.67	78.33	90.00	90.00
multikey_2	91.67	95.00	100.00	100.00
multikey_3	58.33	61.67	80.00	80.00
multiquery	89.17	90.83	85.00	92.50
multivalue	95.42	98.33	97.50	100.00
single_1	100.00	100.00	100.00	100.00
single_2	100.00	100.00	100.00	100.00
single_3	98.33	98.33	100.00	100.00
qa_1	68.33	61.67	70.00	70.00
qa_2	35.00	30.00	50.00	60.00
vt	18.33	41.33	28.00	48.00

Companion half-turn contraction and its reverse score 69.5085/67.8974%. Their exact mean contrast is +1.6111 points with paired interval $[0.2713, 2.9722]$; contraction relative to original RoPE is

+0.1346 points with interval $[-0.6603, 0.9445]$. A separate historical checkpoint-calibrated V1 table scores 72.3846%, a +3.0107-point contrast with interval $[1.2949, 4.7052]$. V1’s calibration distinguishes it from NCP’s public-parameter construction. Point differences here are computed from arm means; the companion report’s bootstrap intervals are retained.

E.3 ADDITIONAL TASK EVALUATION

The new OLMo Full-13 panel has ten source-order records per task (130 pairs), fixed before evaluation. NCP reuses the original public table and unit gain; its task-equal difference is +3.2564 $[0.8077, 5.9231]$ points. The interval uses 20000 paired within-task draws, seed 20260916. Table 19 keeps the original and new task panels separate.

Native-QA contains 99 complete inputs within OLMo’s native window: 24 2WikiMQA, 5 HotpotQA and 70 Qasper questions. Both arms share prompts, BF16 weights, greedy decoding and task-specific generation caps. Normalized whole-response token F1 is averaged over questions within task, then equally over tasks. Paired source-cluster resampling within task retains question weighting, using 20000 draws and seed 20260917.

The task F1 scores (Native/NCP, %) are 37.96/35.05 on 2WikiMQA, 45.00/45.00 on HotpotQA and 35.27/35.32 on Qasper.

Native/NCP F1 is 39.4116/38.4554%, difference $-0.9562 [-4.4410, 1.5666]$ points. The earlier source-equal sensitivity instead gives $-0.8549 [-4.4495, 1.9234]$ points; these estimands remain separate. Natural-QA and controlled RULER diagnostics measure different tasks.

E.4 LANGUAGE MODELING AND BENEFIT FROM EARLIER CONTEXT

We evaluate 128 nonoverlapping windows from 103 documents in the PG-19 and Proof-Pile test corpora. Each 4097-token window defines 4096 next-token inputs. We score only its final 256 target tokens, once with full history and once with 512 preceding tokens plus the same teacher-forced targets. Shared tokens retain identical relative positions; the checkpoint, targets and score positions are paired. Loss is averaged within document and then equally over documents. Paired document intervals use 20000 draws, seed 20260915.

Table 20: Same-target NLL (nat/token), document-equal across 103 documents and 128 windows. Context benefit is recent-history minus full-history NLL; positive benefit differences favor NCP.

Context	Native	NCP	NCP-Native	95% interval
Recent history	3.009922	3.005077	-0.004845	$[-0.005670, -0.004007]$
Full history	2.915269	2.902483	-0.012786	$[-0.015168, -0.010440]$
Context benefit	0.094653	0.102594	+0.007941	$[0.005529, 0.010346]$

NCP improves full-history NLL by 0.012786 nat/token and recent-history NLL by 0.004845. Its additional benefit from earlier context is 0.007941 $[0.005529, 0.010346]$ nat/token. This measures context use at the same prediction targets, separately from generated-task accuracy.

F LEARNING AND EXTRAPOLATION WITH COSH

F.1 MULTI-HEAD LATENT ATTENTION

We report allocation at one fixed rotary budget in a 432M-parameter transformer following the MLA architectural pattern (DeepSeek-AI, 2024a), with 24 layers, width 1024, 16 heads, and per-head width 64. Its decoupled rotary subspace has $d_{\text{rope}}=32$ ($K=16$ frequencies); the remaining $d_{\text{nope}}=32$ dimensions are non-rotary, and the shared KV latent rank is 256. The model trains from scratch on 500M FineWeb-Edu tokens (Penedo et al., 2024) at $L_{\text{train}}=8192$ and base 500K over seeds 42, 43, and 88. EVQ-COSH uses the pre-specified $\tau = 1.414$ setting. At 16K, Geo PPL is 141.1/132.5/142.8 for seeds 42/43/88, whereas EVQ-COSH is 93.7/92.7/100.3; all three paired directions agree.

Evaluation corpus and aggregation. Evaluation uses a common held-out FineWeb-Edu stream of approximately five million GPT-2 tokens. At each length (8/16/20/24/28/32K), all methods and seeds share eight text chunks, each scored on all $L - 1$ next-token targets. Per-seed PPL is the exponential of the mean chunk NLL; reported means and sample standard deviations are across the three training seeds.

For the supporting overlay, let $\lambda_k = 2\pi/\omega_k$ and $t_k = \text{clip}((\lambda_k - L_{\text{train}})/(31L_{\text{train}}), 0, 1)$. The MLA wavelength-blend operator sets $\omega'_k = \omega_k/[1 + t_k(s - 1)]$: it preserves wavelengths below the training length, applies the full factor s above $32L_{\text{train}}$, and interpolates between them.

Table 21: MLA validation (432M, $d_{\text{rope}}=32$, 3-seed mean $\pm$ std). Rows marked “MLA wavelength-blend” use this run’s wavelength-blend operator defined above.

Method	PPL@8K	PPL@16K	PPL@24K	PPL@32K
Geo	35.4 $\pm$ 0.9	138.8 $\pm$ 5.5	241.5 $\pm$ 2.6	323.7 $\pm$ 5.2
Geo + MLA wavelength-blend ($s = 4$)	35.5 $\pm$ 0.9	117.9 $\pm$ 6.5	204.9 $\pm$ 6.9	278.5 $\pm$ 6.3
EVQ-COSH	35.8 $\pm$ 0.8	95.6$\pm$4.1	204.9 $\pm$ 14.7	291.6 $\pm$ 20.5
EVQ-COSH + MLA wavelength-blend ($s = 4$)	35.8 $\pm$ 0.8	71.1$\pm$4.1	153.2$\pm$12.5	236.6$\pm$15.9

EVQ-COSH reduces 16K PPL by 31.1% (138.8 $\rightarrow$ 95.6), while 8K PPL changes by +0.9%. All three training seeds favour EVQ-COSH at 16K. The EVQ-COSH plus MLA wavelength-blend row has the lowest mean PPL at 16K and beyond, showing that the allocation also composes with this wavelength-dependent extension.

F.2 SHARED-START CONTINUATION AND MATCHED ADAPTATION

The 750M pair continues a common 2K geometric checkpoint for 500M tokens at 4K, with 3% passkey and 10% downstream mix (seed 42). Cosh uses $\tau = 1.5$, $K = 32$, inclusive nodes $u_k = k/(K - 1)$ and $\omega_k = \exp[-(K - 1) \log(500000)Q_{1.5}(u_k)/K]$, retaining Geo endpoints. Both arms update all weights. Greedy answer-token exact match generates the expected passkey tokens and closing marker; it does not require a following EOS.

Table 22: 750M continued-pretraining check (seed 42; 2K $\rightarrow$ 4K; 500M continuation tokens). Passkey retrieval is teacher-forced NLL-gap and is saturated at 8K; AR exact is greedy autoregressive exact match over 40 trials per length. The pooled AR row combines 2K, 4K and 8K (120 trials: Geo 80/120, EVQ-COSH 111/120).

Metric	Geo	EVQ-COSH
PPL@2K	25.9	26.2
PPL@4K	22.0	22.3
PPL@8K	23.4	19.6
PPL@16K	45.1	24.4
NLL-gap retrieval @8K	100%	100%
Answer-token exact @8K	0%	77.5%
Pooled answer-token exact (2K/4K/8K)	66.67%	92.5%

For MLA and 750M continuation, AdamW uses (0.9, 0.95), $\epsilon = 10^{-8}$, weight decay 0.1, clipping 1, no dropout, peak LR 2×10^{-4} , 2% linear warmup and cosine decay to 0.1 of peak. Batch sizes are 6 (MLA) and 14 (continuation: microbatch 7, accumulation 2). Text packing is contiguous without cross-document masking. MLA uses unanchored midpoint Cosh; the continuation fixes sampled endpoints.

Larger-scale training and mature-model adaptation. The 1.485B OLMo pair uses the same public step-zero initialization, scientific recipe and reconstructed seed-6198 data prefix: 1000 updates, global batch 512, context 4096, and 2,097,152,000 counted tokens. It has 128-dimensional heads and base 500000. Cosh uses $\tau = 2$, $K = 64$, native nodes $u_k = k/K$, and $\omega'_k = b^{-Q_2(u_k)}$; there is

no affine endpoint anchoring in this study. The released geometric checkpoint uses the upstream distributed trainer; Cosh uses a single-GPU implementation of the same recipe. Evaluation pairs 128 held-out PG-19 book prefixes. At 4/8/16K, Cosh-minus-Geo NLL differences are $+0.0381/-0.0437/-0.1351$, with document-paired intervals $[0.0332, 0.0428]/[-0.0494, -0.0380]/[-0.1420, -0.1281]$.

The Llama-3-8B-Instruct comparison instead applies matched rank-64, alpha-128 Q/K/V/O LoRA for 300 steps, seed 42, learning rate 10^{-4} , BF16 and 8192-token sequences. Native-LoRA uses the original frequency grid; Cosh-LoRA uses midpoint Cosh at $\tau = 1.414$. Rows, order, optimizer, schedule, adapter capacity and evaluator match. Adaptation uses 7476 length-filtered LongAlpaca training rows and 152 validation rows, with effective batch 8. Evaluation uses 24 frozen temporal packs from arXiv, the Federal Register and Stack Overflow; losses are pooled within domain and averaged equally across domains before exponentiation. All 24 packs favor Cosh at both extended lengths. Table 23 keeps these two protocols distinct.

Table 23: **Cosh extrapolation across scale and adaptation.** PPL within matched comparisons; each study retains its own training and evaluation recipe.

Study	Length	Geometric/Native	Cosh
OLMo 1.485B training	2K	177.99	191.36
	4K	161.19	167.45
	8K	163.88	156.87
	16K	182.73	159.64
Llama 8B matched LoRA	8K	6.817	10.068
	16K	108.958	24.068
	32K	991.475	127.911

Matched OLMo routing continuation. OLMo-2-0425-1B-Instruct uses fixed standard-grid Cosh with $\tau = 2$, $b = 500000$ and rank-64, alpha-128 Q/K/V/O LoRA. Native/Cosh share checkpoint, training rows, budgets and evaluation inputs. Starting from each arm’s 300-step routing parent, a matched 100-step query-gap continuation uses peak LR 5×10^{-5} and 20 warmup steps; a 32-step answer-plus-EOS stage uses 10^{-5} and 4 warmup steps. Both use seed 20260728, AdamW (0.9, 0.95), zero weight decay, clipping 1, BF16, batch 8 (micro 4, accumulation 2), and cosine LR decay.

Each cycle has two counterfactual-routing updates and one natural-text update: 67/33 updates in the query-gap stage and 22/10 in the final stage. The 1024 training pairs and 128 calibration pairs use seeds 20293725/20304725 and different query templates. Examples contain eight key-value blocks; changing the queried value changes the target answer. Answer cross-entropy is combined with a counterfactual margin of 1, weighted by 0.5, on positions where correct and swapped answer tokens differ. The final stage supervises three answer tokens and immediate EOS, with EOS weight one; prompt and padding tokens carry no answer loss. Natural replay uses ordinary next-token loss.

Position schedule. Physical inputs stay at most 4096 tokens. Let q be the semantic query-block start, a the answer start, h the end-exclusive source span and n the active sequence length. Both counterfactual variants use $p_i = i + o \mathbf{1}\{i \geq q\}$. Each query-gap update uses four pair offsets: one alternating between 0 and $[1, 4096]$, one in $[4097, 8192]$, and two in $[8193, 12289]$. In the final stage, realized gap targets instead comprise one alternating between contiguous and $[4096, 8191]$, one in $[8192, 12287]$, and two in $[12288, 16079]$. For a noncontiguous target g , the offset is

$$o = \min\{g - (a - h), 12289, 16384 - \min(n, 4095)\}.$$

The resulting gap must remain in its target band and active position IDs below 16384. Natural replay stays contiguous.

Evaluation uses real contiguous numeric-NIAH prompts, 100 per length, and requires the complete expected answer followed by terminal EOS. Native/Cosh scores are 95/100 versus 100/100 at 4K, 18/100 versus 98/100 at 8K, and 0/100 versus 60/100 at 16K. All outputs terminate with EOS. The Cosh long-gap subsets score 48/50 at 8K and 31/66 at 16K. This is matched learned adaptation, with explicit long-phase exposure, distinct from the frozen allocations.

Separate selective-Q/K task continuation. A second matched OLMo study continues the Q/K components of rank-64, alpha-128 Stage-A Q/K/V/O adapters while retaining inherited V/O adapters. It uses 300 updates, batch 8, LR 5×10^{-5} , 20 warmup steps, BF16 and a phase/phase/natural schedule, with explicit 4/8/16K position exposure in ratio 1 : 1 : 2 and answer-plus-EOS supervision. Physical sequences remain at most 4K. The 1536 training and 128 validation examples have no QA-identity overlap with the 200 held-out 2WikiMultiHopQA rows per length. Greedy generation uses a 32-token cap. Inputs retain the official prompt and add answer-filtered distractor contexts to reach approximately 4/8/16K; this is a task-family length-transfer evaluation. Native/Cosh F1 is 25.99/24.84, 0.07/21.48 and 0/8.57%, respectively; exact match is 22/21.5, 0/17.5 and 0/4%. These separately trained task adapters provide generation evidence alongside the Llama language-modeling result.